

Distributional Synthetic Control for Text Outcomes with an Application to Local Government Speech

Amar Venugopal

Stanford University Department of Economics

This version: October 3, 2026

[Link to latest version](#)

Abstract

Applied researchers increasingly study problems in which outcomes of interest include collections of speeches, reviews, and other documents. Standard approaches to making these objects tractable for estimation consist of employing some low-dimensional, interpretable featurization of text, such as word counts, topic shares, or sentiment scores. In this paper, I develop a pipeline for hypothesis generation and evaluation in text-as-outcome settings in causal inference with nonrandom treatment assignment. In the first stage, I develop a distributional synthetic control estimator in which each unit-period outcome is represented as a distribution over document embeddings and weights are estimated to balance the treated unit's pre-treatment distributions against a mixture of donor unit distributions. In an asymptotic regime in which the number of units and time periods is fixed while the number of documents in each cell of the panel matrix grows, I establish consistency and asymptotic normality of these estimated weights. In the second stage, I show that balancing on distributions simultaneously balances on a wide class of summary functions of the text and propose a hypothesis-generation procedure which discovers interpretable aspects of text that drive variation between the observed treated corpus and its constructed counterfactual. I show that a sample-splitting approach allows for consistent estimation and valid inference on the resulting treatment effect estimates, providing analytic standard errors which take into account uncertainty on the counterfactual constructed in the first stage. In an application to California city council meeting public comment periods, I find that the elimination of remote participation options in the aftermath of the COVID-19 pandemic in several cities in northern California lead to increased civility and reduced hostility in public comments.

Contact: amarvenu@stanford.edu. I am deeply grateful to my advisors, Guido Imbens, Jann Spiess, Jiafeng Chen, and Matthew Gentzkow for their invaluable guidance and mentorship. I thank Lea Bottmer, Amir Feder, Han Hong, Lihua Lei, Olivia Martin, Ashesh Rambachan, David Ritzwoller, Yiqing Xu, and attendees of the Stanford Econometrics seminar series for their advice and feedback. I am fortunate to have my research supported by the Stanford Data Science Scholars Program and the B.F. Haley and E.S. Shaw Fellowship for Economics from the Stanford Institute for Economics and Policy Research. Computational support was provided by the Data, Analytics, and Research Computing (DARC) group at the Stanford Graduate School of Business (RRID:SCR_022938).

1 Introduction

Settings with text outcomes are becoming increasingly ubiquitous in the social sciences. Scholars study political speech (Gentzkow, Shapiro, and Taddy, 2019), corporate disclosures (Hoberg and Phillips, 2011), central bank communications (Shapiro and Wilson, 2022), judicial opinions (Ash, Chen, and Naidu, 2026), and consumer reviews (Luca, 2011). In each of these settings, the content and style of the text is economically meaningful.

A standard approach to make these unstructured outcomes amenable to study is to first construct some interpretable representation for each document, with such representations often including word counts, sentiment scores, or topic shares. Researchers can then estimate causal effects on the resulting scalar (or low-dimensional vector) outcomes, producing estimates that have intuitive meanings. However, this strategy also has a key limitation: every such featurization of the text imposes a restriction on the possible causal effects the analysis can detect. For example, a study that represents text documents by topic shares will be unable to detect changes in sentiment or tone within a topic. Conversely, using a sentiment score will leave the researcher blind to changes in topical content. In general, the choice of featurization limits the researcher to detection of causal effects hypothesized to be relevant *ex ante*; if the true nature of the causal effect is not known in advance, the chosen featurization may fail to capture the true scope of effects on the text.

The modern natural language processing literature has given rise to representations of text known as *embeddings*, which are typically the outputs of neural networks and have been shown to preserve semantic content of text in a numeric vector representation in some Euclidean space (Mikolov et al., 2013; Pennington, Socher, and Manning, 2014; Vaswani et al., 2017). These representations much more plausibly capture the full scope of potential causal effects on the text, but are themselves uninterpretable, posing challenges to causal estimation.

Recent work in hypothesis generation for unstructured data seeks to balance these two objectives by using a “complete” representation, such as embeddings, to detect systematic differences between collections of documents, which can then be mapped into interpretable hypotheses that are evaluated on held-out data (Ludwig and Mullainathan, 2024; Ludwig, Mullainathan, and Rambachan, 2025). Modarressi, Spiess, and Venugopal (2025) extend this approach to causal settings to discover and evaluate treatment effects in randomized controlled trials (RCTs), in which the observed control corpus is a valid counterfactual for the treated corpus. In observational settings, however, untreated documents are not automatically plausible counterfactuals for treated documents due to possible unobserved confounding. An algorithm applied directly to treated and untreated text may therefore conflate selection and treatment effects.

In this paper, I develop a two-stage pipeline for hypothesis generation and validation in text-as-outcome causal inference settings with nonrandom treatment assignment in panel data. In the first stage, I develop a distributional synthetic control estimator which allows for construction of a counterfactual corpus of text data, against which the observed treated corpus can be compared in

the second stage. The second stage then generates interpretable hypotheses about how treatment changes language and estimates effects along these learned dimensions. I establish consistency and asymptotic normality for the resulting treatment effect estimates, providing a unified inference framework that takes into account first-stage uncertainty in counterfactual construction.

The first-stage counterfactual estimator treats each unit-period outcome as a probability distribution over document embeddings. Donor weights are chosen to minimize the pre-treatment discrepancy between the treated unit’s sequence of distributions and the corresponding weighted mixtures of donor distributions. Applying the same weights to contemporaneous donor distributions post-treatment produces a synthetic counterfactual distribution, which is itself a mixture over the constituent donor unit distributions. This construction generalizes conventional synthetic control: when attention is restricted to the mean of a single scalar summary, the distributional objective reduces to the usual synthetic control objective for that outcome. Simultaneously, the distributional formulation preserves information that would be discarded by any single scalar featurization. Because the counterfactual is a mixture of donor distributions rather than an averaged embedding, it can also be sampled to obtain an ordinary corpus of documents. The resulting counterfactual is therefore both statistically valid and directly readable.

I provide identification, consistency, and inference results for this counterfactual estimator. If the treated unit’s untreated distribution can be expressed in every period as a common convex mixture of the donor distributions, the population objective recovers the true post-treatment counterfactual. When the synthetic control model is misspecified (that is, the treated unit no longer lies within the convex hull of the donor units), the estimator targets the population mixture of donor unit distributions which is closest in trajectory to the treated unit’s pre-treatment distribution. I term this estimand the pseudo-counterfactual, and provide an approximation error bound which characterizes its distance from the true counterfactual. The asymptotic analysis holds the number of units and periods fixed while the number of documents within each unit-period cell grows. This sampling dimension, similar to that utilized in [Liu \(2025\)](#), is absent from most conventional aggregate synthetic control ([Abadie, Diamond, and Hainmueller, 2010](#); [Almeida et al., 2025](#)). In my setting, I show that this asymptotic regime yields consistency of the estimated weights and induced counterfactual distribution and permits analytic inference.

The second stage converts the distributional effect into interpretable scalar effects without requiring the researcher to specify those dimensions in advance. I split the post-treatment documents into discovery and evaluation samples. A hypothesis generation algorithm, implemented in the application with a large language model (LLM), compares documents sampled from the treated and synthetic discovery distributions and returns a summary function $\hat{g}$ which assigns a scalar score to each document and is intended to describe a difference between the two discovery distributions. This function is then fixed and applied to the held out evaluation documents, where its treatment effect is estimated as the difference between the treated and synthetic mean scores. The resulting

estimator is consistent even if the generated function does not converge to some deterministic limit; the theory instead requires the learning procedure to be honest and its possible outputs to belong to a well-behaved function class. Under bounded misspecification, control of the distributional discrepancy yields uniform bounds on bias over a corresponding class of summary functions, implying that disciplining distributional distance implicitly disciplines distance for many possible summaries.

Inference on the resulting treatment effects must account for both finite sampling of post-treatment documents and estimation of the synthetic-control weights. It must also address a dependence that does not arise with a prespecified outcome: the same estimated weights used in the final effect estimator may have helped construct the corpus from which $\hat{g}$ was learned. I provide a solution that splits the pre-treatment data as well as the post-treatment data, so the weights used for discovery and evaluation are estimated from independent samples. This provides the additional benefit of allowing the hypothesis generation algorithm to ingest pre-treatment observations, allowing for iterative exploration to ensure the chosen summary exhibits good pre-treatment balance between the treated unit and the synthetic control. Feasible standard errors on the treatment effect estimates then incorporate uncertainty from the first-stage estimated counterfactual.

I highlight the method’s practical utility by applying it to a dataset of public comments from California city council meetings. This application builds on [Martin and Venugopal \(2026\)](#), which assembles transcripts from approximately 25,000 meetings in 115 California municipalities between 2015 and 2024 and studies how the staggered removal of remote access options for public participants impacts participant demographics. Their results show that eliminating remote participation reduces the number of speakers but does not clearly change observed speaker composition. Now, I answer a different question: does changing the cost and mode of participation alter what residents say once they speak? I restrict attention to unagendized public comments to reduce concerns about agenda-driven selection and, through two calibrated simulations, show that my distributional estimator correctly identifies true weights and correctly accounts for unobserved confounding when generating hypotheses, significantly improving over a naive, equal-weighting benchmark. I then construct synthetic controls for multiple units which remove remote participation and find requiring public commenters to attend in-person leads to significantly increased civility and reduced hostility in the speech of public comments.

This paper contributes to four strands of related literature. First, it extends synthetic control beyond scalar aggregate outcomes ([Abadie and Gardeazabal, 2003](#); [Abadie, Diamond, and Hainmueller, 2010, 2015](#)). As such, it is related to work using high-dimensional micro-level data or multiple outcomes to improve synthetic controls ([Robbins, Saunders, and Kilmer, 2017](#); [Shi et al., 2022](#); [Sun, Ben-Michael, and Feller, 2025](#); [Liu, 2025](#); [Tian, Lee, and Panchenko, 2026](#)). Its closest antecedent in this strand of the literature is [Gunsilius \(2023\)](#), which develops a distributional synthetic control estimator for scalar outcome distributions via estimating convex averages of quantile functions. We extend this approach by considering settings where within-cell obser-

vations are high-dimensional documents, making quantile functions ill-defined. We instead use a mixture-distribution framework similar to that proposed as a secondary specification in [Gunsilius \(2023\)](#), extending the approach to consider an inferential target that includes summaries selected after inspecting the counterfactual corpus.

Second, this paper contributes to a literature on balancing in causal inference, in which weights are learned to match distributions of covariates between treated and control units ([Zubizarreta, 2015](#); [Doudchenko and Imbens, 2016](#); [Kallus, 2017](#); [Athey, Imbens, and Wager, 2018](#); [Santra, Chen, and Park, 2026](#)). Motivated by this connection, the primary choice of distributional discrepancy used for synthetic control estimation is maximum mean discrepancy (MMD), which measures the difference in kernel mean embeddings between two distributions for a given choice of kernel k . As such, this approach is closely related to that of the kernel balancing literature in particular ([Hazlett, 2020](#); [Hazlett and Xu, 2018](#)).

Third, this paper contributes to the growing literature on causal inference with text and other unstructured data ([Chen and Sarkar, 2020](#); [Feder et al., 2022](#); [Egami et al., 2022](#); [Fong and Grimmer, 2023](#); [Imai and Nakamura, 2024](#); [Sarkar, 2026](#); [Feldman et al., 2026](#); [Freyaldenhoven, 2026](#)). It relates most closely to the text-as-outcome subsection of this literature, with its closest antecedent being [Modarressi, Spiess, and Venugopal \(2025\)](#), which uses sample splitting to generate and validate interpretable outcomes learned from text in randomized trials. I extend this approach to nonrandom treatment assignment by constructing the missing counterfactual corpus and by carrying uncertainty from this first stage into downstream inference. This paper also relates to recent advances in machine learning methods, including LLMs, as tools for hypothesis generation ([Ludwig and Mullainathan, 2024](#); [Ludwig, Mullainathan, and Rambachan, 2025](#)). I extend this literature by applying LLMs as hypothesis generation algorithms in causal settings, and in particular to constructed counterfactuals of unstructured data.

Finally, the application in this paper contributes to two strands of applied literature. The first of these is the study of local government dynamics, including participant demographics and meeting attendance, as well as studies of how remote attendance options impact participation ([Warshaw, 2019](#); [Einstein, Palmer, and Glick, 2019](#); [Einstein et al., 2023](#)). The closest antecedent in this literature is [Martin and Venugopal \(2026\)](#), in which the authors examine the effect of removing remote participation options on participant demographics in California city council meetings. I utilize this data to instead examine how requiring in-person participation changes the *speech* of public comment. Consequently, this study also relates to a literature examining the impacts of remote work more generally ([Bloom et al., 2015](#); [Brucks and Levav, 2022](#); [Bloom, Han, and Liang, 2024](#)). I extend this work by assessing the impacts of remote communication on local government dynamics specifically.

The remainder of the paper proceeds as follows. [Section 2](#) defines potential text distributions and gives conditions for exact and approximate identification of the counterfactual mixture. [Sec-](#)

tion 3 develops the sample estimator, establishing consistency of the estimated weights and induced pseudo-counterfactual as well as inference for the synthetic control weight vector. Section 4 turns to downstream summaries, first establishing consistency and inference for fixed summaries, then showing that post-treatment sample splitting permits consistent estimation for data-driven summaries without requiring convergence of the learned rule. Section 5 applies the method to California local government speech, and Section 6 concludes.

2 Setting & Identification

In this section, we establish notation and define the estimand for counterfactual estimation via distributional synthetic control. We then provide conditions under which this estimand is equal to the true post-treatment counterfactual of the treated unit, and show that a generalization of the synthetic control objective function to the distributional setting yields correct weights in this setting.

Suppose that we have a collection of units indexed by $i = 1, \dots, N$ and time periods $t = 1, \dots, T$. For each unit and in each time period, we observe a collection of K_{it} documents a_{itk} , which are drawn from a corresponding distribution: $a_{itk} \sim F_{it}$. Suppose further we have access to an embedding technology $\mathcal{E} : \mathcal{T} \rightarrow \mathcal{X}$ that maps from text documents in set $\mathcal{T}$, which is understood to contain the supports of all F_{it} , to some subset of d -dimensional Euclidean space:

$$x_{itk} = \mathcal{E}(a_{itk}) \in \mathcal{X} \subseteq \mathbb{R}^d$$

In practice, this embedding function $\mathcal{E}$ will typically be an API call to some pretrained embedding model.¹ There is an extensive literature which highlights the ability of these embedding functions to capture the semantic meaning of text (Mikolov et al., 2013; Pennington, Socher, and Manning, 2014; Devlin et al., 2019; Reimers and Gurevych, 2019). Indeed, embeddings are utilized in our most cutting-edge models of text, including large language models (Vaswani et al., 2017; Brown et al., 2020). We will therefore adopt a standard approach in the natural language processing literature, treating embedding similarity as a proxy for semantic similarity of texts (Muennighoff et al., 2023). Accordingly, our analysis will be defined on the distribution of document embeddings. Henceforth, we will use F_{it} to instead denote the distribution induced by applying the embedding technology $\mathcal{E}$ to documents from unit i in time period t , so that $x_{itk} \sim F_{it}$.

Treatment is then modeled to induce an alteration to these distributions F_{it} from which our observations are drawn. As such, we define the distributional analog to the potential outcomes

¹Further, since text documents are finite combinations of a countable set of words, this embedding map $\mathcal{E}$ is automatically measurable.

framework (Rubin, 1974) as *potential distributions*:

$$F_{it}(w_{it})$$

Here, w_{it} denotes the treatment status for unit i in time period t . Analogous to the canonical synthetic control approach in Abadie, Diamond, and Hainmueller (2010), suppose that unit N is treated in time period T , and all other unit-time cells are untreated:

$$W_{it} = \begin{cases} 1 & i = N \text{ and } t = T \\ 0 & \text{else} \end{cases}$$

We assume throughout that there are no treatment spillovers, and that each unit’s potential distribution depends only on that unit’s contemporaneous treatment assignment (e.g., the stable unit treatment value assumption, or SUTVA, is satisfied). As a consequence, we observe draws from $F_{it}(0)$ for all $(i, t) \neq (N, T)$, and $F_{NT}(1)$ else. We assume that these draws are independent and identically distributed within each cell, and that there are no correlations across cells.

Assumption CS (Comment sampling). *For each unit i and period t , the observed texts $\{a_{itk}\}_{k=1}^{K_{it}}$ are independent and identically distributed and these collections are mutually independent across cells (i, t) . Under the fixed, measurable embedding map $\mathcal{E}$, the corresponding embeddings $x_{itk} = \mathcal{E}(a_{itk})$ are therefore also drawn iid:*

$$x_{itk} \stackrel{\text{iid}}{\sim} F_{it}(w_{it})$$

The collections $\{x_{itk}\}_{k=1}^{K_{it}}$ are also mutually independent across cells.

The key objects of interest for causal effect estimation are the potential distributions of the treated unit in the treated time period, for which we define the following shorthand notation:

$$P_1 := F_{NT}(1) \quad P_0 := F_{NT}(0)$$

Note that we observe draws from P_1 , but not from P_0 ; estimation of this counterfactual will be the primary objective of this distributional synthetic control procedure.

We model the untreated potential distributions of the treated unit as a mixture over contemporaneous untreated potential distributions of the donor units. In order to implement this, we require a metric for inter-distributional distance. Our primary specification for this metric will be maximum mean discrepancy (MMD). Proposed by Gretton et al. (2006), MMD_k is an integral probability metric which measures the distances between two distributions by their maximum distance in expected value of functions expressible within the unit ball of the reproducing kernel

Hilbert space (RKHS) induced by a kernel k :

$$\begin{aligned}\text{MMD}_k(P, Q) &= \sup_{\|f\|_{\mathcal{H}_k} \leq 1} \mathbb{E}_P[f(X)] - \mathbb{E}_Q[f(X)] \\ &= \|\mathbb{E}_P[k(X, \cdot)] - \mathbb{E}_Q[k(X, \cdot)]\|_{\mathcal{H}_k}\end{aligned}$$

An RKHS is defined as a function space endowed with an inner product. As such, functions that are elements of an RKHS have notions of “length” and “angle”, analogous to standard vectors in Euclidean space. The kernel k defines a function $k(x, \cdot)$ for each point x which describes how “similar” x is to the second argument. In our case, we are interested in describing the distance between embedding vectors, and so $x \in \mathcal{X}$. We can then examine all possible linear combinations of these kernel functions:

$$f(\cdot) = \sum_{j=1}^m a_j k(x_j, \cdot) \quad g(\cdot) = \sum_{\ell=1}^r b_\ell k(z_\ell, \cdot)$$

Then $f, g \in \mathcal{H}_k$ and the inner product between them is defined as:

$$\langle f, g \rangle_{\mathcal{H}_k} = \sum_{j=1}^m \sum_{\ell=1}^r a_j b_\ell k(x_j, z_\ell)$$

The Hilbert space is then completed to include limits of all Cauchy sequences. The space is said to be “reproducing” because evaluating any such function $f \in \mathcal{H}_k$ corresponds to taking an inner product with the kernel function:

$$f(x) = \langle f, k(x, \cdot) \rangle_{\mathcal{H}_k}$$

The expectation of these kernel functions, taken with respect to some distribution, is termed the “kernel mean embedding” of the corresponding distribution. The second equality in the definition of MMD_k therefore shows that this inter-distributional distance can be expressed as the norm of the difference in the distributions’ kernel mean embeddings. Applied to our setting, we will henceforth let $\mu_{it}(\cdot)$ denote the kernel mean embedding of $F_{it}(W_{it})$:

$$\mu_{it}(\cdot) = \mathbb{E}_{F_{it}(W_{it})}[k(X, \cdot)] \in \mathcal{H}_k$$

Note that $\mu_{it}(\cdot)$ is a function in the RKHS induced by the kernel k . For notational brevity, we will generally drop the argument and simply refer to this function as μ_{it} .

The choice of MMD_k as an inter-distributional distance metric is also motivated by the kernel balancing literature (Hazlett and Xu, 2018; Hazlett, 2020). In the kernel balancing problem, weights are chosen so that treated and control groups are approximately balanced on a rich, kernel-induced function space. As a consequence, kernel balancing can be equivalently viewed as minimizing

the MMD_k between the treated and control covariate distributions, and its adoption here is an extension to a setting in which we seek to balance trajectories of distributions of high-dimensional text embeddings. Throughout, we will make the following regularity assumptions on the chosen kernel k :

Assumption k-Reg (Kernel regularity). *The kernel k is fixed, measurable, characteristic, and positive definite. It further has a separable RKHS $\mathcal{H}_k$ and is bounded by $\bar{k}$:*

$$\sup_{x \in \mathcal{X}} k(x, x) \leq \bar{k} < \infty$$

Here, characteristicness of a kernel refers to the induced map from distributions to kernel mean embeddings being injective; positive definiteness means that every kernel matrix produced by kernel k is positive semidefinite; having a separable RKHS means that any function $g \in \mathcal{H}_k$ can be approximated arbitrarily closely in RKHS norm by some member of a countable collection of functions $f_1, f_2, \dots \in \mathcal{H}_k$. Measurability and boundedness allow expectations of the kernel applied to random observations to be well-defined and finite. Examples of kernels that satisfy these requirements include the Gaussian RBF kernel, the rational quadratic kernel, the Laplace kernel, and the Matérn kernel (Steinwart and Christmann, 2008; Sriperumbudur et al., 2010). Note that for all of these, fixedness of the kernel (as stated in the assumption) requires fixedness of their relevant parameters (e.g., the bandwidth for the Gaussian RBF and Laplace kernels). In practice, these can be estimated using an independent pilot subsample of data, or specified ex-ante.²

Using MMD_k as our notion of distributional distance, we define the population synthetic-control objective as the sum of squared MMD_k distances between the untreated potential distribution of the treated unit and a mixture over contemporaneous untreated potential distributions of donor units.

$$L(w) = \sum_{t < T} \text{MMD}_k^2 \left(F_{Nt}(0), \sum_{i=1}^{N-1} w_i F_{it}(0) \right) \quad (1)$$

Note that this formulation is a generalization of the standard scalar synthetic control objective function in Abadie and Gardeazabal (2003) and Abadie, Diamond, and Hainmueller (2010). In particular, let $g : \mathcal{X} \rightarrow \mathbb{R}$ be some scalar summary defined on a document’s embedding vector, and replace MMD_k with $\mathcal{D}^{\text{scalar}}$ defined as:

$$\mathcal{D}^{\text{scalar}}(P, Q) = \mathbb{E}_P[g(X)] - \mathbb{E}_Q[g(X)]$$

Equation 1 then precisely reduces to the standard scalar synthetic control objective function, where the scalar outcome is the average value of the summary g across the collection of documents in

²If kernel parameters are selected based on an independent pilot sample and held fixed thereafter, all following assumptions and results are understood to be conditional on that sample, so that the selected kernel and associated population targets are fixed.

each cell.

Returning now to our distributional setting, let $\mathcal{W}$ denote the simplex in $\mathbb{R}^{N-1}$. We can then define the set of optimal weight vectors in the population:

$$\mathcal{W}_0 = \arg \min_{w \in \mathcal{W}} L(w)$$

Our estimand is then the mixture distribution implied by these weights, which in general need not be equal to the true counterfactual. In particular, our general estimand is the *pseudo-counterfactual*:

$$\tilde{P}_0 := \sum_{i=1}^{N-1} w_i^0 F_{iT}(0) \quad w^0 \in \mathcal{W}_0$$

In a setting where $\mathcal{W}_0$ is not a singleton set, in order for a mixture distribution with weights estimated on pre-treatment data to be a valid pseudo-counterfactual post-treatment regardless of the choice of optimal weights, we must impose a temporal stability condition on the induced mixtures.

Assumption PP (Pre-period equivalence implies post-period equivalence). *For any $w, w' \in \mathcal{W}$, if*

$$\sum_{i=1}^{N-1} w_i F_{it}(0) = \sum_{i=1}^{N-1} w'_i F_{it}(0)$$

for all $t < T$, then

$$\sum_{i=1}^{N-1} w_i F_{iT}(0) = \sum_{i=1}^{N-1} w'_i F_{iT}(0).$$

Note that by virtue of the temporal stability in **Assumption PP** and our kernel regularity conditions in **Assumption k-Reg**, all population minimizers $w^0 \in \mathcal{W}_0$ induce the same pseudo-counterfactual $\tilde{P}_0$; see **Appendix A** for details. Furthermore, note that if $\mathcal{W}_0$ is a singleton, then the pseudo-counterfactual is unique trivially and **Assumption PP** is unnecessary.

In general, we will consider the case of $\tilde{P}_0$ lying some nonzero distance away from P_0 : $\text{MMD}_k(\tilde{P}_0, P_0) > 0$. This corresponds to a setting in which the synthetic control model is misspecified; the treated unit does not lie within the donor convex hull. Exact equality of $\tilde{P}_0$ and P_0 is achieved in the setting of correct specification, for which the following, in conjunction with **Assumption PP**, is a sufficient condition:

Assumption CM (Counterfactual mixture distribution). $\exists w^* \in \mathcal{W}$ such that

$$F_{Nt}(0) = \sum_{i=1}^{N-1} w_i^* F_{it}(0)$$

for every t .

This ensures that there is some weight vector $w^* \in \mathcal{W}$ that produces a mixture distribution over donor units that exactly matches the distribution for the treated unit in each period. One possible justification of this assumption is via generalization of the factor model structure used to motivate synthetic control in [Abadie, Diamond, and Hainmueller \(2010\)](#). Suppose that there exists some set of B latent topics, with time evolution in how they are discussed. The distribution Q_{bt} captures how discussion around topic b looks in period t , while λ_{ib} governs the prevalence of discussion of topic b in unit i .

$$F_{it}(0) = \sum_{b=1}^B \lambda_{ib} Q_{bt}$$

Then, the standard synthetic control assumption that the treated unit's loadings λ_N lie in the convex hull of the donor units' loadings implies [Assumption CM](#). However, this is simply one possible justification; see [Lemma 3](#) in [Appendix A](#) for further details.

This brings us to our first result: under correct specification, the exact mixture weights minimize the population objective function and identify the true counterfactual.

Proposition 1 (True weights minimize population objective). *Suppose Assumptions [CM](#) and [k-Reg](#) hold. Then, $L(w^*) = 0$ and*

$$w^* \in \arg \min_{w \in \mathcal{W}} L(w)$$

If additionally [Assumption PP](#) holds, we have

$$w^0 \in \arg \min_{w \in \mathcal{W}} L(w) \implies F_{NT}(0) = \sum_{i=1}^{N-1} w_i^0 F_{iT}(0)$$

This result establishes conditions under which the weights that produce the true counterfactual are precisely those that minimize the objective function in [Equation 1](#), allowing for nonuniqueness of these optimal weights. Therefore, Assumptions [CM](#) and [PP](#) together are sufficient conditions for $\tilde{P}_0 = P_0$. If we impose an additional affine-independence condition on the trajectory of kernel mean embeddings, then we can additionally recover unique optimal weights in the general case, where [Assumption CM](#) need not hold.

Assumption IDT (Affinely independent donor trajectories). *For every $h \in \mathbb{R}^{N-1}$,*

$$\mathbf{1}^\top h = 0 \quad \text{and} \quad \sum_{i=1}^{N-1} h_i F_{it}(0) = 0 \quad \text{for every } t < T \implies h = 0$$

Where $\mathbf{1}$ denotes the $(N-1)$ -dimensional vector of 1s.

Proposition 2 (Unique optimal weights). *Suppose Assumptions [k-Reg](#) and [IDT](#) hold. Then, $\mathcal{W}_0 = \{w^0\}$.*

This concludes the identification for the distributional synthetic control estimator. We have shown that its estimand is the mixture of donor distributions $\tilde{P}_0$ that uses mixture weights which minimize the distance between the treated unit and mixture of donor units' trajectories in the pre-treatment period. Furthermore, if the true counterfactual can be expressed as a mixture of donor distributions, $\tilde{P}_0 = P_0$ and, under some additional stability conditions, all minimizers of [Equation 1](#) induce the correct mixture P_0 in the post-treatment period. We have also provided sufficient conditions for when the optimal weight vector in the population is unique.

3 Distributional Synthetic Control Estimator

Let us now turn our attention to the estimation of the synthetic control model in finite samples. In this section, we will establish conditions under which we achieve consistency of both the estimated synthetic control weights and their induced mixture distribution. We will also provide a characterization of the asymptotic distribution of the weight vector.

Throughout, we will denote empirical distributions as

$$\hat{F}_{it} = \frac{1}{K_{it}} \sum_{k=1}^{K_{it}} \delta_{x_{itk}}$$

and define the sample analog objective function as

$$\hat{L}(w) = \sum_{t < T} \text{MMD}_k^2 \left(\hat{F}_{Nt}, \sum_{i=1}^{N-1} w_i \hat{F}_{it} \right)$$

and the sample estimator as

$$\hat{w} \in \arg \min_{w \in \mathcal{W}} \hat{L}(w)$$

For the remainder of the paper, ties in the weight-estimation problems are broken according to some fixed, measurable selection rule, such as selecting the minimizer with smallest L_2 norm.

3.1 Asymptotic Regime

In the panel data literature, to which standard synthetic control belongs, there are two possible axes along which we can imagine our data growing asymptotically: $N \rightarrow \infty$ and $T \rightarrow \infty$. The choice of asymptotic regime has proved fraught when it comes to providing inference results for synthetic control, with the bulk of current methods relying on placebo tests rather than analytic results ([Almeida et al., 2025](#)). In this paper, we have at our disposal a third axis along which we can imagine our sample growing: $K_{it} \rightarrow \infty$. This regime is closely related to that used in [Liu \(2025\)](#), albeit with a different inferential target. This will be our asymptotic regime of choice for the remainder of this paper, and all subsequent convergence results will be predicated on cell-level

document sample sizes growing arbitrarily large.

This choice of regime comes with some key benefits, chief among them being that we can produce results for consistency and inference of synthetic control weights that are intractable in typical asymptotic regimes. It is important to note, however, that the sources of uncertainty addressed by this asymptotic regime are different than those addressed by more standard regimes, particularly $N \rightarrow \infty$. In particular, our study will consider the size of the panel (N, T) to be fixed. Consequently, this asymptotic regime captures uncertainty based on sampling documents in each unit-period cell, with an inferential target of the treatment effect relative to the pseudo-counterfactual. The inferential target will equal to the true treatment effect only if our synthetic control model is correctly specified, $\tilde{P}_0 = P_0$. This is natural in this setting, where each cell of our panel is a distribution, but if instead our outcome were a scalar aggregate such as GDP, sampling error would be extremely small and this regime would imply that ordinary synthetic control has no uncertainty.

This means that the quality of the synthetic control fit in the population is fixed; if $\text{MMD}_k(\tilde{P}_0, P_0) = \epsilon_0$ for $\epsilon_0 > 0$, then ϵ_0 will remain unchanged as $K_{it} \rightarrow \infty$, as the convex hull of donor units remains static in the population. If instead we considered an asymptotic regime where $N \rightarrow \infty$, then we could imagine the addition of new donor units expanding the convex hull until it eventually encompasses P_0 , leading to $\epsilon_0 \rightarrow 0$. We will be unable to make any such claim in this study, and will correspondingly present results for both cases of population model fit throughout this paper: bounded misspecification ($\epsilon_0 > 0$) and correct specification ($\epsilon_0 = 0$).

With this asymptotic regime established, we define minimum per-cell counts across the entire panel, in the pre-treatment period only, and in the post-treatment period only, respectively:

$$K_{\min} = \min_{i,t} K_{it} \quad K_{\min}^{\text{pre}} = \min_{i,t < T} K_{it} \quad K_{\min}^{\text{post}} = \min_i K_{iT}$$

Note that $K_{\min} \rightarrow \infty$ implies that both $K_{\min}^{\text{pre}} \rightarrow \infty$ and $K_{\min}^{\text{post}} \rightarrow \infty$. In order to ensure well-behaved sample growth for downstream inference results, we introduce the following assumption:

Assumption BSG (Balanced sample growth). *For every $i = 1, \dots, N$ and $t < T$, as $K_{\min}^{\text{pre}} \rightarrow \infty$,*

$$\frac{K_{it}}{K_{\min}^{\text{pre}}} \rightarrow \kappa_{it}^{\text{pre}} \in [1, \infty)$$

Additionally, as $K_{\min} \rightarrow \infty$,

$$\frac{K_{iT}}{K_{\min}} \rightarrow \kappa_{iT}^{\text{post}} \in [1, \infty)$$

This assumption ensures that the samples in each pre-treatment and post-treatment cell do not grow at extremely unbalanced rates.

3.2 Consistent Estimation

In this subsection, we establish consistency of the estimated weights and induced counterfactual estimate. In doing so, we will leverage canonical results for extremum estimators (Newey and McFadden, 1994), of which $\hat{w}$ is an example.

Theorem 1 (Consistency of estimated weights and pseudo-counterfactuals). *Suppose Assumptions CS and k -Reg hold. Let $d(\cdot, \cdot)$ denote Euclidean point-to-set distance:*

$$d(\hat{w}, \mathcal{W}_0) = \inf_{w \in \mathcal{W}_0} \|\hat{w} - w\|$$

Then, as $K_{\min}^{\text{pre}} \rightarrow \infty$, $d(\hat{w}, \mathcal{W}_0) \xrightarrow{p} 0$. If $\mathcal{W}_0 = \{w^0\}$, then $\hat{w} \xrightarrow{p} w^0$.

If Assumption PP holds as well, then as $K_{\min} \rightarrow \infty$,

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \tilde{P}_0 \right) \xrightarrow{p} 0.$$

Therefore, under our standard regularity conditions on the kernel k , we are able to consistently estimate both the synthetic control weights and their induced pseudo-counterfactual. In the distinct cases of bounded misspecification and correct specification, we are able to characterize the probability limit of the estimated counterfactual to the true counterfactual P_0 :

Corollary 1 (Consistency relative to true counterfactual). *Suppose the conditions of Theorem 1 hold, including Assumption PP. Then, as $K_{\min} \rightarrow \infty$,*

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, P_0 \right) \leq \text{MMD}_k(\tilde{P}_0, P_0) + o_p(1)$$

If Assumption CM additionally holds, then $\text{MMD}_k(\tilde{P}_0, P_0) = 0$ and

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, P_0 \right) \xrightarrow{p} 0$$

We have therefore established conditions under which our estimated counterfactual is approximately or exactly equal to the true counterfactual distribution P_0 .

3.3 Weight Inference

Having established consistency of the estimated synthetic control weights $\hat{w}$ in the preceding section, we are now able to leverage our nonstandard asymptotic regime, and characterize their limiting distribution as well. While the consistency results above allow the population-minimizer set $\mathcal{W}_0$

to contain more than one weight vector, provided that all population minimizers induce the same pseudo-counterfactual distribution, a regular first-order limit for the full weight vector requires stronger local identification. We will now proceed to give such limits when the population minimizer is unique and it lies in the relative interior of the donor-weight simplex, once again adapting the approach to establishing asymptotic normality of extremum estimators established in [Newey and McFadden \(1994\)](#).

Let us begin by building some intuition. First, let us define a change of variables to remove the extraneous dimension of the weight vector imposed by the summation restriction. Let Θ denote the feasible set of such $(N - 2)$ -dimensional vectors:

$$\Theta = \left\{ \theta \in \mathbb{R}^{N-2} : \theta_j \geq 0 \text{ for every } j, \sum_{j=1}^{N-2} \theta_j \leq 1 \right\}$$

For a candidate element $\theta \in \Theta$, define $w(\theta)$ as the corresponding projection back to $\mathbb{R}^{N-1}$ and let $\ell(\theta)$ define the objective function in terms of θ :

$$w(\theta) = \left(\theta_1, \dots, \theta_{N-2}, 1 - \sum_{j=1}^{N-2} \theta_j \right)^\top \quad \ell(\theta) = L(w(\theta))$$

We can then derive the Jacobian of this weight map, A , which features in our expression for Σ_w in [Equation 2](#).

$$A = \frac{\partial w(\theta)}{\partial \theta^\top} = \begin{pmatrix} \mathbb{I}_{N-2} \\ -\mathbf{1}_{N-2}^\top \end{pmatrix}$$

Here, $\mathbb{I}_{N-2}$ refers to the $(N - 2)$ -dimensional identity matrix and $\mathbf{1}_{N-2}$ to an $(N - 2)$ -dimensional vector of 1s. The presence of A in the expression of Σ_w therefore yields the claimed singularity of the latter.

Now, let us examine the geometric properties of our induced loss. Define the following quantities, which describe the differences in kernel mean embeddings between the treated unit and the dropped reference donor unit ($j = N - 1$) and between a given donor unit and the reference, respectively:

$$b_t = \mu_{Nt} - \mu_{N-1,t} \quad d_{jt} = \mu_{jt} - \mu_{N-1,t}$$

We can then define the linear operator $D_t : \mathbb{R}^{N-2} \rightarrow \mathcal{H}_k$ which maps a change $h \in \mathbb{R}^{N-2}$ in donor weights θ to the corresponding change in the synthetic distribution's kernel mean embedding, at a given time t :

$$D_t h = \sum_{j=1}^{N-2} h_j d_{jt}$$

By definition of our objective function, we can then rewrite our objective function and derive the corresponding first derivative (score) and second derivative (Hessian):

$$\ell(\theta) = \sum_{t < T} \|b_t - D_t \theta\|_{\mathcal{H}_k}^2 \quad \nabla \ell(\theta) = -2 \sum_{t < T} D_t^* (b_t - D_t \theta) \quad H := \nabla^2 \ell(\theta) = 2 \sum_{t < T} D_t^* D_t$$

Here, $D_t^* : \mathcal{H}_k \rightarrow \mathbb{R}^{N-2}$ denotes the adjoint of D_t , which maps the gradient in the RKHS back to weight-coordinate derivatives. In particular, for every $u \in \mathcal{H}_k$ and $h \in \mathbb{R}^{N-2}$, D_t and D_t^* satisfy the following:

$$\langle u, D_t h \rangle_{\mathcal{H}_k} = \langle D_t^* u, h \rangle_{\mathbb{R}^{N-2}}$$

Here, $\langle x, y \rangle_B$ denotes the inner product of x and y in the space B . Note that these results also highlight the importance of using MMD_k^2 in our objective function; had we used unsquared MMD_k directly, the resulting score would be nonsmooth.

We have therefore demonstrated that our objective function in free parameter donor weights is quadratic, with a well-defined Hessian that admits an explicit kernel-mean interpretation. This geometric parameterization isolates all sampling variation in the empirical kernel means entering $\hat{b}_t$ and $\hat{D}_t$ (the empirical analogs of b_t and D_t , respectively). Boundedness of the kernel via [Assumption k-Reg](#) makes these empirical means root- K consistent by a direct second-moment bound, which yields objective and Hessian convergence.

The key ingredient of this asymptotic normality result lies in showing that the score function admits a central limit theorem. The formal statement and proof of this result is given in [Lemma 11](#) in [Appendix B](#); the following is an abbreviated overview of the intuition behind this result. First, define the residual $r_t = b_t - D_t \theta$ and define the following differences between empirical and population quantities:

$$\Delta b_t = \hat{b}_t - b_t \quad \Delta D_t = \hat{D}_t - D_t \quad \Delta r_t = \Delta b_t - \Delta D_t \theta^0$$

Furthermore, define the residual at the population optimum as $r_t^0 = b_t - D_t \theta^0$. We can then take the difference between the empirical and population score at the population optimum θ^0 :

$$\begin{aligned} \nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) &= -2 \sum_{t < T} (\hat{D}_t^* \hat{r}_t^0 - D_t^* r_t^0) \\ &= -2 \sum_{t < T} [D_t^* \Delta r_t + (\Delta D_t)^* r_t^0 + (\Delta D_t)^* \Delta r_t] \end{aligned}$$

The first two terms of this expansion contain first-order sampling variation, while the final is the product of two such terms, and is therefore second order. In particular, boundedness of the kernel via [Assumption k-Reg](#) implies that each empirical kernel mean satisfies $\hat{\mu}_{it} - \mu_{it} = O_p(K_{it}^{-1/2})$ in

$\mathcal{H}_k$. Under [Assumption BSG](#), this yields

$$\Delta r_t = O_p((K_{\min}^{\text{pre}})^{-1/2}) \quad \Delta D_t = O_p((K_{\min}^{\text{pre}})^{-1/2}) \quad (\Delta D_t)^* \Delta r_t = O_p((K_{\min}^{\text{pre}})^{-1})$$

Consequently, $\sqrt{K_{\min}^{\text{pre}}}(\Delta D_t)^* \Delta r_t = o_p(1)$, and so we can write

$$\sqrt{K_{\min}^{\text{pre}}} \left\{ \nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right\} = -2 \sum_{t < T} \left[D_t^* \sqrt{K_{\min}^{\text{pre}}} \Delta r_t + \left(\sqrt{K_{\min}^{\text{pre}}} \Delta D_t \right)^* r_t^0 \right] + o_p(1)$$

Under our sampling condition in [Assumption CS](#), the collection $\left\{ \sqrt{K_{\min}^{\text{pre}}}(\hat{\mu}_{it} - \mu_{it}) \right\}_{i,t}$ has a joint Hilbert-space Gaussian limit. A Gaussian element of $\mathcal{H}_k$ is the Hilbert-space analogue of a multivariate normal random vector, where instead of being normally distributed itself it rather has normally distributed inner product in every direction. More formally, a random element $G \in \mathcal{H}_k$ is Gaussian if $\langle G, h \rangle_{\mathcal{H}_k}$ is normally distributed for every $h \in \mathcal{H}_k$.

Δr_t and ΔD_t are linear transformations/operators of these embedding errors, and so admit a Gaussian limit via CLT with some limiting variance:

$$\sqrt{K_{\min}^{\text{pre}}} \left(\nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right) \xrightarrow{d} \mathcal{N}(0, \Omega)$$

Furthermore, since the assumed interiority of w^0 implies interiority of θ^0 , we know that $\nabla \ell(\theta^0) = 0$, and so we obtain the following CLT:

$$\sqrt{K_{\min}^{\text{pre}}} \nabla \hat{\ell}(\theta^0) \xrightarrow{d} \mathcal{N}(0, \Omega)$$

Letting ψ_{it} denote the score influence function for unit i at time period t , we can then write a closed-form expression for Ω :

$$\Omega = \sum_{t < T} \sum_{i=1}^N (\kappa_{it}^{\text{pre}})^{-1} \mathbb{E} \left[\psi_{it}(X) \psi_{it}(X)^\top \right]$$

For an exact specification of ψ_{it} , see the proof of [Lemma 11](#) in [Appendix B](#). Armed with this intuition, we can now state our asymptotic normality result formally.

Proposition 3 (Asymptotic normality of unique interior weights). *Suppose Assumptions [CS](#), [k-Reg](#), [IDT](#), and [BSG](#) hold. Suppose further that $w_i^0 > 0$ for every $i = 1, \dots, N-1$. Then,*

$$\sqrt{K_{\min}^{\text{pre}}}(\hat{w} - w^0) \xrightarrow{d} \mathcal{N}(0, \Sigma_w)$$

with

$$\Sigma_w = A H^{-1} \Omega H^{-1} A^\top \tag{2}$$

The covariance matrix Σ_w is singular in $\mathbb{R}^{N-1}$.

Moreover, each of the components of this limiting variance is consistently estimable, allowing us to estimate standard errors in practice.³

Proposition 4 (Feasible covariance estimation). *Suppose the conditions of Proposition 3 hold. Define $\psi_{itr} := \psi_{it}(x_{itr})$, let $\hat{\psi}_{itr}$ denote the feasible plug-in counterpart of ψ_{itr} , and define sample analogs*

$$\hat{H} = \nabla^2 \hat{\ell}(\hat{\theta}) \quad \hat{\Omega} = \sum_{t < T} \sum_{i=1}^N \frac{K_{\min}^{\text{pre}}}{K_{it}} \left(\frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \hat{\psi}_{itr} \hat{\psi}_{itr}^\top \right) \quad \hat{\Sigma}_w = A \hat{H}^{-1} \hat{\Omega} \hat{H}^{-1} A^\top$$

Then, as $K_{\min}^{\text{pre}} \rightarrow \infty$, $\hat{\Sigma}_w \xrightarrow{p} \Sigma_w$.

As a result, we have an analytic expression for the covariance of the synthetic control weights that is itself consistently estimable. This allows us to perform inference on the weights and obtain standard errors, a result that stands in stark contrast to the well-documented difficulties in characterizing standard errors for standard synthetic control (Almeida et al., 2025) and highlights one of the key advantages of our unique asymptotic regime.

3.4 Tractable Two-Stage Estimation & Inference

In the preceding section, we establish asymptotic normality of the synthetic control weight vector $\hat{w}$ under two key conditions: uniqueness of the population minimizer w^0 , and interiority of the same. As we have argued in Proposition 2, uniqueness of w^0 can be established via an affine independence restriction on the donor-kernel mean trajectories. While substantive, this condition is not on its face terribly burdensome; one sufficient condition would be for the donor distributions to be linearly independent in each pre-treatment period.

The interiority restriction, however, is quite strong. Synthetic control weights are typically sparse, and so will therefore not in general lie in the interior of the simplex (Abadie, 2021). However, despite its implausibility, this interiority assumption is critical for the tractable inference results obtained via Proposition 3. The key intuition for this lies in the fact that if the population weights for a given unit are strictly positive, the nonnegativity constraints are locally inactive, and a consistent estimator will satisfy these constraints with probability approaching one. Asymptotically, this estimator can therefore be treated as an unconstrained estimator, with only the summation to one constraint binding. However, when the true weight for a given unit is exactly zero, the nonnegativity constraint is locally active, and negative perturbations are infeasible. We therefore cannot apply the same unconstrained asymptotic approximation, leading to a (potentially) nonstandard limiting distribution. The true limiting distribution of $\sqrt{K_{\min}^{\text{pre}}}(\hat{w} - w^0)$ in this case would generally

³On the event that $\hat{H}$ is singular, define $\hat{\Sigma}_w = 0$.

be a projection of a Gaussian vector onto the critical cone which, while analytically derivable, poses significant challenges to implementation in practice. An inference result of utility to practitioners therefore relies on the (generally implausible) interiority of w^0 .

In this section, we resolve this tension by introducing a two-stage weight estimation procedure that first estimates the active donor set and then estimates weights on the face of the donor simplex where all inactive donors are restricted to have zero weight. As a consequence, the resulting weight estimator lies in the interior of the restricted simplex with probability approaching one, and all preceding consistency and inference results apply naturally. This is analogous to post-model-selection refitting in settings with variable selection (Belloni and Chernozhukov, 2013). However, our approach here yields pointwise asymptotic normality for a fixed population weight vector, rather than uniform validity over sequences of populations.

Under **Assumption IDT**, **Proposition 2** gives us a unique population optimum w^0 and we can define the population donor support as

$$\mathcal{S}_0 = \{i \in \{1, \dots, N-1\} : w_i^0 > 0\}$$

Note that the population zero-weight set is simply the complement of $\mathcal{S}_0$ in the set of donor units.

Our first step in this two-stage procedure is to produce an estimate of this set, which we shall denote $\hat{\mathcal{S}}$. Let $\tilde{w}$ denote the unrestricted estimator, which minimizes the empirical counterpart of **Equation 1** across the full donor simplex:

$$\tilde{w} \in \arg \min_{w \in \mathcal{W}} \hat{L}(w)$$

As $K_{\min}^{\text{pre}}$ becomes arbitrarily large, we can estimate $\hat{\mathcal{S}}$ by examining which elements of $\tilde{w}$ are larger than a cutoff schedule that shrinks to zero asymptotically more slowly than the convergence rate of the weights themselves. In practice, for some threshold λ_K that satisfies these asymptotic conditions, we can estimate:

$$\hat{\mathcal{S}} = \{i \in \{1, \dots, N-1\} : \tilde{w}_i > \lambda_K\}$$

If the estimated set is empty, set $\hat{\mathcal{S}} = \{j\}$, where j is the first index of an element of $\tilde{w}$ that is maximal. As our sample size increases and we can estimate this vector more precisely, the required threshold will tend towards zero. This procedure allows us to consistently estimate the active donor set.

Proposition 5 (Consistent active-donor set recovery). *Suppose Assumptions **CS**, **k-Reg**, and **IDT** hold. Suppose further that there exists a deterministic sequence $\lambda_K > 0$ such that $\lambda_K \rightarrow 0$ and $\sqrt{K_{\min}^{\text{pre}}} \lambda_K \rightarrow \infty$. Then,*

$$\Pr(\hat{\mathcal{S}} = \mathcal{S}_0) \rightarrow 1$$

Here, the condition that $\sqrt{K_{\min}^{\text{pre}}}\lambda_K \rightarrow \infty$ is analogous to similar conditions in the variable selection literature (Pötscher and Schneider, 2011) and implies that λ_K must decline to zero at a rate slower than $(K_{\min}^{\text{pre}})^{-1/2}$. In general, this can be satisfied by $\lambda_K = \alpha (K_{\min}^{\text{pre}})^{-\gamma}$ for $\gamma \in (0, 1/2)$ and $\alpha > 0$ a positive-valued constant.

We have therefore established that we can consistently estimate the set of active donor units by implementing a cutoff on the elements of the unrestricted weight estimate. It now remains to establish that re-estimating weights on this restricted set of donor units retains the previously established consistency and asymptotic normality properties of the unrestricted estimator.

For any nonempty $\mathcal{S} \subseteq \{1, \dots, N-1\}$, define the corresponding face of the donor-weight simplex by

$$\mathcal{W}_{\mathcal{S}} = \{w \in \mathcal{W} : w_i = 0 \text{ for every } i \notin \mathcal{S}\}$$

We now define the feasible second-stage weight estimator as

$$\hat{w}^{\text{ts}} \in \arg \min_{w \in \mathcal{W}_{\mathcal{S}}} \hat{L}(w)$$

where ties are resolved by the same fixed, measurable selection rule as for $\tilde{w}$. Analogous to our approach to Proposition 3, we can perform a change of variables to remove the linearly dependent component of $\hat{w}^{\text{ts}}$ by dropping the final unit as a reference. Let $s_0 = |\mathcal{S}_0|$. When $s_0 \geq 2$, we can then define $\Theta_{\mathcal{S}_0} \subset \mathbb{R}_+^{s_0-1}$, $w_{\mathcal{S}_0}(\theta) \in \mathcal{W}_{\mathcal{S}_0}$, and $\ell_{\mathcal{S}_0}(\theta)$ as direct extrapolations to the case where we restrict attention to the active-donor set $\mathcal{S}_0$. Similarly, let $A_{\mathcal{S}_0}$ denote the Jacobian of the weight map on this restricted face of the donor simplex.

Theorem 2 (Consistency and asymptotic normality of support-refitted weights). *Suppose the conditions of Proposition 5 hold. Then $\hat{w}^{\text{ts}} \xrightarrow{P} w^0$.*

If $s_0 \geq 2$ and Assumption BSG also holds, then

$$\sqrt{K_{\min}^{\text{pre}}}(\hat{w}^{\text{ts}} - w^0) \xrightarrow{d} \mathcal{N}(0, \Sigma_{w, \mathcal{S}_0})$$

where

$$\Sigma_{w, \mathcal{S}_0} = A_{\mathcal{S}_0} H_{\mathcal{S}_0}^{-1} \Omega_{\mathcal{S}_0} H_{\mathcal{S}_0}^{-1} A_{\mathcal{S}_0}^{\top}$$

Note that the $s_0 = 1$ case is omitted - if there is only a single active donor, then it must receive unit weight and the limiting covariance is the zero matrix.

We have thus established that this two-stage weight estimation procedure yields consistent and asymptotically normal weights, without restriction on interiority of w^0 on the full donor simplex. Consequently, this will be our preferred weight estimation method for the remainder of this paper. We conclude by extending the argument from Proposition 4 to this two-stage setting and establishing feasible estimation, where the only additional complication comes from the fact that $\hat{\mathcal{S}}$ is a

random object.⁴

Corollary 2 (Feasible covariance estimation for support-refitted weights). *Suppose the conditions of [Proposition 5](#) hold, $s_0 \geq 2$, and [Assumption BSG](#) holds. Using the parameterization of the selected face $\mathcal{W}_{\hat{S}}$, let $\hat{\psi}_{itr, \hat{S}}$ denote the feasible plug-in counterpart of the corresponding face-specific influence function, and define*

$$\hat{H}_{\hat{S}} = \nabla^2 \hat{\ell}_{\hat{S}}(\hat{\theta}^{\text{ts}}) \quad \hat{\Omega}_{\hat{S}} = \sum_{t < T} \sum_{i \in \hat{S} \cup \{N\}} \frac{K_{\min}^{\text{pre}}}{K_{it}} \left(\frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \hat{\psi}_{itr, \hat{S}} \hat{\psi}_{itr, \hat{S}}^{\top} \right) \quad \hat{\Sigma}_{w, \hat{S}} = A_{\hat{S}} \hat{H}_{\hat{S}}^{-1} \hat{\Omega}_{\hat{S}} \hat{H}_{\hat{S}}^{-1} A_{\hat{S}}^{\top}$$

Then, as $K_{\min}^{\text{pre}} \rightarrow \infty$, $\hat{\Sigma}_{w, \hat{S}} \xrightarrow{p} \Sigma_{w, S_0}$.

4 Hypothesis Generation & Validation

Having established that we can consistently estimate counterfactual distributions and can quantify the uncertainty in their construction via asymptotic normality of the synthetic control weights, we now turn our attention to the second stage of our pipeline: hypothesis generation and validation.

Our objective in this section is to use the observed $\hat{P}_1$ and estimated $\hat{P}_0$ to learn some interpretable scalar-valued summary function of the text, represented as a function $\hat{g} : \mathcal{X} \rightarrow \mathbb{R}$, that varies between those two distributions. This function evaluates the presence of an aspect or theme of the text that is hypothesized to have been impacted by treatment. In practice, learned summary functions may be evaluated on the original text, rather than its embedding. For notational convenience, we will continue to write these summaries as functions of the embeddings, $\hat{g}(x)$. Implicit in this notation is the assumption that the embedding retains all information relevant to the particular summary function at hand, as motivated by the broad literature on the ability of embeddings to capture semantic meaning of text discussed in [Section 2](#).⁵

We can then quantify the effect of treatment on this learned summary via calculating the difference in mean summary values across the two distributions of interest:

$$\hat{\tau}_{\hat{g}} = \mathbb{E}_{\hat{P}_1}[\hat{g}(X)] - \mathbb{E}_{\hat{P}_0}[\hat{g}(X)] = \mathbb{E}_{\hat{F}_{NT}}[\hat{g}(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{\hat{F}_{iT}}[\hat{g}(X)]$$

In the setting in which our synthetic control model is correctly specified (that is, [Assumption CM](#) holds), this is an estimate of the true average treatment effect on the treated unit (ATT), as in the

⁴On the event that $\hat{H}$ is singular or $|\hat{S}| = 1$, define $\hat{\Sigma}_{w, \hat{S}} = 0$.

⁵Specifically, if $\hat{s}$ denotes the learned text-level scoring function, we assume that $\hat{s}(a) = \hat{g}(\mathcal{E}(a))$ almost surely under the relevant potential text distributions.

broader synthetic control literature (Abadie, Diamond, and Hainmueller, 2010):

$$\tau_{\hat{g}} = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{P_0}[\hat{g}(X)]$$

More generally, however, $\hat{\tau}_{\hat{g}}$ estimates the effect relative to the pseudo-counterfactual:

$$\tilde{\tau}_{\hat{g}} = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{\tilde{P}_0}[\hat{g}(X)]$$

In this section, we will establish consistency and inference for $\hat{\tau}_{\hat{g}}$, in both the general case of bounded misspecification and the special case of correct specification.

4.1 Fixed Summary

We begin by considering the simpler case where we are simply interested in evaluating the effect of treatment on some pre-specified, measurable summary function $g : \mathcal{X} \rightarrow \mathbb{R}$. Conceptually, this corresponds to g measuring some aspect of the text that has been hypothesized by the researcher without exposure to any data. As a consequence, this function is considered to be a fixed, nonrandom object unconditionally. The estimated weights $\hat{w}^{\text{ts}}$, on the other hand, remain random and are estimated from pre-treatment data. Although we make no assumptions on how the researcher specified this summary function (other than having done so without examining the data), we do assume that the produced summary itself is bounded:

Assumption BFS (Bounded fixed summary). *There exists a finite constant $\bar{g} < \infty$ such that*

$$\sup_{x \in \mathcal{X}} |g(x)| \leq \bar{g}$$

We redefine our estimator for the case of a pre-specified g as:

$$\hat{\tau}_g = \mathbb{E}_{\hat{P}_1}[g(X)] - \mathbb{E}_{\hat{P}_0}[g(X)] = \mathbb{E}_{\hat{F}_{NT}}[g(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{\hat{F}_{iT}}[g(X)]$$

Our population parameters of interest correspondingly become:

$$\tau_g = \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{P_0}[g(X)] \quad \tilde{\tau}_g = \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{\tilde{P}_0}[g(X)]$$

We can now establish consistency of the difference-in-means estimator for the ATT relative to the pseudo-counterfactual.

Theorem 3 (Fixed summary consistency). *Suppose the conditions of [Proposition 5](#) hold, and that [Assumption BFS](#) holds as well. Then, as $K_{\min} \rightarrow \infty$, $\hat{\tau}_g \xrightarrow{P} \tilde{\tau}_g$.*

This result relies on an algebraic decomposition of the estimation error $\hat{\tau}_g - \tilde{\tau}_g$, with certain

terms going to zero by a law of large numbers and others due to consistency of weight estimation via invocation of [Theorem 2](#). In particular, no smoothness conditions on g are required. We can now characterize the probability limit of our estimator in terms of the true ATT τ_g in the distinct cases of bounded misspecification and correct specification. For the bounded misspecification case, we need to introduce an extra bridge to map the distributional distance ϵ_0 to the difference in summary difference-in-means values induced by $\tilde{P}_0$ versus P_0 .

Assumption ASF (Fixed-summary stability under approximation). *There exists a finite constant $C_g < \infty$ such that, for all distributions P, Q over $\mathcal{X}$ for which the relevant expectations are well-defined:*

$$|\mathbb{E}_P[g(X)] - \mathbb{E}_Q[g(X)]| \leq C_g \text{MMD}_k(P, Q)$$

Corollary 3 (Fixed summary consistency relative to τ_g). *Suppose the conditions of [Theorem 3](#) hold. Then, the following hold:*

1. *If [Assumption ASF](#) holds as well, then $|\hat{\tau}_g - \tau_g| \leq C_g \text{MMD}_k(\tilde{P}_0, P_0) + o_p(1)$.*
2. *If $\tilde{P}_0 = P_0$, then $\text{MMD}_k(\tilde{P}_0, P_0) = 0$ and $\hat{\tau}_g \xrightarrow{p} \tau_g$.*

Observe that [Assumption ASF](#) implies that proximity of the underlying distributions $\tilde{P}_0$ and P_0 (as measured by MMD_k) implies proximity in mean summary scores applied to each distribution. We can extend this argument further: our chosen inter-distributional distance metric MMD_k in fact automatically controls a wide class of summary functions. In particular, controlling for MMD_k distance implicitly controls mean inter-distributional distance for all functions expressible in the RKHS $\mathcal{H}_k$, with bound corresponding to the RKHS norm of the chosen function.

Proposition 6 (Automatic summary control). *Suppose [Assumption k-Reg](#) holds. Let $g \in \mathcal{H}_k$. Then,*

$$|\tilde{\tau}_g - \tau_g| \leq \|g\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0)$$

This is a key feature of our distributional matching approach: even if our set of donor units is insufficiently rich to allow for correct specification of the synthetic control model, by controlling for the distance between the true counterfactual P_0 and its approximation $\tilde{P}_0$ we can simultaneously control for a broad class of scalar summaries. The exact class of functions that are expressible in the RKHS, and therefore simultaneously controlled, depends on the choice of kernel k . For a linear kernel, this would be linear functions; for a polynomial kernel, this would be polynomials up to the degree of the kernel. However, these are in general not characteristic, and so would not satisfy [Assumption k-Reg](#). Among the previously discussed suitable kernels which satisfy those necessary conditions, the Gaussian RBF and rational quadratic kernels have respective RKHSs containing analytic functions ([Sun and Zhou, 2008](#)), while the Matérn kernel has RKHS containing Sobolev functions with finite smoothness ([Porcu et al., 2024](#)).

However, even in cases where the summary g does not itself lie within this RKHS, we can bound the induced bias in the difference in means estimator for the ATT by the distance (in sup-norm) between $g \notin \mathcal{H}_k$ and a reference function $g_k \in \mathcal{H}_k$, if one exists.

Corollary 4 (Approximate summary control). *Suppose **Assumption k-Reg** holds. Let $g, g_k : \mathcal{X} \rightarrow \mathbb{R}$, with $g_k \in \mathcal{H}_k$. Then,*

$$\sup_{x \in \mathcal{X}} |g(x) - g_k(x)| \leq \delta \implies |\tilde{\tau}_g - \tau_g| \leq \|g_k\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) + 2\delta$$

Having established consistency of our difference-in-means estimator for the ATT relative to the pseudo-counterfactual, and bounded the bias relative to the true ATT in the case of bounded misspecification of the synthetic control model, we now turn our attention to performing inference on $\hat{\tau}_g$ in a manner that allows us to take into account uncertainty from the first-stage weight estimation. We begin by making an assumption governing the rate of growth of the pre-treatment period relative to the full sample.

Assumption RPG (Relative pre-treatment growth). *As $K_{\min} \rightarrow \infty$,*

$$\rho_K := \frac{K_{\min}}{K_{\min}^{\text{pre}}} \longrightarrow \rho \in [0, 1]$$

This assumption allows us to track how fast the pre-treatment subset is growing in relation to the overall rate of growth of the whole data. Note that since $K_{\min} \leq K_{\min}^{\text{pre}}$, this ratio is bounded between 0 and 1 by construction. The case $\rho = 0$ corresponds to pre-treatment sample size used to estimate the weights growing faster than the post-treatment samples used to estimate $\hat{\tau}_g$. With this assumption in hand, we can now state our inference result.

Theorem 4 (Asymptotic normality of fixed summary treatment effect). *Suppose the conditions of **Theorem 2** hold, and that Assumptions **BFS**, **BSG**, and **RPG** hold as well. Then, as $K_{\min} \rightarrow \infty$,*

$$\sqrt{K_{\min}} (\hat{\tau}_g - \tilde{\tau}_g) \xrightarrow{d} N(0, V_g)$$

where

$$V_g = \frac{\sigma_N^2(g)}{\kappa_{NT}^{\text{post}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(g)}{\kappa_{iT}^{\text{post}}} + \rho \mu(g)^\top \Sigma_{w, S_0} \mu(g)$$

$$\sigma_i^2(g) = \text{Var}_{F_{iT}(W_{iT})}(g(X)) \quad \mu(g) = \left(\mathbb{E}_{F_{1T}(0)}[g(X)], \dots, \mathbb{E}_{F_{(N-1)T}(0)}[g(X)] \right)^\top$$

We have therefore established asymptotic normality of the difference in means estimator. Examining the expression for the limiting variance V_g we observe that it contains three terms: the first two correspond to the variance in the summary g in the treated and donor unit distributions,

respectively, while the third term captures the uncertainty in the first-stage weight estimation problem, with Σ_{w,S_0} corresponding to the limiting variance of $\hat{w}^{\text{ts}}$ established in [Theorem 2](#). Finally, we can provide an extension to [Corollary 2](#) and characterize a feasible, consistent estimator of this limiting variance.

Corollary 5 (Feasible fixed summary variance estimation). *Suppose the conditions of [Theorem 4](#) hold, and that the number of active donors $s_0 \geq 2$. Let V_g be defined as in [Theorem 4](#), $\hat{\Sigma}_{w,\hat{S}}$ as defined in [Corollary 2](#), and define the following:*

$$\begin{aligned}\bar{g}_i &= \frac{1}{K_{iT}} \sum_{k=1}^{K_{iT}} g(x_{iT k}) & \hat{\sigma}_i^2(g) &= \frac{1}{K_{iT} - 1} \sum_{k=1}^{K_{iT}} (g(x_{iT k}) - \bar{g}_i)^2 & \hat{\mu}(g) &= (\bar{g}_1, \dots, \bar{g}_{N-1})^\top \\ \hat{V}_g &= K_{\min} \left(\frac{\hat{\sigma}_N^2(g)}{K_{NT}} + \sum_{i=1}^{N-1} (\hat{w}_i^{\text{ts}})^2 \frac{\hat{\sigma}_i^2(g)}{K_{iT}} \right) + \rho_K \hat{\mu}(g)^\top \hat{\Sigma}_{w,\hat{S}} \hat{\mu}(g)\end{aligned}$$

Then, as $K_{\min} \rightarrow \infty$, $\hat{V}_g \xrightarrow{p} V_g$.

We have therefore shown that distributional matching via MMD_k simultaneously controls for a wide class of summary functions automatically, and that for any bounded summary function chosen ex-ante we can consistently estimate the ATT on this summary (relative to the pseudo-counterfactual) and perform valid inference on the resulting causal estimate (provided $V_g > 0$), taking into account first-stage uncertainty in the constructed counterfactual.

4.2 Data-Driven Summary

Having built intuition and established key results in the context of an ex-ante fixed summary function g , we now turn our attention to the case where the researcher seeks to *learn* what features of the text were impacted by treatment, and to then quantify the treatment effects along those learned dimensions. This process of hypothesizing aspects of text sensitive to treatment requires providing data from $\hat{P}_1$ and $\hat{P}_0$ to some algorithm, so that it might learn what differences arise by virtue of treatment being imposed. In practice, LLMs would be a compelling choice of algorithm, given their ability to process and reason about text at scale. However, it is unclear how we can reason about the properties of a function $\hat{g}$ returned by an LLM, particularly with respect to how we can conceptualize its convergence properties to some (potentially known) “true” g . Consequently, we will proceed by providing results for consistency and inference conditional on the hypothesized summary $\hat{g}$; we will continue to take into account first-stage uncertainty in counterfactual construction and second-stage sampling uncertainty, but will not make any statements about possible errors made in the formulation of the learned $\hat{g}$ itself.

This modeling decision, while necessary to incorporate modern “black-box” AI tools such as LLMs, immediately gives rise to one challenge that was absent from the fixed-summary setting:

evaluating the learned summary $\hat{g}$ on the same data from which it was discovered cannot admit tractable inference for $\hat{\tau}_{\hat{g}}$. However, this is easily resolved by splitting the post-treatment dataset into two subsets, separately for each cell of the panel matrix. We will call these subsets “discovery” (abbreviated “disc”) and “evaluation” (abbreviated “eval”).

More formally, we will partition the observed post-treatment document indices for each unit $i = 1, \dots, N$ into two disjoint sets:

$$\mathcal{I}_{iT}^{\text{disc}} \dot{\cup} \mathcal{I}_{iT}^{\text{eval}} = \{1, \dots, K_{iT}\}$$

For $a \in \{\text{disc}, \text{eval}\}$, let $K_{iT}^a = |\mathcal{I}_{iT}^a|$ and define the subset-specific empirical distributions as

$$\hat{F}_{iT}^a = \frac{1}{K_{iT}^a} \sum_{k \in \mathcal{I}_{iT}^a} \delta_{x_{iT k}}$$

We now introduce an assumption that governs this split process, ensuring that the partitions themselves are constructed honestly and in a consistent manner.

Assumption PoTS (Post-treatment sample splitting). *The partitions $\{\mathcal{I}_{iT}^{\text{disc}}, \mathcal{I}_{iT}^{\text{eval}}\}_{i=1}^N$ are constructed independently of the observed comments. There exist constants*

$$\pi_T^{\text{disc}}, \pi_T^{\text{eval}} \in (0, 1) \quad \pi_T^{\text{disc}} + \pi_T^{\text{eval}} = 1$$

such that, for every $i = 1, \dots, N$ and $a \in \{\text{disc}, \text{eval}\}$,

$$\frac{K_{iT}^a}{K_{iT}} \rightarrow \pi_T^a$$

Consequently, conditional on the split indices, the discovery and evaluation samples are mutually independent across folds and unit-period cells. We now define the minimum number of comments across units in the post-treatment period as:

$$K_{\min}^a := \min_{1 \leq i \leq N} K_{iT}^a \quad a \in \{\text{disc}, \text{eval}\}$$

Under this sample-splitting regime, we can redefine our target distributions of interest for each fold $a \in \{\text{disc}, \text{eval}\}$:

$$\hat{P}_1^a = \hat{F}_{NT}^a \quad \hat{P}_0^a = \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \hat{F}_{iT}^a$$

A summary function measuring some hypothesized aspect of the text can then be generated by an algorithm which is given access to $\hat{P}_1^{\text{disc}}$ and $\hat{P}_0^{\text{disc}}$. The resulting function will itself then be a function of the post-treatment discovery-fold data, denoted $D_{\text{post}}^{\text{disc}}$, from which these empirical distributions are constructed. However, construction of $\hat{P}_0^{\text{disc}}$ relies on the weights $\hat{w}^{\text{ts}}$, which are

functions of the pre-treatment data, denoted D_{pre} . As a result, the algorithm is (indirectly) exposed to pre-treatment data as well. We formalize the hypothesis generation algorithm and its produced summary function in the following assumption, in which we also impose a boundedness condition on the learned function analogous to **Assumption BFS**.

Assumption SG (Summary generation). *There exists an algorithm $\mathcal{A}_K$ and a class $\mathcal{G}$ of measurable functions $g : \mathcal{X} \rightarrow \mathbb{R}$ such that*

$$\hat{g} = \mathcal{A}_K \left(D_{\text{pre}}, D_{\text{post}}^{\text{disc}}; \mathcal{R}_K \right) \in \mathcal{G}$$

where $\mathcal{R}_K$ captures auxiliary randomization inherent to the algorithm $\mathcal{A}_K$ and is independent of data and split assignments. The generated summary function $\hat{g}$ is therefore measurable with respect to the discovery sigma field

$$\mathcal{D}_K = \mathcal{S}_K \vee \sigma \left(D_{\text{pre}}, D_{\text{post}}^{\text{disc}}, \mathcal{R}_K \right)$$

where $\mathcal{S}_K$ captures post-treatment sample-split assignments. Additionally, there exists a finite constant $\bar{g} < \infty$ such that

$$\sup_{g \in \mathcal{G}} \sup_{x \in \mathcal{X}} |g(x)| \leq \bar{g}$$

This assumption makes explicit that the post-treatment evaluation fold is not used to construct, tune, revise, or select $\hat{g}$. The boundedness condition is imposed uniformly over the possible outputs of the discovery algorithm $\mathcal{A}_K$. The algorithm-specific auxiliary randomization $\mathcal{R}_K$ captures any and all sources of randomness that are introduced in the algorithmic process; in the case of an LLM-based algorithm, this includes the parameters which govern randomness in the generative process.

With this learned summary function in hand, we construct estimates of the ATT on the evaluation fold data:

$$\hat{\tau}_{\hat{g}}^{\text{eval}} = \mathbb{E}_{\hat{P}_1^{\text{eval}}}[\hat{g}(X)] - \mathbb{E}_{\hat{P}_0^{\text{eval}}}[\hat{g}(X)] = \mathbb{E}_{\hat{F}_{NT}^{\text{eval}}}[\hat{g}(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{\hat{F}_{iT}^{\text{eval}}}[\hat{g}(X)]$$

Note that this object is now defined conditional on $\hat{g}$. We can similarly define the new population quantities of interest, the true ATT and the ATT relative to the pseudo-counterfactual, conditional on this learned summary:

$$\tau_{\hat{g}} = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{P_0}[\hat{g}(X)] \quad \tilde{\tau}_{\hat{g}} = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{\tilde{P}_0}[\hat{g}(X)]$$

We can now establish consistency of our estimated ATT $\hat{\tau}_{\hat{g}}^{\text{eval}}$, analogous to **Theorem 3**. In this case, we now must first condition on the discovery sigma-field $\mathcal{D}_K$, in order to treat $\hat{g}$ as fixed. Henceforth, convergence in distribution of a statistic T_K conditional on $\mathcal{D}_K$ is understood to mean,

for every bounded and continuous f :

$$\mathbb{E}[f(T_K)|\mathcal{D}_K] \xrightarrow{p} \mathbb{E}[f(Z)]$$

Here, Z is distributed according to the distributional limit of T_K , conditional on $\mathcal{D}_K$.

Theorem 5 (Generated summary consistency). *Suppose the conditions of [Proposition 5](#) hold, and that Assumptions [PoTS](#) and [SG](#) hold as well. Then, conditional on $\mathcal{D}_K$, as $K_{\min} \rightarrow \infty$, $\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}} \xrightarrow{p} 0$.*

Proceeding in the same fashion, we can go on to characterize the probability limit of our estimator in terms of the true ATT $\tau_{\hat{g}}$ in the special cases of bounded misspecification and correct specification.

Corollary 6 (Generated summary consistency relative to $\tau_{\hat{g}}$). *Suppose the conditions of [Theorem 5](#) hold. Then, the following hold:*

1. *If $\hat{g} \in \mathcal{H}_k$, then $|\hat{\tau}_{\hat{g}}^{\text{eval}} - \tau_{\hat{g}}| \leq \|\hat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) + o_p(1)$.*
2. *If $\tilde{P}_0 = P_0$, then $\text{MMD}_k(\tilde{P}_0, P_0) = 0$ and $\hat{\tau}_{\hat{g}}^{\text{eval}} - \tau_{\hat{g}} \xrightarrow{p} 0$.*

That is, as long as the learned summary function $\hat{g}$ is expressible within the RKHS of our chosen kernel k , we can bound the bias induced in the treatment effect estimate by misspecification of the synthetic control model. By minimizing the distance between the pre-treatment trajectories of the treated unit and the mixture of donor units, we are therefore finding weights that simultaneously bound misspecification-induced pre-treatment discrepancies for a large, kernel-dependent class of summary functions.

We now turn our attention to inference for this estimated $\hat{\tau}_{\hat{g}}^{\text{eval}}$. This will be a more complicated exercise than in the fixed-summary case, in which the summary itself was completely independent of pre-treatment data. Suppose, for example, that the same pre-treatment weight estimator $\hat{w}^{\text{ts}}$ is used both to construct the discovery pseudo-counterfactual and to estimate the final treatment effect. Then, $\hat{g}$ is some function of $\hat{w}^{\text{ts}}$, and conditioning on the complete discovery information would also condition on $\hat{w}^{\text{ts}}$, eliminating the sampling variation used to derive the first-stage weight limit. Consequently, we must eliminate the (first-order) dependence of $\hat{g}$ on the pre-treatment data used to estimate the weights used to construct $\hat{\tau}_{\hat{g}}^{\text{eval}}$.

4.3 Inference for Data-Driven Summaries via Pre-Treatment Sample Splitting

In this approach to inference, we remove the dependence of $\hat{\tau}_{\hat{g}}^{\text{eval}}$ and $\hat{g}$ on a common set of weights by splitting the pre-treatment dataset and estimating two sets of weights on independent data, one of which is used for constructing the counterfactual used to learn $\hat{g}$ while the other is used for estimating $\hat{\tau}_{\hat{g}}^{\text{eval}}$.

Maintaining consistent notation for the different folds used when splitting the post-treatment sample, we partition the observed pre-treatment ($t < T$) document indices for each unit $i = 1, \dots, N$ into two disjoint sets:

$$\mathcal{I}_{it}^{\text{disc}} \cup \mathcal{I}_{it}^{\text{eval}} = \{1, \dots, K_{it}\}$$

For $a \in \{\text{disc}, \text{eval}\}$, let $K_{it}^a = |\mathcal{I}_{it}^a|$ and define the subset-specific empirical distributions as

$$\hat{F}_{it}^a = \frac{1}{K_{it}^a} \sum_{k \in \mathcal{I}_{it}^a} \delta_{x_{itk}}$$

We now introduce an assumption analogous to **Assumption PoTS**.

Assumption PrTS (Pre-treatment sample splitting). *The partitions $\{\mathcal{I}_{it}^{\text{disc}}, \mathcal{I}_{it}^{\text{eval}}\}_{i=1, \dots, N; t < T}$ are constructed independently of the observed comments and of the post-treatment split. There exist constants*

$$\pi_{\text{pre}}^{\text{disc}}, \pi_{\text{pre}}^{\text{eval}} \in (0, 1), \quad \pi_{\text{pre}}^{\text{disc}} + \pi_{\text{pre}}^{\text{eval}} = 1$$

such that, for every $i = 1, \dots, N$ and $a \in \{\text{disc}, \text{eval}\}$,

$$\frac{K_{it}^a}{K_{it}} \longrightarrow \pi_{\text{pre}}^a$$

We can now further define the minimum number of comments across units in the pre-treatment period as:

$$K_{\min}^{\text{pre}, a} := \min_{1 \leq i \leq N, t < T} K_{it}^a \quad a \in \{\text{disc}, \text{eval}\}$$

We will now proceed to estimate two sets of weights, one on each of these pre-treatment folds, using for each our two-stage estimation procedure to ensure interiority. For fold $a \in \{\text{disc}, \text{eval}\}$, define the first-stage weight and active donor set estimates:

$$\tilde{w}^a \in \arg \min_{w \in \mathcal{W}} \hat{L}^a(w) \quad \hat{\mathcal{S}}^a = \{i \in \{1, \dots, N-1\} : \tilde{w}_i^a > \lambda_K^a\}$$

Here, $\hat{L}^a(w)$ denotes the objective function evaluated on the given fold of pre-treatment data:

$$\hat{L}^a(w) = \sum_{t < T} \text{MMD}_k^2 \left(\hat{F}_{Nt}^a, \sum_{i=1}^{N-1} w_i \hat{F}_{it}^a \right)$$

We can then estimate the two-stage weights on each fold as:

$$\hat{w}^{\text{ts}, a} \in \arg \min_{w \in \mathcal{W}_{\hat{\mathcal{S}}^a}} \hat{L}^a(w)$$

For both $\tilde{w}^a$ and $\hat{w}^{\text{ts}, a}$, ties are resolved via the same fixed measurable rule (as outlined in **subsec-**

tion 3.4). Previous arguments for consistency and asymptotic normality of the two-stage weights on the full sample lead naturally to analogous results for these fold-specific estimators; see [Appendix C](#) for further details.

We can now use these weights to construct the relevant counterfactuals for the corresponding folds. In particular, we alter how we construct the summary $\hat{g}$, resulting in the following modified version of [Assumption SG](#):

Assumption PrSG (Summary generation with pre-treatment splitting). *There exists an algorithm $\mathcal{A}_K$ and a class $\mathcal{G}$ of measurable functions $g : \mathcal{X} \rightarrow \mathbb{R}$ such that*

$$\hat{g} = \mathcal{A}_K \left(D_{pre}^{disc}, D_{post}^{disc}, \mathcal{R}_K \right) \in \mathcal{G}$$

where $\mathcal{R}_K$ captures auxiliary randomization in inherent to the algorithm $\mathcal{A}_K$ and is independent of data and split assignments. The generated summary function $\hat{g}$ is therefore measurable with respect to the discovery sigma field

$$\mathcal{D}_K = \mathcal{S}_K \vee \sigma \left(D_{pre}^{disc}, D_{post}^{disc}, \mathcal{R}_K \right)$$

where $\mathcal{S}_K$ captures both pre- and post-treatment sample-split assignments. Additionally, there exists a finite constant $\bar{g} < \infty$ such that

$$\sup_{g \in \mathcal{G}} \sup_{x \in \mathcal{X}} |g(x)| \leq \bar{g}.$$

The key modification here is that $\hat{g}$ is now fixed conditioning on both pre- and post-treatment data only in the discovery fold; crucially, we do not have to condition on D_{pre}^{eval} . We then use weights learned on that evaluation fold of data to estimate the treatment effect:

$$\hat{\tau}_{\hat{g}}^{eval, split} = \mathbb{E}_{\hat{F}_{NT}^{eval}}[\hat{g}(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{ts, eval} \mathbb{E}_{\hat{F}_{iT}^{eval}}[\hat{g}(X)]$$

Our two population parameters of interest, $\tilde{\tau}_{\hat{g}}$ and $\tau_{\hat{g}}$, are insensitive to this modification and remain unchanged from before. First, we verify that our previous consistency results for $\hat{\tau}_{\hat{g}}^{eval}$ are unaffected by this pre-treatment sample splitting.

Proposition 7 (Consistency with pre-treatment sample splitting). *Suppose the conditions of [Proposition 5](#) hold for each pre-treatment fold, and that Assumptions [PoTS](#), [PrTS](#), and [PrSG](#) hold as well. Then, conditional on $\mathcal{D}_K$, as $K_{min} \rightarrow \infty$, the following hold:*

1. $\hat{\tau}_{\hat{g}}^{eval, split} - \tilde{\tau}_{\hat{g}} \xrightarrow{p} 0$
2. If $\hat{g} \in \mathcal{H}_k$, then $|\hat{\tau}_{\hat{g}}^{eval, split} - \tau_{\hat{g}}| \leq \|\hat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) + o_p(1)$
3. If $\tilde{P}_0 = P_0$, then $\text{MMD}_k(\tilde{P}_0, P_0) = 0$ and $\hat{\tau}_{\hat{g}}^{eval, split} - \tau_{\hat{g}} \xrightarrow{p} 0$

The intuition behind this result is simple: since the fold-specific weights are both consistent for the population weights, the same is true for the fold-specific counterfactuals, and so the same argument underlying [Theorem 5](#) and [Corollary 6](#) holds in this setting. With consistency established, we can show that our treatment effect estimator $\hat{\tau}_{\hat{g}}^{\text{eval,split}}$ is asymptotically normal, with a limiting variance that takes into account the first-stage weight-estimation error, analogous to the result in [Theorem 4](#).

Assumption PEFG (Evaluation-fold growth with pre-treatment sample splitting). *As $K_{\min}^{\text{eval}} \rightarrow \infty$, there exist constants $\kappa_{iT}^{\text{eval}} \in (0, \infty)$ for all $i = 1, \dots, N$ such that*

$$\frac{K_{iT}^{\text{eval}}}{K_{\min}^{\text{eval}}} \rightarrow \kappa_{iT}^{\text{eval}}$$

Moreover,

$$\rho_K^{\text{eval}} := \frac{K_{\min}^{\text{eval}}}{K_{\min}^{\text{pre,eval}}} \rightarrow \rho^{\text{eval}} \in [0, \infty)$$

Theorem 6 (Treatment effect inference via pre-treatment sample splitting). *Suppose the conditions of [Proposition 7](#) hold, and that Assumptions [BSG](#) and [PEFG](#) hold as well. Suppose further that $V_{\hat{g}}^{-1} = O_p(1)$. Then, conditional on $\mathcal{D}_K$, as $K_{\min}^{\text{eval}} \rightarrow \infty$,*

$$\frac{\sqrt{K_{\min}^{\text{eval}}} \left(\hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tilde{\tau}_{\hat{g}} \right)}{\sqrt{V_{\hat{g}}}} \xrightarrow{d} N(0, 1)$$

where

$$V_{\hat{g}} = \frac{\sigma_N^2(\hat{g})}{\kappa_{NT}^{\text{eval}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(\hat{g})}{\kappa_{iT}^{\text{eval}}} + \rho^{\text{eval}} \mu(\hat{g})^\top \Sigma_{w, \mathcal{S}_0} \mu(\hat{g})$$

$$\sigma_i^2(\hat{g}) = \text{Var}_{F_{iT}(W_{iT})}(\hat{g}(X)) \quad \mu(\hat{g}) = \left(\mathbb{E}_{F_{1T}(0)}[\hat{g}(X)], \dots, \mathbb{E}_{F_{(N-1)T}(0)}[\hat{g}(X)] \right)^\top$$

The form of this limiting variance is effectively identical to that of [Theorem 4](#), with the first two terms corresponding to variances of the learned summary function $\hat{g}$ in both treated and donor units while the third term captures uncertainty from the weights estimated in the first stage. The primary difference is that now, all of these quantities must be estimated on the evaluation fold of data, both pre- and post-treatment. The following result extends [Corollary 5](#) to this setting, establishing consistency of the feasible variance estimators.

Corollary 7 (Feasible variance estimation via pre-treatment sample splitting). *Suppose the conditions of [Proposition 7](#) hold, Assumptions [BSG](#) and [PEFG](#) hold, and the number of active donors $s_0 \geq 2$. Let $V_{\hat{g}}$ be defined as in [Theorem 6](#), $\hat{\Sigma}_{w, \hat{\mathcal{S}}^{\text{eval}}}$ as defined in [Corollary 2](#) using weights estimated*

on the pre-treatment evaluation fold, and define the following:

$$\begin{aligned}\bar{g}_i &= \frac{1}{K_{iT}^{\text{eval}}} \sum_{k \in \mathcal{I}_{iT}^{\text{eval}}} \hat{g}(x_{iT k}) & \hat{\sigma}_i^2(\hat{g}) &= \frac{1}{K_{iT}^{\text{eval}} - 1} \sum_{k \in \mathcal{I}_{iT}^{\text{eval}}} (\hat{g}(x_{iT k}) - \bar{g}_i)^2 & \hat{\mu}(\hat{g}) &= (\bar{g}_1, \dots, \bar{g}_{N-1})^\top \\ \hat{V}_{\hat{g}} &= K_{\min}^{\text{eval}} \left(\frac{\hat{\sigma}_N^2(\hat{g})}{K_{NT}^{\text{eval}}} + \sum_{i=1}^{N-1} (\hat{w}_i^{\text{ts,eval}})^2 \frac{\hat{\sigma}_i^2(\hat{g})}{K_{iT}^{\text{eval}}} \right) + \rho_K^{\text{eval}} \hat{\mu}(\hat{g})^\top \hat{\Sigma}_{w, \hat{\mathcal{S}}^{\text{eval}}} \hat{\mu}(\hat{g})\end{aligned}$$

Then, as $K_{\min}^{\text{eval}} \rightarrow \infty$, $\hat{V}_{\hat{g}} - V_{\hat{g}} \xrightarrow{P} 0$.

Thus, we have established comprehensive inference results for treatment effect estimators for learned summaries via pre-treatment sample splitting. This approach is highly appealing due to its flexibility; it makes no assumptions on how the algorithm $\mathcal{A}_K$ ingests or uses the discovery-fold data. In particular, this allows for iterative “tuning” of summaries by the researcher, who can show $\mathcal{A}_K$ pre-treatment documents in order to guide search towards summaries that do not differentiate between the treated unit and the synthetic control in the pre-period. This can help ensure that the generated summary is well-balanced in practice by the synthetic control while still enabling valid inference. Indeed, a researcher can generate a hypothesis, assess its pre-period balance between treated unit and synthetic control on the pre-treatment discovery fold, and then proceed to estimate the treatment effect on the evaluation fold. If the generated summary has undesirable pre-period behavior, the researcher can provide this information to $\mathcal{A}_K$ as well and obtain an alternative $\hat{g}$, which can be similarly evaluated. Proceeding iteratively in this manner can ensure that the final learned summary function $\hat{g}$ is well-behaved pre-treatment, while still maintaining valid inference on the evaluation-fold estimate of the treatment effect.

This iterative approach is highly appealing due to the fact that our consistency and inference results for $\hat{\tau}_{\hat{g}}^{\text{eval,split}}$ are relative to the pseudo-treatment effect $\tilde{\tau}_{\hat{g}}$. Consequently, the valid confidence intervals produced via this method do not provide coverage for the true treatment effect, but rather for the difference in mean summary scores between the treated unit and the pseudo-counterfactual. As established in [Proposition 7](#), if the synthetic control model is misspecified (that is, if $\text{MMD}_k(P_0, \tilde{P}_0) > 0$), this estimand may not be equal to the true treatment effect $\tau_{\hat{g}}$. If this misspecification is sufficiently extreme, then any learned summary function $\hat{g}$ may not generate an estimated treatment effect $\hat{\tau}_{\hat{g}}^{\text{eval,split}}$ that is close to $\tau_{\hat{g}}$. Moreover, even if this misspecification is mild and $\text{MMD}_k(P_0, \tilde{P}_0)$ is small, it may be that a given learned summary has large RKHS norm $\|\hat{g}\|_{\mathcal{H}_k}$, or may lie far from $\mathcal{H}_k$, in which case the bias between our estimand $\tilde{\tau}_{\hat{g}}$ and $\tau_{\hat{g}}$ may still be large.

Iterating in the aforementioned manner will provide a diagnostic that may be helpful in guarding against these cases. Poor pre-period balance between the synthetic control and the treated unit, which we can identify by examining pre-period balance of a given summary on the pre-period discovery fold, may be suggestive of either synthetic control misspecification or large norm of the

learned summary. Simply looking at post-period confidence intervals will only assess the null of no effect relative to the pseudo-counterfactual, making these diagnostics crucial for causal interpretation: if pre-period balance does not appear strong for a given summary, it may be an indication that the resulting causal estimate should not be interpreted as a treatment effect relative to the true counterfactual. Pre-period sample splitting enables this diagnostic exercise, allowing for assessment of pre-period behavior while maintaining statistical validity on held-out data. It will therefore be our preferred means of conducting inference.

5 Remote Participation and Local Government Speech

We will now apply our distributional synthetic control estimator to the dataset of California city council meetings compiled by [Martin and Venugopal \(2026\)](#). In that paper, the authors compile a dataset of over 25,000 transcripts of city council meetings from 115 cities in California from 2015-2024, studying the structure and content of meetings and the characteristics of the public participant pool. Their primary causal question is how the staggered removal of remote participation options for public commenters in the aftermath of the COVID-19 pandemic impacted participant demographics. Applying the difference-in-differences estimator from [Callaway and Sant’Anna \(2021\)](#), the authors find that requiring in-person participation led to lower turnout but did not meaningfully change the makeup of the participant pool. In this paper, we ask a natural follow-up question: how does shifting to in-person participation shape *how* constituents speak? Throughout, as our kernel k for weight estimation we use a Gaussian RBF kernel:

$$k(x, x') = \exp\left(-\frac{\|x - x'\|_2^2}{2\sigma^2}\right)$$

Here, the parameter σ is termed the “bandwidth”; we set it equal to the median inter-embedding distance on a pilot sample of data. This kernel is attractive due to its ubiquity and the fact that it is “universal”: any continuous function can be uniformly approximated arbitrarily well by elements of its RKHS, on any compact subset of Euclidean space ([Steinwart and Christmann, 2008](#)).

The section proceeds as follows: we begin by reviewing the institutional setting and key features of the data from the study of [Martin and Venugopal \(2026\)](#). We then conduct two calibrated simulations using real public comment data to demonstrate the ability of our distributional synthetic control estimator to correctly identify true counterfactual distributions and dimensions of treatment effects in the presence of confounding. We conclude by applying our method in two synthetic control configurations encompassing seven treated units, primarily in northern California, in which we find that requiring in-person participation leads to a decrease in hostile, combative, and provocative speech in the following months.

5.1 Institutional Setting and Data

We begin with a review of the institutional setting and data collected in [Martin and Venugopal \(2026\)](#). Under the Brown Act, all local legislative bodies in California are required to deliberate in public and allocate time for the members of the public to provide comment ([Brown, 1953](#)). Crucially, this requirement is not limited to items on the given meeting agenda or subject to a vote by the given body; public participants must be allowed to speak on any issue that falls under the given legislative body’s jurisdiction. As a result, city council meetings in California must include regular public comment periods, in which residents and stakeholders can be heard by the council on virtually any issue of their choosing pertaining to local government or city conditions.

Indeed, this form of participation is ubiquitous: in the data collected in [Martin and Venugopal \(2026\)](#), the authors find over 250,000 distinct public comments, of which over 40% (nearly 115,000) are *unagendized* (that is, unrelated to any issue on the agenda of the given city council meeting). At the onset of the COVID-19 pandemic in March 2020, Governor Newsom issued statewide emergency orders allowing local legislative bodies, including city councils, to meet and allow public comment virtually. Consequently, most cities in the collected dataset switched to remote participation shortly thereafter, allowing members of the public to participate via remote options such as Zoom or telephone. While those cities that shifted online did so at approximately the same time, there is considerable heterogeneity in when cities remove synchronous remote participation options and require that public commenters attend in-person. In [Martin and Venugopal \(2026\)](#), the authors argue that this heterogeneity is poorly explained by city-level observables, and that this source of variation therefore constitutes a natural experiment to examine the impact of the removal of remote participation options on participation.

In this paper, we analyze the causal effect of requiring in-person participation on the speech of public commenters. In order to account for potential effects of selection into participation based on meeting agendas, we will consider only the effect on unagendized public comments. Given occasionally sparse frequency of unagendized comments in meetings, we will aggregate across all meetings in a three-month period within a given municipality as a single “observation”. Each unit-period observation is then given by the empirical distribution over embedding vectors of the unagendized public comments that were present in that city (unit) during that three-month period; embeddings are produced by EmbeddingGemma-300M ([Vera et al., 2025](#)). Treated units will consist of municipalities that remove remote public participation options at some point during the covered window. The treatment date for each treated city is defined as the date of the first city council meeting in which public commenters were required to attend in-person. The donor pool will consist of “never-treated” units, defined as those cities that either still offer remote participation options or removed them after June 2024 (the end of the data coverage).

5.2 Calibrated Simulations

Before applying our estimator to assess the causal effect of requiring in-person public participation on public comment speech, we first perform two calibrated simulations to demonstrate that the method performs well and accurately recovers true weights, counterfactuals, and treatment effect dimensions. In order to make these studies as realistic as possible, both simulations will use embedding vectors for real unagendized public comments in never-treated municipalities. In the first simulation, we will use real donor empirical distributions and construct a synthetic treated unit with known weights, allowing us to assess the estimator’s ability to accurately recover the true counterfactual distribution. In the second simulation, we will construct a “fake” panel of cities where a single city is treated. This treated city will differ from the majority of the donor pool in a selected dimension (mimicking selection into treatment) and will have its post-treatment comments modified to inject a controlled treatment. The goal of this simulation is to show that our estimator adequately accounts for unknown selection effects and returns hypotheses that pertain to the true treatment dimension. In both studies, our primary benchmark model will be one where a counterfactual distribution is constructed via equal weighting of the donor units. In the first simulation, we show that this naive benchmark is unable to recover the true counterfactual except for when the true weights happen to be approximately equal; in the second simulation, it is both unable to distinguish between selection and treatment and, if the selection effect is stronger than treatment, is unable to hypothesize true treatment dimensions at all.

5.2.1 Calibrated Simulation 1: Learning the Correct Counterfactual Distribution

We begin by configuring the simulation. We will implement a pseudo-treatment date of July 1, 2023, with the eight (three-month) periods immediately preceding that date constituting the “pre-treatment”; the two periods immediately following this date are considered “post-treatment”. For our donor pool, we select the ten largest never-treated cities in our sample by their volume of unagendized public comments in the pre-treatment period, spanning July 1, 2021 – July 1, 2023. We then split this comment panel within each unit-period cell equally into *construction* and *estimation* folds.

In order to synthesize our treated unit, we first simulate a true weight vector w^* . In order to mimic varying degrees of sparsity in these true weights, we construct a weight vector via an exponential weighting of city-level, non-text observables, with a controllable correlation strength parameter. Let x_j denote the observable profile of donor unit j , and let x_r^* denote the profile of the treated unit. Since the treated unit is synthetic, this will correspond to the true observable profile of some other non-donor, never-treated reference unit r . Then, define a donor similarity score as the normalized distance between donor unit j and reference (treated) unit r :

$$s_{rj} = -(x_j - x_r^*)^\top \widehat{\Sigma}_x^{-1} (x_j - x_r^*)$$

Here, $\widehat{\Sigma}_x^{-1}$ refers to the estimated covariance matrix of the non-text town-level observables contained in x . Then, the true weights are defined through exponential normalization (softmax):

$$w_{rj}^*(\beta) = \frac{\exp(\beta s_{rj})}{\sum_{\ell=1}^{10} \exp(\beta s_{r\ell})}$$

That is, donor unit j receives weight w_{rj}^* correlated with its observable-profile proximity to the treated unit. The strength of this correlation is governed by the parameter β ; when $\beta = 0$, the produced weights are equal, while increasing β results in the weights concentrating on those donor units with observable profiles most similar to the treated unit. In practice, for these observables we use the following variables obtained from the US Census 2024 American Community Survey (ACS) 1-Year Estimates: population; population density; median age; percentage of population over the age of 65; percentage of the population that is White, Black, and Hispanic; and median income.

For a given reference town r and correlation strength value β , we then sample comments from the construction fold of data from each donor unit, in both the pre- and post-treatment period. This is accomplished by first selecting donor unit j with probability w_{rj}^* and then sampling uniformly at random from its period t construction fold of comments. We repeat this for each time period, constructing a treated unit with 50 sampled comments in each period. We then estimate our distributional weights using the estimation-fold donor unit empirical distributions and the constructed treated unit distribution in each pre-period. Critically, the distributional estimator has no knowledge of the true weights w^* or the underlying observable profiles $x_r, \{x_j\}_{j=1}^{10}$.

The primary measure of performance for the method is then given by the distributional distance, calculated as the 1-Wasserstein distance, between the true treated distribution and the induced mixture over donor distributions in the held-out post-treatment periods. This directly measures the quality of the constructed counterfactual, our primary estimand of interest for the weight-estimation stage. Additionally, in order to validate the ability of this estimator to implicitly balance on summary themes of the kind produced by our downstream hypothesis generation engine, we score each treated unit comment and estimation-fold donor unit comment with a battery of 15 ex-ante specified summary functions, measuring 10 topical themes identified in [Martin and Venugopal \(2026\)](#) and 5 additional stylistic themes encompassing aspects such as sentiment and formality. We can then plot these two performance metrics as we vary the concentration of the true weights w^* by varying β . The results are shown in [Figure 1](#), in which concentration of the weights is measured on the x-axis by the L_2 norm of w^* , which in this setting is equivalent to the square root of the commonly used Herfindahl-Hirschmann Index (HHI).

The results demonstrate that a counterfactual constructed via equal weighting of the donor units only performs well when the true weights are approximately equally distributed. The weights learned by minimizing pre-period MMD, however, perform well across both metrics regardless of the concentration of the weights, with the MMD curve in both panels of [Figure 1](#) appearing




Figure 1: 1-Wasserstein distance between true and estimated counterfactuals (left) and summary score RMSE (right) for estimated counterfactuals produced via equal weighting (orange) and distributional MMD weighting (black). Each point is an average over 20 simulated draws, corresponding to 5 reference profiles r and 4 comment sampling seeds for each. Shaded regions correspond to $\pm$ one standard deviation from the mean.

nearly flat. This demonstrates that our proposed distributional estimator is able to construct accurate counterfactual distributions regardless of how concentrated the weights are, and that these weights also implicitly balance on summary themes in the text that are not provided to the estimator explicitly, in keeping with [Proposition 6](#). For additional details on the simulation, including results showing that the weights and observable profiles themselves are well-approximated by our distributional estimator, see [Appendix D](#).

5.2.2 Calibrated Simulation 2: Recovering an Unknown Treatment Effect

Having established that our distributional, MMD-based estimator accurately recovers weights and implicitly balances on held-out summary themes, we now turn our attention to validating the interface between our first-stage counterfactual construction and second-stage hypothesis generation. Our goal here is to mimic selection into treatment by constructing a synthetic panel of “fake” cities that vary consistently and at baseline according to some chosen aspect of text. For this “selection aspect” we choose the “Housing and Homelessness” topic identified in [Martin and Venugopal \(2026\)](#) as being a significant driver of meeting attendance. We score all comments in 15 never-treated units chosen for their comment volume according to the 10-topic classification outlined in [Martin and Venugopal \(2026\)](#) and construct a panel of synthetic units that have stable topic shares of “Housing and Homelessness”.

We then construct a synthetic panel by sampling comments (without replacement) period-by-period from this pool of never-treated cities. In doing so, we construct two types of units: “low”

types, with “Housing and Homelessness” topic shares no more than 25%, and “high” types, with “Housing and Homelessness” topic shares of at least 60%. We construct our donor pool of 7 low-types and 3 high-types, with our treated unit being a high-type with target “Housing and Homelessness” topic share of 70%. As in the first calibrated simulation, we choose July 1, 2023 as our “treatment onset”, now keeping 6 pre-treatment periods and 2 post-treatment periods, where each period comprises a three-month window, as before. Each donor unit has 20 sampled comments in every period, while the treated unit has 20 sampled comments in each pre-treatment period and 50 in each post-treatment period, to ensure sufficient sample size for hypothesis generation after post-treatment sample splitting. Figure 2 shows the distribution of realized “Housing and Homelessness” topic shares in this synthetic panel.




Figure 2: Realized “Housing and Homelessness” topic shares in the synthetic panel

For post-treatment documents in the treated unit, we impose a treatment of increased formality in the text. This is implemented by asking an LLM (gemini-3.5-flash) to rewrite the sampled untreated comments to make them slightly more formal. This rewritten corpus will now play the role of our observed comment set for the treated unit post-treatment, while the original corpus constitutes our true counterfactual. We observe a measured treatment effect of an increase in formality score of 0.60, measured on a discrete $\{0, 1, 2, 3\}$ scale, where 0 is extremely informal and 3 is extremely formal. By contrast, when we measure the presence of discussion on “Housing and Homelessness” on the same discrete scale (now with 0 corresponding to not at all about “Housing and Homelessness” while 3 corresponds to exclusively about “Housing and Homelessness”), we observe a true mean difference of 1.22 in the original (counterfactual) corpus and 1.14 in the rewritten (observed) corpus, relative to an equal-weighted pool of donor comments. This demonstrates that the treatment injection did not have a meaningful impact on the concentration of “Housing and Homelessness” in the rewritten documents, and this corresponds to a setting with strong selection:

the mean difference in “Housing and Homelessness” is approximately twice as large as the true treatment effect.

As in the previous simulation, we now estimate two sets of synthetic control weights: one via our distributional MMD estimator, and another via equal weighting. We then implement post-treatment sample splitting for our panel, creating discovery and evaluation folds. We then implement our hypothesis generation procedure, using as our algorithm $\mathcal{A}_K$ an LLM that is provided with samples of documents from the treated unit and the estimated counterfactual and is asked to hypothesize differences between them. Each returned hypothesis should be accompanied by a scoring scale, providing a learned function $\hat{g}$ that can then be applied to future documents, in keeping with our theoretical framework. The resulting learned hypotheses can be found in [Table 1](#).

Counterfactual Weighting	Hypothesized Themes
Equal	Focus on housing policy; policy and legal specificity; rhetorical tone; speaker identity and institutional alignment; geographic and operational scope
MMD	Transcription fidelity and cleanliness; formal vs. colloquial register; procedure and interactive meta-talk; argumentation style: policy vs personal narrative

Table 1: Hypotheses generated by equal- and MMD-weighted post-treatment counterfactuals

The hypotheses generated by the distributional MMD-weighted counterfactual include the true treatment of formality and do not include the selection dimension, “Housing and Homelessness”. By contrast, the hypotheses generated by the equal-weighted counterfactual include the selection dimension but not the true treatment dimension. This highlights the key advantage of distributional weighting: by constructing a counterfactual to minimize the distributional distance in pre-treatment text embeddings, we are able to correctly account for and discard selection effects, allowing us to accurately recover treatment dimensions even when the strength of the treatment effect is much less than that of selection. A naive researcher who instead directly compares the post-treatment treated corpus to the full donor pool (corresponding to an equal-weighted counterfactual) will instead be unable to distinguish between selection and treatment effects. Moreover, as in this setting, if the selection effect is sufficiently strong relative to the treatment effect, it may “crowd out” treatment entirely and hypotheses generated from the equal-weighted counterfactual may be unable to recover the true treatment at all.

We can see this dynamic at play in this simulation. [Figure 3](#) plots an event study of the “focus on housing policy” theme learned from the equal-weighted counterfactual, which comes with a discrete $\{0, 1, 2, 3\}$ scoring scale analogous to the oracle selection theme. The left panel shows the equal-weighted counterfactual which exhibits correspondingly poor pre-period balance, leading to a spurious conclusion of a significant effect in the post-period. The right panel, meanwhile, shows

that the text distribution-weighted counterfactual correctly accounts for this baseline difference and shows no significant effect, consistent with the hypothesis generation algorithm not proposing any “Housing and Homelessness”-related themes when given the MMD-weighted counterfactual.⁶




Figure 3: Event studies for learned “focus on housing policy” theme. Post-period estimates are calculated on the evaluation fold, with error bars that are valid 95% confidence intervals conditional on the learned theme and take into account uncertainty in the weights (in the MMD case). Pre-period error bars are diagnostic. Vertical dashed line denotes simulated treatment onset.

A naive researcher who does not construct a counterfactual to account for selection would be confronted with the left panel, realize that it does not show a true treatment effect, and would erroneously conclude that there is no treatment effect (since no formality-related rule was discovered). Implementing the proposed method, however, shows no significant effect for a selection dimension and indeed does not even hypothesize this dimension in the first place. Indeed, it instead accurately hypothesizes the true treatment dimension. This highlights both the importance of constructing intelligent counterfactuals in general and the ability of the proposed method to do so in particular.

5.3 The Effects of Removing Remote Participation

Having established in the preceding simulations that our synthetic control estimator correctly constructs appropriate counterfactuals and leads to learned hypotheses that correctly identify treatment effects and discard selection, we now turn our attention to applying the method to real treated-donor configurations in the data. As previously discussed, our treatment now is the removal of remote participation options for public commenters and our donor pool is the set of cities that implemented remote participation and have not yet removed it (as of the end of our dataset coverage).

⁶Note that since the provided theme was not learned by the MMD-weighted estimator, the resulting error bars are valid even without pre-treatment sample splitting.

Given that we are in a setting with staggered adoption of treatment (i.e. many cities remove remote participation options at various times), there are several different treated units that we could utilize for our synthetic control study. In order to pare this down, we search among the set of possible candidate specifications to identify those that achieve good pre-period balance and generate economically interesting hypotheses. In order to ensure statistical validity of our final treatment effect estimates, we accomplish this by implementing pre-treatment sample splitting and performing this screening step entirely within the discovery half of the data. This is analogous to the iterative hypothesis generation procedure outlined in [subsection 4.3](#). All configurations use three-month aggregated windows of comments within each city as a single period, keeping only those donor units which have a minimum volume of comments per period pre-treatment, to ensure accurate weight estimation. For our two-stage weight estimation procedure, we use a threshold of $\lambda_K = 0.02$, defining as an active donor any unit which has an estimated weight of at least that value and setting all others to zero.⁷ After removing inactive donors which fall below this threshold, we then re-estimate the weights on the remaining active donor set.

Our primary chosen configuration examines the effect of requiring public commenters to attend in-person on public comments in six cities across California. Redwood City, San Ramon, Santa Rosa, San Rafael, Ventura, and Novato, each of which may have insufficient comment volume for feasible estimation on their own, remove remote participation in a one-month period around October 2023. We therefore combine them into a single, aggregated treated unit. In this configuration, we require a minimum of 20 comments per unit-period cell to be an eligible donor and are able to construct a continuous panel with this restriction stretching back to March 2021.




Figure 4: Estimated second-stage MMD weights for donor units for aggregated treated unit. Error bars denote 95% confidence intervals. Donors with estimated first-stage weight under 0.02 (horizontal dashed line) are determined to be inactive and assigned weight 0 in the second stage refit.

⁷This cutoff can be viewed as a specific element of the asymptotic sequence described in [subsection 3.4](#), at the observed sample size.

Figure 4 shows estimated weights for each of the 13 eligible donors for this aggregated treated unit. The estimated weights are highly intuitive: the treated unit is comprised of smaller cities and towns, mostly from northern California, and these are precisely the donor units that receive highest weight (e.g., Palo Alto, San Leandro). Conversely, Los Angeles, a large city in southern California, receives zero weight. We then generate hypotheses by showing documents from this aggregated treated unit and its constructed counterfactual from the post-treatment discovery fold to an LLM (gemini-3.5-flash); for details on this hypothesis generation procedure, including prompting, please see Appendix D. A hypothesis returned by the LLM that is economically interesting and well-balanced in the pre-treatment period in the discovery fold is “confrontational informal tone”. The corresponding event study, displaying the per-period difference-in-means of the learned theme on the held-out evaluation fold, can be found in Figure 5.




Figure 5: Event studies for treatment effect estimates for aggregated treated unit for learned theme “confrontational informal tone”. Post-treatment error bars are valid 95% CIs which incorporate uncertainty in the estimated weights. Pre-period error bars are diagnostic. Vertical dashed line denotes treatment onset, with period 0 corresponding to post-treatment. All estimates are computed on evaluation fold.

We can also repeat the same discovery and estimation procedure for San Jose, which removed remote participation options on January 30, 2024. We construct a synthetic control using 15 cities across the state, each of which has a minimum of 10 unagendized public comments in each of the 8 available pre-treatment periods, stretching back to January 2022. Figure 6 shows the estimated donor weights, with associated uncertainties, for this configuration. San Jose is a large city in the Bay Area, and we immediately see that the learned weights are correspondingly intuitive: most weight is placed on cities in the Bay Area (e.g., San Leandro, Palo Alto, Santa Clara) and larger cities (e.g., Los Angeles). By contrast, the cities we would expect to be most dissimilar to San Jose are smaller cities in southern California, such as Glendale, Burbank, and Pasadena, all of which receive weight zero in this synthetic control.




Figure 6: Estimated second-stage MMD weights for donor units for San Jose. Error bars denote 95% confidence intervals. Donors with estimated first-stage weight under 0.02 (horizontal dashed line) are determined to be inactive and assigned weight 0 in the second stage refit.

Passing the discovery-fold post-treatment corpuses from San Jose and its synthetic control to the LLM yields two hypotheses that are economically interesting and (reasonably) well-balanced in the pre-treatment period: “conspiracy and provocation” and “hostility and profanity”. Event studies for the resulting difference-in-means estimators for each theme can be seen in [Figure 7](#).




Figure 7: Event studies for treatment effect estimates for San Jose for learned themes “conspiracy and provocation” (left) and “hostility and profanity” (right). Post-treatment error bars are valid 95% CIs which incorporate uncertainty in the estimated weights. Pre-period error bars are diagnostic. Vertical dashed lines denote treatment onset, with period 0 corresponding to post-treatment. All estimates are computed on evaluation fold.

Both themes show a marked decline in the post-treatment period, and are fairly well-balanced by the synthetic control estimator in the pre-treatment periods, with most pre-treatment mean differences in the periods preceding treatment indistinguishable from zero. The exception to this is period -8 in the right panel of [Figure 7](#), corresponding to theme “hostility and profanity”. This

point captures comments made between January and April 2022, and the dramatic spike is due to a spate of comments in San Jose regarding the shooting of K’aun Green by the San Jose Police Department on March 27, 2022.⁸ This led to a documented public outcry in subsequent San Jose city council meetings, with more than 40 residents calling into the April 5, 2022 city council meeting to advocate for police reform and accountability.⁹ Inspection of the comments in the impacted period confirm that this advocacy increased the prevalence of comments labeled “hostile and profane”. This constitutes a transitory shock to “hostility and profanity” that dissipates by the next period. Inspection of the post-treatment period comment corpus identifies no comparable transient shock.

Details on the meaning and scoring of each learned theme (provided by the LLM) across these two configurations can be found in [Table 2](#). For example texts exhibiting each theme, see [Appendix D](#).

Learned Theme	Description	Scoring Scale
Confrontational Informal Tone	Score the document based on the level of informal, direct, or confrontational language used by the speaker. Look for slang, expressions of open anger/contempt, or calling out individuals directly with informal phrases.	{0, 1, 2}
Conspiracy and Provocation	Evaluate whether the comment contains bizarre, highly provocative, or unverified conspiracy theories, hate-fueled claims, or extreme historical revisionism.	{0, 1}
Hostility and Profanity	Evaluate the level of hostility, profanity, or abusive language in the public comment.	{0, 1, 2}

Table 2: Hypotheses generated by selected configurations for the treatment effect of the removal of remote participation options. “Confrontational informal tone” is obtained from the study with aggregated treated unit, while “conspiracy and provocation” and “hostility and profanity” are learned from treated unit San Jose.

In all three themes, lower scores indicate a lesser degree of provocative, hostile, or combative speech. As [Figures 5](#) and [7](#) show, all three have estimated negative treatment effects in the post-treatment period that are significant at the 5% level. This fact, coupled with the breadth of cities covered between the two configurations, suggests that the removal of remote participation options led to increased civility and reduced hostility in public comment speech in the months

⁸See related [news coverage](#) for additional information.

⁹See further [news coverage](#) for additional information on the council meeting.

after enactment in the treated cities. We now go on to show that the significance of the effects on “confrontational informal tone” and “hostility and profanity” is further supported by unit-level permutation analyses, in the style of the unit-placebo analyses proposed by [Abadie, Diamond, and Hainmueller \(2010\)](#).

5.4 Supplemental Permutation Studies

In the standard synthetic control literature with scalar outcomes, placebo analyses are a widely-accepted means of examining whether the estimated treatment effect is likely to be spurious ([Abadie, Diamond, and Hainmueller, 2010](#); [Bohn, Lofstrom, and Raphael, 2014](#); [Pinotti, 2015](#); [Cunningham and Shah, 2018](#); [Abadie, 2021](#)). Given that the method in this paper provides analytic inference on the estimated treatment effects, placebo tests are not necessary to conduct inference, unlike standard synthetic control configurations. For robustness, however, we present the following permutation study, which serves as a proxy for a unit-placebo approach to inference. Note that since our inferential target is the treatment effect relative to the pseudo-counterfactual, our causal estimand may not be equal to zero under the null even when no treatment exists. The produced empirical permutation-based distributions throughout this section therefore need not be centered at zero, and diagnostic support for the significance of the observed treatment effect is provided by its relative position to the permutation distribution, not the magnitude of the test statistic.

In this permutation analysis, we select a donor unit and fit a synthetic control for it in the true pre-treatment period, using the rest of the donor pool. The true treated unit is held out throughout. We then construct a counterfactual distribution in the post-treatment period using this synthetic control and pass the resulting corpus of text to the LLM, which then hypothesizes summary functions. Maintaining pre- and post-treatment sample splitting throughout, as detailed in [subsection 4.3](#), we then construct valid standard errors for these estimated effects. We then repeat iteratively across the donor units, permuting the treated label. In order to take the varying precision and scales of each of these estimates into account, we compare studentized values of the permutation-based and true treatment effects, dividing each estimate by its corresponding standard error. The resulting permutation distribution is then at the configuration-level, with each point in that distribution corresponding to the measured value of a (studentized) treatment effect on a particular hypothesized summary function learned for a particular permuted unit. [Figure 8](#) shows permutation distributions of event studies and post-treatment estimated effects for the theme learned on the aggregated treated unit comprised of six smaller cities across California.




Figure 8: Permutation distributions of studentized effects for aggregated treated unit. Left panel: permuted event study trajectories. Dark black line shows the event study for the observed treatment effect, while each light grey line corresponds to a theme learned from a permuted donor unit. Vertical dashed line denotes treatment onset. Right panel: empirical CDF in the post-treatment period (blue curve). Red line shows observed studentized treatment effect estimate for the treated unit.

The left panel of **Figure 8** shows the resulting event studies for each permuted donor unit against that of the true treated unit. The observed trajectory for the treated unit lies in the interior of the permutation distribution throughout, with a sharp drop in the post-treatment period. This reinforces the conclusion provided by our analytical standard errors in **subsection 5.3** that the removal of remote participation options led to a significant decrease in the prevalence of confrontational and informal language in the aggregated treated unit. The post-treatment permutation distribution can be viewed in more detail in the right panel, which shows how the real treatment effect compares to the empirical permutation CDF.

We can repeat the same procedure for San Jose, constructing an analogous permutation distribution and comparing to the observed treatment effects. The resulting event studies and post-treatment period empirical CDFs for the two learned themes are found in **Figure 9**. For both themes, the observed trajectory of the treated unit (shown in the top row) lies at a greater extreme of the permutation distribution in the post-treatment period than in nearly all pre-treatment periods. The empirical CDFs (shown in the bottom row) provide further evidence suggestive of significance, with both “conspiracy and provocation” and “hostility and profanity” exhibiting studentized effects that are unusually extreme relative to the full permutation distribution.




Figure 9: Top: permutation event study trajectories for studentized treatment effect estimates for San Jose. Dark black lines show the event study for the observed treatment effect, while each light grey line corresponds to a theme learned for a permuted donor unit. Vertical dashed lines denote treatment onset. Bottom: empirical permutation CDF for studentized treatment effect estimates for San Jose in the post-treatment period. Blue curve shows the empirical CDF of the permutation distribution, while the vertical red line shows the observed studentized value of the treatment effect on San Jose.

Furthermore, we can extend this permutation study to account for pre-treatment fit for each of the permutation unit-theme pairs. [Abadie, Diamond, and Hainmueller \(2010\)](#) observe that placebo runs with poor pre-treatment fit are not informative as to the likelihood of observing a strong post-treatment effect and good pre-treatment balance, as is the case in our true configurations. We therefore follow their procedure and filter out permutation unit-theme pairs with a mean squared prediction error (MSPE)¹⁰ greater than 5 times that of the treated configuration for the given theme. A full summary of empirical permutation p -values, including for this filtering approach, can be found in [Table 3](#).¹¹ The MSPE-filtered results suggest even more strongly that the observed

¹⁰MSPE is defined as the average squared error between the “treated” unit and synthetic control. In our setting, this corresponds to the average of squared per-period difference-in-means estimates.

¹¹Note that these p -values are not established to be calibrated under the causal null, for the reasons previously stated. They should rather be interpreted as diagnostic empirical lower-tailed proportions. For related discussion on the limitations of placebo-based inference for synthetic control, see [Ferman and Pinto \(2017\)](#).

effects are extreme relative to the permutation distribution, as two of the three learned themes lie even further in the tails after filtering out permutation unit-theme pairs with poor pre-period balance.

Learned Theme	MSPE Filtering	Empirical Permutation p -Value
Confrontational Informal Tone	No	0.077
	Yes	0.000
Conspiracy and Provocation	No	0.200
	Yes	0.100
Hostility and Profanity	No	0.089
	Yes	0.111

Table 3: One-sided empirical permutation p -values with and without MSPE filtering. Each p -value is calculated as the fraction of permutation unit-theme pairs with studentized effect less than or equal to the observed studentized treatment effect.

Collectively, the results of this permutation study reinforce the story told by the analytical standard errors provided in [subsection 5.3](#), in which [Figure 5](#) suggests a strong treatment effect on “confrontational informal tone” and [Figure 7](#) suggests a moderate effect on “conspiracy and provocation” and a strong effect on “hostility and profanity”. This permutation study captures both steps of our pipeline, from counterfactual construction to hypothesis generation and evaluation, and provides diagnostic evidence that the speech in public comments in both treated units became less confrontational and more civil after the removal of remote public participation options. This conclusion is similarly reinforced by an alternative permutation study which conditions on the themes learned from the treated units; see [Appendix D](#) for details.

6 Conclusion

In many settings in social science, researchers are increasingly interested in studying causal effects on text outcomes. Yet, standard approaches typically rely on an ex-ante featurization of text that may strongly restrict the scope of possible causal effects that can be estimated. Concurrently, a burgeoning literature in hypothesis generation has developed methods for learning how treatments have impacted outcomes captured by unstructured data, but applications to text have primarily been limited to settings where treatment assignment is randomized. In the observational settings that prevail in the social sciences, where it is typically infeasible to randomize treatment assignment, construction of counterfactuals that can account for selection and confounding effects is of paramount importance, but constructing interpretable counterfactuals for text data remains an open problem.

In this paper, I address this gap in the literature by developing a two-stage pipeline for hy-

pothesis generation and validation in text-as-outcome causal inference settings with nonrandom treatment assignment in panel data. In the first stage, I propose a distributional synthetic control estimator in which observations are distributions over text embeddings and a counterfactual is estimated as a mixture over contemporaneous donor distributions, with weights estimated to minimize average pre-period squared maximum mean discrepancy (MMD). I show that these estimated weights, and the induced counterfactual distribution, are consistent for their population analogs under correct specification of the synthetic control model, and provide bias bounds under misspecification. In an unconventional asymptotic regime in which the size of the panel matrix remains fixed while the number of observations per cell grows large, I establish asymptotic normality of the estimated weights, a novel result in the context of synthetic control.

In the second stage, I propose an algorithm for hypothesis generation and validation. Real text documents are sampled from the estimated counterfactual and, along with documents from the treated unit, passed to a hypothesis generation technology, such as an LLM. This technology then returns summary functions of the text which capture aspects that are believed to differentiate between these two corpuses. Utilizing sample splitting in the post-treatment period, the mean differences of these hypothesized summary functions can be calculated on held-out post-treatment data, yielding estimates of the treatment effect which I show are consistent under correct specification of the synthetic control model and have bounded bias otherwise. Additionally, I establish asymptotic normality of these treatment effect estimates via pre-treatment sample splitting, providing analytic inference and standard errors which are consistently estimable and take into account uncertainty in the first-stage counterfactual construction.

I highlight the practical appeal of my method in an application to California city council meeting comment periods. In two calibrated simulations, I demonstrate that my proposed method correctly recovers true weights and counterfactuals, and that it constructs counterfactuals that correct for selection effects, recovering true treatment effect hypotheses in the second stage. In two synthetic control configurations with a real treatment of the removal of remote participation options for public commenters, I find that requiring in-person participation leads to increased civility and reduced hostility in public comments. These results are novel findings at this scale and demonstrate the ability of this method to provide new insight into previously intractable problems in the social sciences.

References

- Abadie, Alberto (2021). Using Synthetic Controls: Feasibility, Data Requirements, and Methodological Aspects. *Journal of Economic Literature*, 59(2):391–425. (Cited on pages 17 and 44.)
- Abadie, Alberto, Alexis Diamond, and Jens Hainmueller (2010). Synthetic Control Methods for Comparative Case Studies: Estimating the Effect of California’s Tobacco Control Program. *Journal of the American Statistical Association*, 105(490):493–505. (Cited on pages 2, 3, 6, 8, 10, 21, 44, and 46.)
- Abadie, Alberto, Alexis Diamond, and Jens Hainmueller (2015). Comparative Politics and the Synthetic Control Method. *American Journal of Political Science*, 59(2):495–510. (Cited on page 3.)
- Abadie, Alberto and Javier Gardeazabal (2003). The Economic Costs of Conflict: A Case Study of the Basque Country. *American Economic Review*, 93(1):113–132. (Cited on pages 3 and 8.)
- Almeida, Alexander, Susan Athey, Guido Imbens, Eva Lestant, and Alexia Olaizola (2025). Estimating Variances for Causal Panel Data Estimators. arXiv:2510.11841. (Cited on pages 2, 11, and 17.)
- Ash, Elliott, Daniel L Chen, and Suresh Naidu (2026). Ideas Have Consequences: The Impact of Law and Economics on American Justice. *The Quarterly Journal of Economics*, 141(1):845–887. (Cited on page 1.)
- Athey, Susan, Guido W. Imbens, and Stefan Wager (2018). Approximate Residual Balancing: Debiased Inference of Average Treatment Effects in High Dimensions. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 80(4):597–623. (Cited on page 4.)
- Belloni, Alexandre and Victor Chernozhukov (2013). Least squares after model selection in high-dimensional sparse models. *Bernoulli*, 19(2):521–547. (Cited on page 18.)
- Bloom, Nicholas, Ruobing Han, and James Liang (2024). Hybrid working from home improves retention without damaging performance. *Nature*, 630(8018):920–925. (Cited on page 4.)
- Bloom, Nicholas, James Liang, John Roberts, and Zhichun Jenny Ying (2015). Does Working from Home Work? Evidence from a Chinese Experiment*. *The Quarterly Journal of Economics*, 130(1):165–218. (Cited on page 4.)
- Bohn, Sarah, Magnus Lofstrom, and Steven Raphael (2014). Did the 2007 Legal Arizona Workers Act Reduce the State’s Unauthorized Immigrant Population? *The Review of Economics and Statistics*, 96(2):258–269. (Cited on page 44.)

- Brown, Ralph M. (1953). Ralph M. Brown Act. (Cited on page 33.)
- Brown, Tom, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei (2020). Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc. (Cited on page 5.)
- Brucks, Melanie S. and Jonathan Levav (2022). Virtual communication curbs creative idea generation. *Nature*, 605(7908):108–112. (Cited on page 4.)
- Callaway, Brantly and Pedro H.C. Sant’Anna (2021). Difference-in-Differences with multiple time periods. *Journal of Econometrics*, 225(2):200–230. (Cited on page 32.)
- Chen, Jiafeng and Suproteem Sarkar (2020). A Semantic Approach to Financial Fundamentals. In Chen, Chung-Chi, Hen-Hsen Huang, Hiroya Takamura, and Hsin-Hsi Chen, editors, *Proceedings of the Second Workshop on Financial Technology and Natural Language Processing*, pages 22–26, Kyoto, Japan. -. (Cited on page 4.)
- Cunningham, Scott and Manisha Shah (2018). Decriminalizing Indoor Prostitution: Implications for Sexual Violence and Public Health. *The Review of Economic Studies*, 85(3):1683–1715. (Cited on page 44.)
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics. (Cited on page 5.)
- Doudchenko, Nikolay and Guido W. Imbens (2016). Balancing, Regression, Difference-In-Differences and Synthetic Control Methods: A Synthesis. (Cited on page 4.)
- Egami, Naoki, Christian J. Fong, Justin Grimmer, Margaret E. Roberts, and Brandon M. Stewart (2022). How to make causal inferences using texts. *Science Advances*, 8(42):eabg2652. (Cited on page 4.)
- Einstein, Katherine Levine, David Glick, Luisa Godinez Puig, and Maxwell Palmer (2023). Still Muted: The Limited Participatory Democracy of Zoom Public Meetings. *Urban Affairs Review*, 59(4):1279–1291. (Cited on page 4.)

- Einstein, Katherine Levine, Maxwell Palmer, and David M. Glick (2019). Who Participates in Local Government? Evidence from Meeting Minutes. *Perspectives on Politics*, 17(1):28–46. (Cited on page 4.)
- Feder, Amir, Katherine A. Keith, Emaad Manzoor, Reid Pryzant, Dhanya Sridhar, Zach Wood-Doughty, Jacob Eisenstein, Justin Grimmer, Roi Reichart, Margaret E. Roberts, Brandon M. Stewart, Victor Veitch, and Diyi Yang (2022). Causal Inference in Natural Language Processing: Estimation, Prediction, Interpretation and Beyond. *Transactions of the Association for Computational Linguistics*, 10:1138–1158. (Cited on page 4.)
- Feldman, Omri, Amar Venugopal, Jann Spiess, and Amir Feder (2026). Causal Effect Estimation with Latent Textual Treatments. (Cited on page 4.)
- Ferman, Bruno and Cristine Pinto (2017). Placebo Tests for Synthetic Controls. (Cited on page 46.)
- Fong, Christian and Justin Grimmer (2023). Causal Inference with Latent Treatments. *American Journal of Political Science*, 67(2):374–389. (Cited on page 4.)
- Freyaldenhoven, Simon (2026). When Can We Work in Embedding Space? What Text Embeddings Preserve. arXiv:2608.31059 [econ.EM]. (Cited on page 4.)
- Gentzkow, Matthew, Jesse M. Shapiro, and Matt Taddy (2019). Measuring Group Differences in High-Dimensional Choices: Method and Application to Congressional Speech. *Econometrica*, 87(4):1307–1340. (Cited on page 1.)
- Gretton, Arthur, Karsten Borgwardt, Malte Rasch, Bernhard Schölkopf, and Alex Smola (2006). A Kernel Method for the Two-Sample-Problem. In *Advances in Neural Information Processing Systems*, volume 19. MIT Press. (Cited on page 6.)
- Gunsilius, F. F. (2023). Distributional Synthetic Controls. *Econometrica*, 91(3):1105–1117. (Cited on pages 3 and 4.)
- Hazlett, Chad (2020). Kernel Balancing: A Flexible Non-Parametric Weighting Procedure for Estimating Causal Effects. *Statistica Sinica*, 30(3):1155–1189. (Cited on pages 4 and 7.)
- Hazlett, Chad and Yiqing Xu (2018). Trajectory Balancing: A General Reweighting Approach to Causal Inference With Time-Series Cross-Sectional Data. *SSRN Electronic Journal*. (Cited on pages 4 and 7.)
- Hoberg, Gerard and Gordon M. Phillips (2011). Text-Based Network Industries and Endogenous Product Differentiation. *SSRN Electronic Journal*. (Cited on page 1.)
- Imai, Kosuke and Kentaro Nakamura (2024). Causal Representation Learning with Generative Artificial Intelligence: Application to Texts as Treatments. arXiv:2410.00903. (Cited on page 4.)

- Kallus, Nathan (2017). A Framework for Optimal Matching for Causal Inference. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, pages 372–381. PMLR. (Cited on page 4.)
- Liu, Yiqi (2025). Synthetic Parallel Trends. arXiv:2511.05870 [econ.EM]. (Cited on pages 2, 3, and 11.)
- Luca, Michael (2011). Reviews, Reputation, and Revenue: The Case of Yelp.Com. *SSRN Electronic Journal*. (Cited on page 1.)
- Ludwig, Jens and Sendhil Mullainathan (2024). Machine Learning as a Tool for Hypothesis Generation. *The Quarterly Journal of Economics*, 139(2):751–827. (Cited on pages 1 and 4.)
- Ludwig, Jens, Sendhil Mullainathan, and Ashesh Rambachan (2025). Large Language Models: An Applied Econometric Framework. arXiv:2412.07031. (Cited on pages 1 and 4.)
- Martin, Olivia and Amar Venugopal (2026). Participation and Representation in Local Government Speech. arXiv:2604.21202. (Cited on pages 3, 4, 32, 33, 35, and 36.)
- Mikolov, Tomas, Kai Chen, Greg Corrado, and Jeffrey Dean (2013). Efficient Estimation of Word Representations in Vector Space. arXiv:1301.3781 [cs]. (Cited on pages 1 and 5.)
- Modarressi, Iman, Jann Spiess, and Amar Venugopal (2025). Causal Inference on Outcomes Learned from Text. arXiv:2503.00725. (Cited on pages 1 and 4.)
- Muennighoff, Niklas, Nouamane Tazi, Loic Magne, and Nils Reimers (2023). MTEB: Massive Text Embedding Benchmark. In Vlachos, Andreas and Isabelle Augenstein, editors, *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 2014–2037, Dubrovnik, Croatia. Association for Computational Linguistics. (Cited on page 5.)
- Newey, Whitney K. and Daniel McFadden (1994). Chapter 36 Large sample estimation and hypothesis testing. In *Handbook of Econometrics*, volume 4, pages 2111–2245. Elsevier. (Cited on pages 13 and 14.)
- Pennington, Jeffrey, Richard Socher, and Christopher Manning (2014). GloVe: Global Vectors for Word Representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, Doha, Qatar. Association for Computational Linguistics. (Cited on pages 1 and 5.)
- Pinotti, Paolo (2015). The Economic Costs of Organised Crime: Evidence from Southern Italy. *The Economic Journal*, 125(586):F203–F232. (Cited on page 44.)
- Porcu, Emilio, Moreno Bevilacqua, Robert Schaback, and Chris J. Oates (2024). The Matérn Model: A Journey Through Statistics, Numerical Analysis and Machine Learning. *Statistical Science*, 39(3):469–492. (Cited on page 22.)

- Pötscher, Benedikt M. and Ulrike Schneider (2011). Distributional results for thresholding estimators in high-dimensional Gaussian regression models. *Electronic Journal of Statistics*, 5(none):1876–1934. (Cited on page 19.)
- Reimers, Nils and Iryna Gurevych (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In Inui, Kentaro, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China. Association for Computational Linguistics. (Cited on page 5.)
- Robbins, Michael W., Jessica Saunders, and Beau Kilmer (2017). A Framework for Synthetic Control Methods With High-Dimensional, Micro-Level Data: Evaluating a Neighborhood-Specific Crime Intervention. *Journal of the American Statistical Association*, 112(517):109–126. (Cited on page 3.)
- Rubin, Donald B. (1974). Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66(5):688–701. (Cited on page 6.)
- Santra, Diptanil, Guanhua Chen, and Chan Park (2026). Distributional Balancing for Causal Inference: A Unified Framework via Characteristic Function Distance. arXiv:2601.15449. (Cited on page 4.)
- Sarkar, Suproteem (2026). Economic Representations. (Cited on page 4.)
- Shapiro, Adam Hale and Daniel J Wilson (2022). Taking the Fed at its Word: A New Approach to Estimating Central Bank Objectives using Text Analysis. *The Review of Economic Studies*, 89(5):2768–2805. (Cited on page 1.)
- Shi, Claudia, Dhanya Sridhar, Vishal Misra, and David Blei (2022). On the Assumptions of Synthetic Control Methods. In *Proceedings of The 25th International Conference on Artificial Intelligence and Statistics*, pages 7163–7175. PMLR. (Cited on page 3.)
- Sriperumbudur, Bharath K., Arthur Gretton, Kenji Fukumizu, Bernhard Schölkopf, and Gert R. G. Lanckriet (2010). Hilbert Space Embeddings and Metrics on Probability Measures. *Journal of Machine Learning Research*, 11(50):1517–1561. (Cited on page 8.)
- Steinwart, Ingo and Andreas Christmann (2008). Kernels and Reproducing Kernel Hilbert Spaces. In Steinwart, Ingo and Andreas Christmann, editors, *Support Vector Machines*, pages 110–163. Springer, New York, NY. (Cited on pages 8 and 32.)
- Sun, Hong-Wei and Ding-Xuan Zhou (2008). Reproducing Kernel Hilbert Spaces Associated with Analytic Translation-Invariant Mercer Kernels. *Journal of Fourier Analysis and Applications*, 14(1):89–101. (Cited on page 22.)

- Sun, Liyang, Eli Ben-Michael, and Avi Feller (2025). Using Multiple Outcomes to Improve the Synthetic Control Method. *Review of Economics and Statistics*, pages 1–29. (Cited on page 3.)
- Tian, Wei, Seojeong Lee, and Valentyn Panchenko (2026). Synthetic controls with multiple outcomes. *The Econometrics Journal*, page utag005. (Cited on page 3.)
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin (2017). Attention is All you Need. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc. (Cited on pages 1 and 5.)
- Vera, Henrique Schechter, Sahil Dua, Biao Zhang, Daniel Salz, Ryan Mullins, Sindhu Raghuram Panyam, Sara Smoot, Iftexhar Naim, Joe Zou, Feiyang Chen, Daniel Cer, Alice Lisak, Min Choi, Lucas Gonzalez, Omar Sanseviero, Glenn Cameron, Ian Ballantyne, Kat Black, Kaifeng Chen, Weiyi Wang, Zhe Li, Gus Martins, Jinhyuk Lee, Mark Sherwood, Juyeong Ji, Renjie Wu, Jingxiao Zheng, Jyotinder Singh, Abheesht Sharma, Divyashree Sreepathihalli, Aashi Jain, Adham Elarabawy, A. J. Co, Andreas Doumanoglou, Babak Samari, Ben Hora, Brian Potetz, Dahun Kim, Enrique Alfonseca, Fedor Moiseev, Feng Han, Frank Palma Gomez, Gustavo Hernández Ábrego, Heseng Zhang, Hui Hui, Jay Han, Karan Gill, Ke Chen, Koert Chen, Madhuri Shanbhogue, Michael Boratko, Paul Suganthan, Sai Meher Karthik Duddu, Sandeep Mariserla, Setareh Ariafer, Shanfeng Zhang, Shijie Zhang, Simon Baumgartner, Sonam Goenka, Steve Qiu, Tanmaya Dabral, Trevor Walker, Vikram Rao, Waleed Khawaja, Wenlei Zhou, Xiaoli Ren, Ye Xia, Yichang Chen, Yi-Ting Chen, Zhe Dong, Zhongli Ding, Francesco Visin, Gaël Liu, Jiageng Zhang, Kathleen Kenealy, Michelle Casbon, Ravin Kumar, Thomas Mesnard, Zach Gleicher, Cormac Brick, Olivier Lacombe, Adam Roberts, Qin Yin, Yunhsuan Sung, Raphael Hoffmann, Tris Warkentin, Armand Joulin, Tom Duerig, and Mojtaba Seyedhosseini (2025). EmbeddingGemma: Powerful and Lightweight Text Representations. arXiv:2509.20354 [cs.CL]. (Cited on page 33.)
- Warshaw, Christopher (2019). Local Elections and Representation in the United States. *Annual Review of Political Science*, 22(1):461–479. (Cited on page 4.)
- Zubizarreta, José R. (2015). Stable Weights that Balance Covariates for Estimation With Incomplete Outcome Data. *Journal of the American Statistical Association*, 110(511):910–922. (Cited on page 4.)

A Proofs for Section 2

A.1 Useful Lemmas

Lemma 1 (Objective function continuity). *Suppose **Assumption k-Reg** holds. Then the population objective $L(w)$ is continuous in w on $\mathcal{W}$.*

Proof. By definition of MMD_k and linearity of the kernel mean embedding, the population objective can be rewritten as the sum of squared norms:

$$L(w) = \sum_{t < T} \left\| \mu_{Nt} - \sum_{i=1}^{N-1} w_i \mu_{it} \right\|_{\mathcal{H}_k}^2$$

For each $t < T$, define A_t such that $A_t w = \sum_{i=1}^{N-1} w_i \mu_{it}$. Since the μ_{it} are fixed elements of $\mathcal{H}_k$, A_t is a bounded linear map from $\mathbb{R}^{N-1}$ to $\mathcal{H}_k$. Consequently, the map

$$w \mapsto \mu_{Nt} - \sum_{i=1}^{N-1} w_i \mu_{it}$$

is affine and continuous. Since the squared norm is continuous on $\mathcal{H}_k$, each element of the above sum is a continuous quadratic function of w . Since the number of pre-treatment periods is fixed, L is the finite sum of continuous functions and is itself continuous on $\mathcal{W}$. $\square$

A.2 Uniqueness of induced pseudo-counterfactual

Lemma 2 (Unique pseudo-counterfactual). *Suppose **Assumptions k-Reg** and **PP** hold. Then, $\mathcal{W}_0$ is nonempty and convex, and for every pair of population minimizers $w, w' \in \mathcal{W}_0$:*

$$\sum_{i=1}^{N-1} w_i F_{iT}(0) = \sum_{i=1}^{N-1} w'_i F_{iT}(0)$$

Proof. First, note that by **Lemma 1**, L is continuous, while $\mathcal{W}$ is compact by definition of the simplex. L therefore attains its minimum on $\mathcal{W}$ by the extreme value theorem, and so $\mathcal{W}_0$ is nonempty.

Now, let $\mu_k(P)$ denote the kernel mean embedding of distribution P and define:

$$\Psi(w) = \left(\sum_{i=1}^{N-1} w_i \mu_k(F_{it}(0)) \right)_{t < T}$$

Similarly, we can stack the kernel mean embeddings of the treated unit:

$$m = (\mu_k(F_{Nt}(0)))_{t < T}$$

Our objective function [Equation 1](#) can then be written as:

$$L(w) = \sum_{t < T} \|m_t - \Psi_t(w)\|_{\mathcal{H}_k}^2$$

Take $w, w' \in \mathcal{W}_0$. Convexity of $\mathcal{W}$ implies that:

$$\bar{w} = \frac{w + w'}{2} \in \mathcal{W}$$

We can then write:

$$\begin{aligned} L(\bar{w}) &= \sum_{t < T} \|m_t - \Psi_t(\bar{w})\|_{\mathcal{H}_k}^2 \\ &= \sum_{t < T} \left\| m_t - \frac{1}{2}\Psi_t(w) - \frac{1}{2}\Psi_t(w') \right\|_{\mathcal{H}_k}^2 \\ &= \sum_{t < T} \left\| \frac{1}{2}(m_t - \Psi_t(w)) + \frac{1}{2}(m_t - \Psi_t(w')) \right\|_{\mathcal{H}_k}^2 \\ &= \underbrace{\frac{1}{2} \sum_{t < T} \|m_t - \Psi_t(w)\|_{\mathcal{H}_k}^2}_{L(w)} + \underbrace{\frac{1}{2} \sum_{t < T} \|m_t - \Psi_t(w')\|_{\mathcal{H}_k}^2}_{L(w')} - \frac{1}{4} \sum_{t < T} \|\Psi_t(w) - \Psi_t(w')\|_{\mathcal{H}_k}^2 \\ &= \min_{\nu \in \mathcal{W}} L(\nu) - \frac{1}{4} \sum_{t < T} \|\Psi_t(w) - \Psi_t(w')\|_{\mathcal{H}_k}^2 \end{aligned}$$

Where we have used linearity of Ψ_t , the squared-norm identity for Hilbert spaces, and the fact that since $w, w' \in \mathcal{W}_0$, $L(w) = L(w')$. Since $\bar{w} \in \mathcal{W}$, it must be the case that the final term above is equal to zero:

$$\frac{1}{4} \sum_{t < T} \|\Psi_t(w) - \Psi_t(w')\|_{\mathcal{H}_k}^2 = 0$$

Since this is a sum of nonnegative values, it must be that every element of the sum is equal to zero, which in turn implies that $\Psi_t(w) = \Psi_t(w')$ for every $t < T$. For any $\alpha \in [0, 1]$, this then yields:

$$\Psi_t(\alpha w + (1 - \alpha)w') = \alpha \Psi_t(w) + (1 - \alpha) \Psi_t(w') = \Psi_t(w) = \Psi_t(w')$$

Therefore, any convex combination of minimizers achieves the same value of the objective function, and so is itself a minimizer. The set $\mathcal{W}_0$ is therefore convex.

Finally, because k is characteristic by [Assumption k-Reg](#), $\Psi(w)$ and $\Psi(w')$ being identical in each coordinate implies that their underlying distributions are identical, yielding for every $t < T$:

$$\sum_{i=1}^{N-1} w_i F_{it}(0) = \sum_{i=1}^{N-1} w'_i F_{it}(0)$$

By **Assumption PP**, this then implies that:

$$\sum_{i=1}^{N-1} w_i F_{iT}(0) = \sum_{i=1}^{N-1} w'_i F_{iT}(0)$$

□

A.3 Factor model for correct specification

Lemma 3 (Factor model). *Suppose that untreated distributions for every unit i and period t are generated by latent mixture components:*

$$F_{it}(0) = \sum_{b=1}^B \lambda_{ib} Q_{bt}.$$

The Q_{bt} are time-varying component distributions and $\lambda_i \in \Delta^B$ are unit-varying mixture shares. Suppose further that $\exists w^ \in \mathcal{W}$ such that*

$$\lambda_N = \sum_{i=1}^{N-1} w_i^* \lambda_i.$$

*Then, **Assumption CM** holds.*

Proof of Lemma 3. By our first assumption, we can write the untreated potential distribution for the treated unit at time t as:

$$F_{Nt}(0) = \sum_{b=1}^B \lambda_{Nb} Q_{bt}$$

By our second assumption, we can express each dimension of the mixture shares for unit N as:

$$\lambda_{Nb} = \sum_{i=1}^{N-1} w_i^* \lambda_{ib}$$

Substituting this into our first equation yields:

$$F_{Nt}(0) = \sum_{b=1}^B \left(\sum_{i=1}^{N-1} w_i^* \lambda_{ib} \right) Q_{bt}$$

Swapping the order of summation yields:

$$F_{Nt}(0) = \sum_{i=1}^{N-1} w_i^* \underbrace{\left(\sum_{b=1}^B \lambda_{ib} Q_{bt} \right)}_{F_{it}(0)}$$

Where the parenthetical term is equal to $F_{it}(0)$ by assumption. This establishes the desired result. $\square$

A.4 Proof of Proposition 1

Proof. By definition of our objective function,

$$L(w^*) = \sum_{t < T} \text{MMD}_k^2 \left(F_{Nt}(0), \sum_{i=1}^{N-1} w_i^* F_{it}(0) \right)$$

By **Assumption CM**, for each $t < T$, the arguments of MMD_k^2 in the above expression are equal. We then have:

$$L(w^*) = \sum_{t < T} \text{MMD}_k^2(F_{Nt}(0), F_{Nt}(0))$$

By definition of MMD_k , we have:

$$\text{MMD}_k(P, Q) = \|\mu_P - \mu_Q\|_{\mathcal{H}_k}$$

Where μ_P refers to the kernel mean embedding of P . Consequently, $\text{MMD}_k(P, Q) = 0 \iff \mu_P = \mu_Q$. Since k is characteristic by **Assumption k-Reg**, the kernel mean embedding mapping is injective:

$$\mu_P = \mu_Q \iff P = Q$$

Therefore, $\text{MMD}_k(P, Q) = 0 \iff P = Q$. Therefore, each term in the above sum is 0, and so $L(w^*) = 0$. Since the above is a sum of squared norms, it is always nonnegative. Therefore, for every $w \in \mathcal{W}$, $L(w^*) = 0 \leq L(w)$. This means that

$$w^* \in \arg \min_{w \in \mathcal{W}} L(w)$$

Establishing the first claim. Now, consider an arbitrary population minimizer

$$w^0 \in \arg \min_{w \in \mathcal{W}} L(w)$$

We must then have $L(w^0) = L(w^*) = 0$. Since $L(w^0)$ is the sum of nonnegative terms, as previously argued, it must be that every element of the sum is 0. By the previous injectivity argument, this in turn means that for every $t < T$:

$$F_{Nt}(0) = \sum_{i=1}^{N-1} w_i^0 F_{it}(0)$$

Applying **Assumption CM** to the left hand side yields:

$$\sum_{i=1}^{N-1} w_i^* F_{it}(0) = \sum_{i=1}^{N-1} w_i^0 F_{it}(0)$$

Assumption PP then yields the same equality in period T :

$$\sum_{i=1}^{N-1} w_i^* F_{iT}(0) = \sum_{i=1}^{N-1} w_i^0 F_{iT}(0)$$

Applying **Assumption CM** to the left hand side once again yields:

$$F_{NT}(0) = \sum_{i=1}^{N-1} w_i^0 F_{iT}(0)$$

Since the choice of w^0 was arbitrary, this proves the second claim. □

A.5 Proof of **Proposition 2**

Lemma 1 establishes continuity of L on $\mathcal{W}$. Since $\mathcal{W}$ is the simplex, it is compact, and so L attains its minimum on $\mathcal{W}$. To establish uniqueness, take distinct $w, w' \in \mathcal{W}$, let $h = w' - w \neq 0$, and fix $\alpha \in (0, 1)$. Because both weight vectors belong to the simplex, $\mathbf{1}^\top h = 0$. Since $h \neq 0$, the contrapositive of **Assumption IDT** and the fact that $\mathbf{1}^\top h = 0$ imply that for at least one period $t < T$,

$$\sum_{i=1}^{N-1} h_i F_{it}(0) \neq 0 \implies \sum_{i=1}^{N-1} w_i F_{it}(0) \neq \sum_{i=1}^{N-1} w'_i F_{it}(0)$$

Since k is characteristic by **Assumption k-Reg**, this in turn implies:

$$\sum_{i=1}^{N-1} w_i \mu_{it} \neq \sum_{i=1}^{N-1} w'_i \mu_{it}$$

Where μ_{it} is the kernel mean embedding of $F_{it}(0)$. This then means that

$$\left\| \sum_{i=1}^{N-1} h_i \mu_{it} \right\|_{\mathcal{H}_k}^2 = \left\| \sum_{i=1}^{N-1} w_i \mu_{it} - \sum_{i=1}^{N-1} w'_i \mu_{it} \right\|_{\mathcal{H}_k}^2 > 0$$

We can then write:

$$\begin{aligned}
L(\alpha w + (1 - \alpha)w') &= \sum_{t < T} \text{MMD}_k^2 \left(F_{Nt}(0), \sum_{i=1}^{N-1} (\alpha w_i + (1 - \alpha)w'_i) F_{it}(0) \right) \\
&= \sum_{t < T} \left\| \mu_{Nt} - \sum_{i=1}^{N-1} (\alpha w_i + (1 - \alpha)w'_i) \mu_{it} \right\|_{\mathcal{H}_k}^2 \\
&= \sum_{t < T} \left\| \alpha \left(\mu_{Nt} - \sum_{i=1}^{N-1} w_i \mu_{it} \right) + (1 - \alpha) \left(\mu_{Nt} - \sum_{i=1}^{N-1} w'_i \mu_{it} \right) \right\|_{\mathcal{H}_k}^2 \\
&= \alpha L(w) + (1 - \alpha) L(w') - \alpha(1 - \alpha) \sum_{t < T} \left\| \sum_{i=1}^{N-1} h_i \mu_{it} \right\|_{\mathcal{H}_k}^2 \\
&< \alpha L(w) + (1 - \alpha) L(w')
\end{aligned}$$

Note that the final equality uses the squared-norm identity for Hilbert spaces. Thus, L is strictly convex on $\mathcal{W}$. It therefore has exactly one minimizer on $\mathcal{W}$, establishing uniqueness.

B Proofs for Section 3

B.1 Useful Lemmas

Lemma 4 (Uniform convergence of empirical objective). *Suppose Assumptions CS and k-Reg hold. Then, as $K_{\min}^{\text{pre}} \rightarrow \infty$,*

$$\sup_{w \in \mathcal{W}} \left| \widehat{L}(w) - L(w) \right| \xrightarrow{p} 0$$

Proof of Lemma 4. Let $\phi(x) = k(x, \cdot) \in \mathcal{H}_k$ denote the feature map induced by kernel k . By Assumption k-Reg, k is bounded and measurable and $\mathcal{H}_k$ is separable. This means that ϕ is measurable, and that

$$\|\phi(x)\|_{\mathcal{H}_k}^2 = k(x, x) \leq \bar{k}$$

By Assumption CS, for every pre-treatment cell (i, t) all observations are iid draws from $F_{it}(0)$. Centering this feature map gives:

$$u_{it} = \widehat{\mu}_{it} - \mu_{it} = \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \phi(x_{itr}) - \mu_{it}$$

Where we use the definition of $\widehat{\mu}_{it}$ in the second equality. We can then write:

$$\mathbb{E}[\|u_{it}\|_{\mathcal{H}_k}^2] = \frac{1}{K_{it}^2} \sum_{r=1}^{K_{it}} \sum_{s=1}^{K_{it}} \mathbb{E}[\langle \phi(X_{itr}) - \mu_{it}, \phi(X_{its}) - \mu_{it} \rangle_{\mathcal{H}_k}]$$

If $r \neq s$, independence and centering yield:

$$\begin{aligned}
\mathbb{E}[\langle \phi(X_{itr}) - \mu_{it}, \phi(X_{its}) - \mu_{it} \rangle_{\mathcal{H}_k}] &= \mathbb{E}[\mathbb{E}[\langle \phi(X_{itr}) - \mu_{it}, \phi(X_{its}) - \mu_{it} \rangle_{\mathcal{H}_k} | \phi(X_{itr}) - \mu_{it}]] \\
&= \mathbb{E}[\langle \phi(X_{itr}) - \mu_{it}, \mathbb{E}[\phi(X_{its}) - \mu_{it} | \phi(X_{itr}) - \mu_{it}] \rangle_{\mathcal{H}_k}] \\
&= \mathbb{E}[\langle \phi(X_{itr}) - \mu_{it}, \mathbb{E}[\phi(X_{its}) - \mu_{it}] \rangle_{\mathcal{H}_k}] \\
&= 0
\end{aligned}$$

Thus all cross terms vanish, yielding:

$$\begin{aligned}
\mathbb{E}[\|u_{it}\|_{\mathcal{H}_k}^2] &= \frac{1}{K_{it}^2} \sum_{r=1}^{K_{it}} \mathbb{E}[\|\phi(X_{itr}) - \mu_{it}\|_{\mathcal{H}_k}^2] \\
&= \frac{1}{K_{it}} \mathbb{E}[\|\phi(X) - \mu_{it}\|_{\mathcal{H}_k}^2] \\
&\leq \frac{1}{K_{it}} \mathbb{E}[\|\phi(X)\|_{\mathcal{H}_k}^2] \\
&\leq \frac{\bar{k}}{K_{it}}
\end{aligned}$$

Where we have made use of the Hilbert space variance identity and boundedness of k via **Assumption k-Reg**. Therefore, as $K_{\min}^{\text{pre}} \rightarrow \infty$, $\|u_{it}\|_{\mathcal{H}_k} \xrightarrow{p} 0$ for every i and $t < T$. Because N and T are fixed, this convergence holds jointly over the finite collection of pre-treatment cells:

$$\max_{\substack{i=1,\dots,N \\ t < T}} \|u_{it}\|_{\mathcal{H}_k} \xrightarrow{p} 0$$

Now, for each $t < T$ and $w \in \mathcal{W}$, define the following:

$$r_t(w) = \mu_{Nt} - \sum_{i=1}^{N-1} w_i \mu_{it} \quad \hat{r}_t(w) = \hat{\mu}_{Nt} - \sum_{i=1}^{N-1} w_i \hat{\mu}_{it}$$

The triangle inequality then yields:

$$\begin{aligned}
\sup_{w \in \mathcal{W}} \|\hat{r}_t(w) - r_t(w)\|_{\mathcal{H}_k} &\leq \|u_{Nt}\|_{\mathcal{H}_k} + \sup_{w \in \mathcal{W}} \sum_{i=1}^{N-1} w_i \|u_{it}\|_{\mathcal{H}_k} \\
&\leq \|u_{Nt}\|_{\mathcal{H}_k} + \max_{i < N} \|u_{it}\|_{\mathcal{H}_k} \\
&=: \delta_t
\end{aligned}$$

As the sum of two terms which have been shown to converge to 0, $\delta_t \xrightarrow{p} 0$. Moreover, the triangle

inequality provides:

$$\begin{aligned} \sup_{w \in \mathcal{W}} \|r_t(w)\|_{\mathcal{H}_k} &\leq \|\mu_{Nt}\|_{\mathcal{H}_k} + \sup_{w \in \mathcal{W}} \sum_{i=1}^{N-1} w_i \|\mu_{it}\|_{\mathcal{H}_k} \\ &\leq 2\sqrt{k} \end{aligned}$$

For any two elements a, b of a normed space,

$$|\|a\|^2 - \|b\|^2| \leq \|a - b\|(\|a\| + \|b\|)$$

Applying this inequality with $a = \hat{r}_t(w)$ and $b = r_t(w)$, and using

$$\|\hat{r}_t(w)\|_{\mathcal{H}_k} \leq \|r_t(w)\|_{\mathcal{H}_k} + \delta_t$$

gives

$$\sup_{w \in \mathcal{W}} |\|\hat{r}_t(w)\|_{\mathcal{H}_k}^2 - \|r_t(w)\|_{\mathcal{H}_k}^2| \leq \delta_t (4\sqrt{k} + \delta_t) \xrightarrow{p} 0$$

Finally, note that we can write:

$$\hat{L}(w) - L(w) = \sum_{t < T} \|\hat{r}_t(w)\|_{\mathcal{H}_k}^2 - \|r_t(w)\|_{\mathcal{H}_k}^2$$

Because the number of pre-treatment periods is fixed and finite, this yields:

$$\sup_{w \in \mathcal{W}} |\hat{L}(w) - L(w)| \leq \sum_{t < T} \sup_{w \in \mathcal{W}} |\|\hat{r}_t(w)\|_{\mathcal{H}_k}^2 - \|r_t(w)\|_{\mathcal{H}_k}^2| \xrightarrow{p} 0$$

□

Lemma 5 (Mixture stability of MMD_k^2). *Suppose **Assumption k-Reg** holds. Let $\{P_i\}_{i=1}^{N-1}, \{Q_i\}_{i=1}^{N-1}$ be sequences of distributions defined on $\mathcal{X}$. Then, for all $w \in \mathcal{W}$,*

$$\text{MMD}_k^2 \left(\sum_{i=1}^{N-1} w_i P_i, \sum_{i=1}^{N-1} w_i Q_i \right) \leq \sum_{i=1}^{N-1} w_i \text{MMD}_k^2(P_i, Q_i)$$

Proof of Lemma 5. Fix $w \in \mathcal{W}$ and probability distributions $P_1, \dots, P_{N-1}$ and $Q_1, \dots, Q_{N-1}$. Linearity of kernel mean embeddings yields:

$$\text{MMD}_k^2 \left(\sum_{i=1}^{N-1} w_i P_i, \sum_{i=1}^{N-1} w_i Q_i \right) = \left\| \sum_{i=1}^{N-1} w_i (\mu_{P_i} - \mu_{Q_i}) \right\|_{\mathcal{H}_k}^2$$

Because $w_i \geq 0$ and $\sum_{i=1}^{N-1} w_i = 1$, Jensen's inequality for the convex squared norm function yields:

$$\begin{aligned} \left\| \sum_{i=1}^{N-1} w_i (\mu_{P_i} - \mu_{Q_i}) \right\|_{\mathcal{H}_k}^2 &\leq \sum_{i=1}^{N-1} w_i \|\mu_{P_i} - \mu_{Q_i}\|_{\mathcal{H}_k}^2 \\ &= \sum_{i=1}^{N-1} w_i \text{MMD}_k^2(P_i, Q_i) \end{aligned}$$

□

Lemma 6 (Consistency of empirical distributions). *Suppose Assumptions **CS** and **k-Reg** hold. Then, for each donor $i = 1, \dots, N-1$, as $K_{\min}^{\text{post}} \rightarrow \infty$*

$$\text{MMD}_k \left(\widehat{F}_{iT}, F_{iT}(0) \right) \xrightarrow{p} 0,$$

Proof. Suppose **Assumption CS** holds. Fix a donor $i \in \{1, \dots, N-1\}$, and write

$$P = F_{iT}(0) \quad K = K_{iT} \quad X_r = x_{iT_r}$$

Because donor i is untreated in period T , **Assumption CS** implies

$$X_1, \dots, X_K \stackrel{\text{iid}}{\sim} P$$

Let $\phi(x) = k(x, \cdot) \in \mathcal{H}_k$ denote the feature map induced by kernel k and define

$$Z_r = \phi(X_r) - \mu_P$$

Then, $\mathbb{E}[Z_r] = 0$ and

$$\widehat{\mu}_P - \mu_P = \frac{1}{K} \sum_{r=1}^K Z_r$$

Expanding the squared Hilbert-space norm yields

$$\mathbb{E}[\|\widehat{\mu}_P - \mu_P\|_{\mathcal{H}_k}^2] = \frac{1}{K^2} \sum_{r=1}^K \sum_{s=1}^K \mathbb{E}[\langle Z_r, Z_s \rangle_{\mathcal{H}_k}]$$

If $r \neq s$, independence and centering imply

$$\begin{aligned}
\mathbb{E}[\langle Z_r, Z_s \rangle_{\mathcal{H}_k}] &= \mathbb{E}[\mathbb{E}[\langle Z_r, Z_s \rangle_{\mathcal{H}_k} \mid Z_r]] \\
&= \mathbb{E}[\langle Z_r, \mathbb{E}[Z_s \mid Z_r] \rangle_{\mathcal{H}_k}] \\
&= \mathbb{E}[\langle Z_r, \mathbb{E}[Z_s] \rangle_{\mathcal{H}_k}] \\
&= 0
\end{aligned}$$

Thus all cross terms vanish, leaving

$$\begin{aligned}
\mathbb{E}[\|\hat{\mu}_P - \mu_P\|_{\mathcal{H}_k}^2] &= \frac{1}{K^2} \sum_{r=1}^K \mathbb{E}[\|Z_r\|_{\mathcal{H}_k}^2] \\
&= \frac{1}{K} \mathbb{E}[\|\phi(X) - \mu_P\|_{\mathcal{H}_k}^2]
\end{aligned}$$

Moreover,

$$\begin{aligned}
\mathbb{E}[\|\phi(X) - \mu_P\|_{\mathcal{H}_k}^2] &= \mathbb{E}[\|\phi(X)\|_{\mathcal{H}_k}^2] - \|\mu_P\|_{\mathcal{H}_k}^2 \\
&= \mathbb{E}[k(X, X)] - \|\mu_P\|_{\mathcal{H}_k}^2 \\
&\leq \bar{k}
\end{aligned}$$

Where the last line uses boundedness of the kernel, by [Assumption k-Reg](#). It follows that

$$\mathbb{E}[\text{MMD}_k^2(\hat{F}_{iT}, F_{iT}(0))] \leq \frac{\bar{k}}{K_{iT}}$$

Hence, for every $\varepsilon > 0$, Markov's inequality gives

$$\begin{aligned}
\Pr\{\text{MMD}_k(\hat{F}_{iT}, F_{iT}(0)) > \varepsilon\} &= \Pr\{\text{MMD}_k^2(\hat{F}_{iT}, F_{iT}(0)) > \varepsilon^2\} \\
&\leq \frac{\bar{k}}{K_{iT}\varepsilon^2}
\end{aligned}$$

Since $K_{iT} \geq K_{\min}^{\text{post}} \rightarrow \infty$, the right-hand side converges to zero. Therefore,

$$\text{MMD}_k(\hat{F}_{iT}, F_{iT}(0)) \xrightarrow{P} 0$$

Since the argument applies to every donor i , this proves the desired result. $\square$

Lemma 7 (Triangle inequality). *Suppose [Assumption k-Reg](#) holds. Then, for any distributions P, Q, R over $\mathcal{X}$,*

$$\text{MMD}_k(P, R) \leq \text{MMD}_k(P, Q) + \text{MMD}_k(Q, R)$$

Proof. For any probability distributions P, Q, R ,

$$\begin{aligned}
\text{MMD}_k(P, R) &= \|\mu_P - \mu_R\|_{\mathcal{H}_k} \\
&= \|(\mu_P - \mu_Q) + (\mu_Q - \mu_R)\|_{\mathcal{H}_k} \\
&\leq \|\mu_P - \mu_Q\|_{\mathcal{H}_k} + \|\mu_Q - \mu_R\|_{\mathcal{H}_k} \\
&= \text{MMD}_k(P, Q) + \text{MMD}_k(Q, R),
\end{aligned}$$

where the inequality is the triangle inequality in $\mathcal{H}_k$. $\square$

Lemma 8 (Continuity of post-treatment synthetic distribution map). *Suppose [Assumption k-Reg](#) holds. Then, the map*

$$w \mapsto \sum_{i=1}^{N-1} w_i F_{iT}(0)$$

is continuous under metric MMD_k on $\mathcal{W}$. That is, if $w^{(n)} \rightarrow w$, then

$$\text{MMD}_k\left(\sum_{i=1}^{N-1} w_i^{(n)} F_{iT}(0), \sum_{i=1}^{N-1} w_i F_{iT}(0)\right) \rightarrow 0$$

Since $\mathcal{W}$ is compact, this further implies uniform continuity on $\mathcal{W}$.

Proof. Define the post-treatment synthetic-distribution map

$$S_T(w) = \sum_{i=1}^{N-1} w_i F_{iT}(0).$$

For any $w, w' \in \mathcal{W}$, linearity of the kernel mean embedding, boundedness via [Assumption k-Reg](#), and [Lemma 7](#) imply

$$\begin{aligned}
\text{MMD}_k(S_T(w'), S_T(w)) &= \left\| \sum_{i=1}^{N-1} (w'_i - w_i) \mu_{iT} \right\|_{\mathcal{H}_k} \\
&\leq \sum_{i=1}^{N-1} |w'_i - w_i| \|\mu_{iT}\|_{\mathcal{H}_k} \\
&\leq \sqrt{k} \sum_{i=1}^{N-1} |w'_i - w_i| \\
&= \sqrt{k} \|w' - w\|_1
\end{aligned}$$

Therefore, if $w' \rightarrow w$, then $\|w' - w\|_1 \rightarrow 0$ and

$$\text{MMD}_k(S_T(w'), S_T(w)) \rightarrow 0$$

Furthermore, the displayed bound shows that the map is globally Lipschitz under the ℓ_1 norm. Because $\mathcal{W}$ is compact, the map is also uniformly continuous on $\mathcal{W}$. $\square$

Lemma 9 (Donor mixture empirical convergence). *Suppose that Assumptions **CS** and **k-Reg** hold. Then, for any random $\hat{w} \in \mathcal{W}$, as $K_{\min}^{\text{post}} \rightarrow \infty$,*

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) \xrightarrow{p} 0.$$

Proof. For each donor $i = 1, \dots, N-1$, define

$$Z_i = \text{MMD}_k \left(\hat{F}_{iT}, F_{iT}(0) \right).$$

By **Lemma 6**, $Z_i \xrightarrow{p} 0$ for every donor $i = 1, \dots, N-1$.

Now consider the discrepancy between the empirical donor mixture and its population counterpart, evaluated at the random weight vector $\hat{w}$:

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right)$$

Applying **Lemma 5** gives

$$\begin{aligned} \text{MMD}_k^2 \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) &\leq \sum_{i=1}^{N-1} \hat{w}_i \text{MMD}_k^2 \left(\hat{F}_{iT}, F_{iT}(0) \right) \\ &= \sum_{i=1}^{N-1} \hat{w}_i Z_i^2 \end{aligned}$$

Since $\hat{w}_i \geq 0$ and $\sum_{i=1}^{N-1} \hat{w}_i = 1$,

$$\sum_{i=1}^{N-1} \hat{w}_i Z_i^2 \leq \max_{1 \leq i \leq N-1} Z_i^2$$

Hence

$$\text{MMD}_k^2 \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) \leq \max_{1 \leq i \leq N-1} Z_i^2$$

It remains to show that the right-hand side converges to zero in probability. Since N is fixed and $Z_i \xrightarrow{p} 0$ for every donor i , the continuous mapping theorem yields

$$\max_{1 \leq i \leq N-1} Z_i^2 \xrightarrow{p} 0$$

Therefore,

$$\text{MMD}_k^2 \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) \xrightarrow{p} 0$$

Applying the continuous mapping theorem once again and taking square roots yields

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) \xrightarrow{p} 0$$

This proves the claim. $\square$

Lemma 10 (Local smoothness of the population objective). *Suppose Assumptions **k-Reg** and **IDT** hold. Then, the population objective ℓ is twice continuously differentiable on $\mathbb{R}^{N-2}$ and $H = \nabla^2 \ell(\theta^0)$ is positive definite.*

Proof. Under the parameterization $w = w(\theta)$,

$$\sum_{i=1}^{N-1} w_i(\theta) \mu_{it} = \mu_{N-1,t} + D_t \theta$$

Consequently,

$$\ell(\theta) = \sum_{t < T} \|b_t - D_t \theta\|_{\mathcal{H}_k}^2$$

While ℓ is originally defined on Θ , we may extend it to all of $\mathbb{R}^{N-2}$ via the above equation. Note that when $\theta \in \Theta$, the definition of ℓ remains unchanged. This is a quadratic function of θ , with gradient

$$\nabla \ell(\theta) = 2 \sum_{t < T} D_t^* (D_t \theta - b_t)$$

and Hessian

$$\nabla^2 \ell(\theta) = 2 \sum_{t < T} D_t^* D_t =: H$$

Thus ℓ is twice continuously differentiable on all of $\mathbb{R}^{N-2}$.

It remains to establish that H is positive definite. For every $h \in \mathbb{R}^{N-2}$,

$$h^\top H h = 2 \sum_{t < T} \|D_t h\|_{\mathcal{H}_k}^2 \geq 0$$

Therefore, H is positive semi-definite. Towards a contradiction, suppose that this expression were exactly zero for some $h \neq 0$. Then, we would have $D_t h = 0$ for every $t < T$. Let $a = Ah \in \mathbb{R}^{N-1}$ denote the corresponding perturbation of the full weight vector. We would then have $\mathbf{1}^\top a = 0$ and,

for every $t < T$:

$$\begin{aligned}
\sum_{i=1}^{N-1} a_i \mu_{it} &= \sum_{j=1}^{N-2} h_j \mu_{jt} - \left(\sum_{j=1}^{N-2} h_j \right) \mu_{N-1,t} \\
&= \sum_{j=1}^{N-2} h_j (\mu_{jt} - \mu_{N-1,t}) \\
&= D_t h \\
&= 0
\end{aligned}$$

Because k is characteristic by **Assumption k-Reg**, its mean embedding is injective on finite signed measures of total mass zero. It therefore follows that, for every $t < T$:

$$\sum_{i=1}^{N-1} a_i F_{it}(0) = 0$$

Assumption IDT therefore implies $a = 0$. Since A has full column rank, this implies $h = 0$, a contradiction. Thus, H must be positive definite. $\square$

Lemma 11 (Interior score central limit theorem). *Suppose Assumptions **CS**, **k-Reg**, and **BSG** hold, and that $w_i^0 > 0$ for every $i = 1, \dots, N-1$. Then, the empirical objective $\hat{\ell}$ is differentiable on a neighborhood of θ^0 , and*

$$\sqrt{K_{\min}^{\text{pre}}} \left(\nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right) \xrightarrow{d} N(0, \Omega)$$

Additionally, because θ^0 is an interior minimizer, $\nabla \ell(\theta^0) = 0$.

Proof. For notational convenience, write $K = K_{\min}^{\text{pre}}$. Define

$$\Delta b_t = \hat{b}_t - b_t = u_{Nt} - u_{Jt}$$

and

$$\Delta D_t = \hat{D}_t - D_t$$

so that

$$\Delta D_t h = \sum_{j=1}^{N-2} h_j (u_{jt} - u_{Jt})$$

Here, $J = N-1$ denotes the index of the reference donor unit. Expanding the empirical score at

θ^0 gives

$$\nabla \widehat{\ell}(\theta^0) - \nabla \ell(\theta^0) = 2 \sum_{t < T} \left\{ (\Delta D_t)^* (D_t \theta^0 - b_t) + D_t^* (\Delta D_t \theta^0 - \Delta b_t) \right\} + \underbrace{2 \sum_{t < T} (\Delta D_t)^* (\Delta D_t \theta^0 - \Delta b_t)}_{R_K},$$

Under Assumptions **CS** and **BSG**,

$$\|u_{it}\|_{\mathcal{H}_k} = O_p(K_{it}^{-1/2}) = O_p(K^{-1/2})$$

for every pre-treatment cell. Since the number of cells is fixed,

$$\|\Delta b_t\|_{\mathcal{H}_k} = O_p(K^{-1/2}) \quad \text{and} \quad \|\Delta D_t\|_{\text{op}} = O_p(K^{-1/2})$$

Consequently,

$$\begin{aligned} \|R_K\|_2 &\leq 2 \sum_{t < T} \|\Delta D_t\|_{\text{op}} (\|\Delta D_t\|_{\text{op}} \|\theta^0\|_2 + \|\Delta b_t\|_{\mathcal{H}_k}) \\ &= O_p(K^{-1}) \end{aligned}$$

and hence

$$\sqrt{K} R_K = o_p(1)$$

It remains to characterize the linear part of the score expansion. For each pre-treatment cell (i, t) , define

$$\zeta_{it}(x) = \phi(x) - \mu_{it}$$

Note that $\zeta_{it}(x)$ is centered: $\mathbb{E}[\zeta_{it}(X)] = 0$. For $z \in \mathcal{H}_k$, define

$$\Delta b_t^{(i)}(z) = (\mathbf{1}\{i = N\} - \mathbf{1}\{i = J\}) z$$

and the linear operator $\Delta D_t^{(i)}(z) : \mathbb{R}^{N-2} \rightarrow \mathcal{H}_k$ by

$$\Delta D_t^{(i)}(z) h = \sum_{j=1}^{N-2} h_j (\mathbf{1}\{i = j\} - \mathbf{1}\{i = J\}) z$$

Note that by definition,

$$\Delta b_t = \sum_{i=1}^N \Delta b_t^{(i)}(u_{it}) \quad \Delta D_t = \sum_{i=1}^N \Delta D_t^{(i)}(u_{it})$$

Now, define the $\mathbb{R}^{N-2}$ -valued influence function

$$\psi_{it}(x) = 2 \left\{ (\Delta D_t^{(i)}(\zeta_{it}(x)))^* (D_t \theta^0 - b_t) + D_t^* (\Delta D_t^{(i)}(\zeta_{it}(x)) \theta^0 - \Delta b_t^{(i)}(\zeta_{it}(x))) \right\}$$

This influence function is linear in $\zeta_{it}(x)$. Therefore,

$$\mathbb{E}_{F_{it}(0)}[\psi_{it}(X)] = 0$$

Moreover, by **Assumption k-Reg**,

$$\|\zeta_{it}(x)\|_{\mathcal{H}_k} \leq \|\phi(x)\|_{\mathcal{H}_k} + \|\mu_{it}\|_{\mathcal{H}_k} \leq 2\sqrt{k}$$

so $\psi_{it}(X)$ is bounded and has a finite covariance matrix. Write

$$V_{it} = \text{Var}_{F_{it}(0)}(\psi_{it}(X))$$

We can now write the linear component of the score expansion as

$$\sum_{t < T} \sum_{i=1}^N \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr})$$

Hence

$$\sqrt{K} \left\{ \nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right\} = \sum_{t < T} \sum_{i=1}^N \sqrt{\frac{K}{K_{it}}} \left\{ \frac{1}{\sqrt{K_{it}}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr}) \right\} + o_p(1)$$

For every fixed cell (i, t) , the multivariate central limit theorem gives

$$\frac{1}{\sqrt{K_{it}}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr}) \xrightarrow{d} N(0, V_{it})$$

Assumption CS gives independence across cells, while **Assumption BSG** gives

$$\sqrt{\frac{K}{K_{it}}} \longrightarrow \kappa_{it}^{-1/2}$$

Since the number of pre-treatment cells is fixed, joint convergence and Slutsky's theorem imply

$$\sqrt{K} \left\{ \nabla \hat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right\} \xrightarrow{d} N(0, \Omega)$$

where

$$\Omega = \sum_{t < T} \sum_{i=1}^N \kappa_{it}^{-1} V_{it}$$

Finally, θ^0 is an interior population minimizer, so

$$\nabla \ell(\theta^0) = 0$$

The empirical objective is quadratic and therefore differentiable on an open neighborhood of θ^0 . $\square$

Lemma 12 (Interior Hessian convergence). *Suppose Assumptions [CS](#) and [k-Reg](#) hold. Then, as $K_{\min}^{\text{pre}} \rightarrow \infty$,*

$$\nabla^2 \widehat{\ell}(\theta) \xrightarrow{p} H$$

uniformly in $\theta \in \mathbb{R}^{N-2}$.

Proof. The empirical squared-MMD objective is

$$\widehat{\ell}(\theta) = \sum_{t < T} \|\widehat{b}_t - \widehat{D}_t \theta\|_{\mathcal{H}_k}^2$$

Hence it is quadratic on all of $\mathbb{R}^{N-2}$, and its Hessian is

$$\nabla^2 \widehat{\ell}(\theta) = 2 \sum_{t < T} \widehat{D}_t^* \widehat{D}_t$$

Note that the population analog of this expression provides for existence of H in the population. In particular, the empirical Hessian does not depend on θ .

The empirical-mean convergence established in the proof of [Lemma 4](#) gives

$$\widehat{d}_{jt} - d_{jt} = u_{jt} - u_{Jt} \xrightarrow{p} 0$$

in $\mathcal{H}_k$ for every j and $t < T$. For every $h \in \mathbb{R}^{N-2}$ with $\|h\|_2 \leq 1$,

$$\|(\widehat{D}_t - D_t)h\|_{\mathcal{H}_k} = \left\| \sum_{j=1}^{N-2} h_j (\widehat{d}_{jt} - d_{jt}) \right\|_{\mathcal{H}_k} \leq \left(\sum_{j=1}^{N-2} h_j^2 \right)^{1/2} \left(\sum_{j=1}^{N-2} \|\widehat{d}_{jt} - d_{jt}\|_{\mathcal{H}_k}^2 \right)^{1/2}$$

Therefore,

$$\|\widehat{D}_t - D_t\|_{\text{op}} \leq \left(\sum_{j=1}^{N-2} \|\widehat{d}_{jt} - d_{jt}\|_{\mathcal{H}_k}^2 \right)^{1/2} \xrightarrow{p} 0$$

It follows that

$$\begin{aligned} \|\widehat{\mathbf{D}}_t^* \widehat{\mathbf{D}}_t - \mathbf{D}_t^* \mathbf{D}_t\|_{\text{op}} &\leq \|\widehat{\mathbf{D}}_t - \mathbf{D}_t\|_{\text{op}} \left(\|\widehat{\mathbf{D}}_t\|_{\text{op}} + \|\mathbf{D}_t\|_{\text{op}} \right) \\ &\xrightarrow{p} 0 \end{aligned}$$

Since the number of pre-treatment periods is fixed,

$$\left\| \nabla^2 \widehat{\ell}(\theta) - H \right\|_{\text{op}} \xrightarrow{p} 0$$

uniformly in θ . In particular, for every possibly random sequence $\widetilde{\theta} \xrightarrow{p} \theta^0$,

$$\nabla^2 \widehat{\ell}(\widetilde{\theta}) \xrightarrow{p} H$$

This proves the result. □

Lemma 13 (Unique interior minimizer on the oracle face). *Suppose Assumptions **k-Reg** and **IDT** hold. Then*

$$w^0 \in \text{relint}(\mathcal{W}_{\mathcal{S}_0})$$

and

$$\arg \min_{w \in \mathcal{W}_{\mathcal{S}_0}} L(w) = \{w^0\}$$

Consequently, on the event $\{\widehat{\mathcal{S}} = \mathcal{S}_0\}$,

$$w^0 \in \text{relint}(\mathcal{W}_{\widehat{\mathcal{S}}})$$

and

$$\arg \min_{w \in \mathcal{W}_{\widehat{\mathcal{S}}}} L(w) = \{w^0\}$$

If $|\mathcal{S}_0| \geq 2$, then $\mathcal{W}_{\mathcal{S}_0}$ has dimension $|\mathcal{S}_0| - 1$, and the restricted problem has a unique interior minimizer. If $|\mathcal{S}_0| = 1$, then $\mathcal{W}_{\mathcal{S}_0} = \{w^0\}$.

Proof of Lemma 13. By definition of $\mathcal{S}_0$, $w_i^0 > 0$ for every $i \in \mathcal{S}_0$ and $w_i^0 = 0$ for every $i \notin \mathcal{S}_0$. Therefore, $w^0 \in \mathcal{W}_{\mathcal{S}_0}$ by definition. The relative interior of this face is

$$\text{relint}(\mathcal{W}_{\mathcal{S}_0}) = \{w \in \mathcal{W}_{\mathcal{S}_0} : w_i > 0 \text{ for every } i \in \mathcal{S}_0\}$$

It follows that $w^0 \in \text{relint}(\mathcal{W}_{\mathcal{S}_0})$.

We next establish uniqueness on the oracle face. Because $\mathcal{W}_{\mathcal{S}_0} \subseteq \mathcal{W}$ and w^0 is the unique

minimizer of L over $\mathcal{W}$ (by **Proposition 2**), every $w \in \mathcal{W}_{\mathcal{S}_0}$ satisfies

$$L(w) \geq L(w^0)$$

Moreover, $w^0 \in \mathcal{W}_{\mathcal{S}_0}$, so

$$\min_{w \in \mathcal{W}_{\mathcal{S}_0}} L(w) = L(w^0)$$

If $w \in \mathcal{W}_{\mathcal{S}_0}$ attains this restricted minimum, then

$$L(w) = L(w^0)$$

Such a w is also a population minimizer over $\mathcal{W}$. Since $\mathcal{W}_0 = \{w^0\}$, it follows that $w = w^0$. Hence

$$\arg \min_{w \in \mathcal{W}_{\mathcal{S}_0}} L(w) = \{w^0\}$$

On the event

$$\{\widehat{\mathcal{S}} = \mathcal{S}_0\}$$

we have

$$\mathcal{W}_{\widehat{\mathcal{S}}} = \mathcal{W}_{\mathcal{S}_0}$$

The preceding conclusions therefore imply

$$w^0 \in \text{relint}(\mathcal{W}_{\widehat{\mathcal{S}}})$$

and

$$\arg \min_{w \in \mathcal{W}_{\widehat{\mathcal{S}}}} L(w) = \{w^0\}$$

Finally, if $|\mathcal{S}_0| \geq 2$, then $\mathcal{W}_{\mathcal{S}_0}$ is a simplex of dimension $|\mathcal{S}_0| - 1$, and w^0 is its unique relative-interior minimizer. If $|\mathcal{S}_0| = 1$, the adding-up restriction uniquely determines the weight on the active donor, so

$$\mathcal{W}_{\mathcal{S}_0} = \{w^0\}$$

□

B.2 Proof of **Theorem 1**

Proof. We begin by proving consistency of the weights. First, note that the donor-weight set

$$\mathcal{W} = \left\{ w \in \mathbb{R}_+^{N-1} : \sum_{i=1}^{N-1} w_i = 1 \right\}$$

is a closed and bounded subset of finite-dimensional Euclidean space, and is therefore compact. By [Lemma 1](#), L is continuous on $\mathcal{W}$. Hence, by the extreme value theorem, L attains its minimum on $\mathcal{W}$, so

$$\mathcal{W}_0 = \arg \min_{w \in \mathcal{W}} L(w)$$

is nonempty. Moreover, since L is continuous, $\mathcal{W}_0$ is a closed subset of the compact set $\mathcal{W}$, and is therefore compact.

Let

$$L_0 = \min_{w \in \mathcal{W}} L(w)$$

Fix $\varepsilon > 0$, and define the set of feasible weights that are at least ε away from the population minimizer set:

$$\mathcal{W}_\varepsilon = \{w \in \mathcal{W} : d(w, \mathcal{W}_0) \geq \varepsilon\}$$

Since $d(\cdot, \mathcal{W}_0)$ is continuous and $\mathcal{W}$ is compact, $\mathcal{W}_\varepsilon$ is compact. If $\mathcal{W}_\varepsilon$ is empty, then $d(\hat{w}, \mathcal{W}_0) < \varepsilon$ automatically, so there is nothing to prove. Suppose therefore that $\mathcal{W}_\varepsilon$ is nonempty.

By continuity of L and compactness of $\mathcal{W}_\varepsilon$, L attains its minimum over $\mathcal{W}_\varepsilon$. Define

$$m_\varepsilon = \min_{w \in \mathcal{W}_\varepsilon} L(w)$$

We claim that $m_\varepsilon > L_0$. To see this, suppose instead that $m_\varepsilon = L_0$. Then there exists $w_\varepsilon \in \mathcal{W}_\varepsilon$ such that $L(w_\varepsilon) = L_0$, which implies $w_\varepsilon \in \mathcal{W}_0$. But then $d(w_\varepsilon, \mathcal{W}_0) = 0$, contradicting the definition of $\mathcal{W}_\varepsilon$, which requires $d(w_\varepsilon, \mathcal{W}_0) \geq \varepsilon > 0$. Therefore $m_\varepsilon > L_0$.

Now, let $\gamma_\varepsilon = m_\varepsilon - L_0 > 0$ and consider the event

$$\mathcal{E}_\varepsilon = \left\{ \sup_{w \in \mathcal{W}} \left| \hat{L}(w) - L(w) \right| \leq \frac{\gamma_\varepsilon}{3} \right\}$$

By [Lemma 4](#), $\Pr(\mathcal{E}_\varepsilon) \rightarrow 1$. We now show that, on $\mathcal{E}_\varepsilon$, $\hat{w} \notin \mathcal{W}_\varepsilon$. Suppose, toward a contradiction, that $\hat{w} \in \mathcal{W}_\varepsilon$. Then

$$L(\hat{w}) \geq m_\varepsilon = L_0 + \gamma_\varepsilon$$

Hence, on $\mathcal{E}_\varepsilon$,

$$\hat{L}(\hat{w}) \geq L(\hat{w}) - \frac{\gamma_\varepsilon}{3} \geq L_0 + \frac{2\gamma_\varepsilon}{3}$$

On the other hand, take any $w^0 \in \mathcal{W}_0$. Since $L(w^0) = L_0$, on $\mathcal{E}_\varepsilon$,

$$\hat{L}(w^0) \leq L(w^0) + \frac{\gamma_\varepsilon}{3} = L_0 + \frac{\gamma_\varepsilon}{3}.$$

Therefore, $\widehat{L}(\hat{w}) > \widehat{L}(w^0)$, which contradicts

$$\hat{w} \in \arg \min_{w \in \mathcal{W}} \widehat{L}(w)$$

Thus, on $\mathcal{E}_\varepsilon$, $d(\hat{w}, \mathcal{W}_0) < \varepsilon$. Consequently,

$$\begin{aligned} \Pr(d(\hat{w}, \mathcal{W}_0) \geq \varepsilon) &\leq 1 - \Pr(\mathcal{E}_\varepsilon) \\ &= \Pr\left(\sup_{w \in \mathcal{W}} \left| \widehat{L}(w) - L(w) \right| > \frac{\gamma_\varepsilon}{3}\right) \\ &\rightarrow 0 \end{aligned}$$

where the convergence follows from [Lemma 4](#). Since this holds for every $\varepsilon > 0$, $d(\hat{w}, \mathcal{W}_0) \xrightarrow{p} 0$. Finally, if $\mathcal{W}_0 = \{w^0\}$, then the point-to-set distance reduces to ordinary Euclidean distance:

$$d(\hat{w}, \mathcal{W}_0) = \inf_{w \in \{w^0\}} \|\hat{w} - w\| = \|\hat{w} - w^0\|$$

Hence,

$$d(\hat{w}, \mathcal{W}_0) \xrightarrow{p} 0 \implies \|\hat{w} - w^0\| \xrightarrow{p} 0$$

This is equivalent to $\hat{w} \xrightarrow{p} w^0$, proving the first result.

Now, we turn our attention to consistency of the estimated pseudo-counterfactual. Define the post-treatment synthetic distribution map

$$S(w) = \sum_{i=1}^{N-1} w_i F_{iT}(0)$$

By [Lemma 2](#), every population minimizer $w^0 \in \mathcal{W}_0$ induces the same post-treatment synthetic distribution:

$$S(w^0) = \sum_{i=1}^{N-1} w_i^0 F_{iT}(0) = \tilde{P}_0$$

We first show that

$$\text{MMD}_k(S(\hat{w}), \tilde{P}_0) \xrightarrow{p} 0$$

Fix $\varepsilon > 0$. By [Lemma 8](#), the map $S(\cdot)$ is uniformly continuous on the compact set $\mathcal{W}$ under MMD_k . Therefore, there exists $\delta > 0$ such that, for any $w, w' \in \mathcal{W}$,

$$\|w - w'\| < \delta \implies \text{MMD}_k(S(w), S(w')) < \varepsilon$$

On the event $\{d(\hat{w}, \mathcal{W}_0) < \delta\}$, there exists $w^0 \in \mathcal{W}_0$ such that $\|\hat{w} - w^0\| < \delta$. Hence, on this

event,

$$\text{MMD}_k \left(S(\hat{w}), \tilde{P}_0 \right) = \text{MMD}_k \left(S(\hat{w}), S(w^0) \right) < \varepsilon$$

where the equality uses [Lemma 2](#). Therefore,

$$\left\{ \text{MMD}_k \left(S(\hat{w}), \tilde{P}_0 \right) \geq \varepsilon \right\} \subseteq \{d(\hat{w}, \mathcal{W}_0) \geq \delta\}$$

Taking probabilities gives

$$\begin{aligned} \Pr \left[\text{MMD}_k \left(S(\hat{w}), \tilde{P}_0 \right) \geq \varepsilon \right] &\leq \Pr [d(\hat{w}, \mathcal{W}_0) \geq \delta] \\ &\rightarrow 0 \end{aligned}$$

The convergence comes from our previous conclusion that $d(\hat{w}, \mathcal{W}_0) \xrightarrow{p} 0$. Thus,

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0), \tilde{P}_0 \right) \xrightarrow{p} 0$$

Now decompose the discrepancy between the empirical synthetic distribution and the pseudo-counterfactual. By [Lemma 7](#),

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \tilde{P}_0 \right) \leq \text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0) \right) + \text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i F_{iT}(0), \tilde{P}_0 \right)$$

We have just shown that the second term on the right hand side converges to 0 in probability, and the first term does so as well by [Lemma 9](#). Therefore,

$$\text{MMD}_k \left(\sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}, \tilde{P}_0 \right) \xrightarrow{p} 0$$

This establishes the second claim and completes the proof. □

B.3 Proof of [Corollary 1](#)

Proof. Define the estimated post-treatment synthetic distribution:

$$\hat{P}_0 = \sum_{i=1}^{N-1} \hat{w}_i \hat{F}_{iT}$$

By [Theorem 1](#), $\text{MMD}_k(\hat{P}_0, \tilde{P}_0) \xrightarrow{p} 0$. By the triangle inequality established in [Lemma 7](#),

$$\begin{aligned} \text{MMD}_k(\hat{P}_0, P_0) &\leq \text{MMD}_k(\hat{P}_0, \tilde{P}_0) + \text{MMD}_k(\tilde{P}_0, P_0) \\ &= o_p(1) + \text{MMD}_k(\tilde{P}_0, P_0) \end{aligned}$$

This proves the misspecification result.

Now, suppose [Assumption CM](#) holds. Then, $\text{MMD}_k(\tilde{P}_0, P_0) = 0$ and, equivalently, $\tilde{P}_0 = P_0$. Direct application of [Theorem 1](#) then yields

$$\text{MMD}_k(\hat{P}_0, P_0) \xrightarrow{p} 0$$

This establishes the correct specification result. □

B.4 Proof of [Proposition 3](#)

Proof. By [Proposition 2](#) and the conditions of [Theorem 1](#), $\hat{w} \xrightarrow{p} w^0$ and, therefore, $\hat{\theta} \xrightarrow{p} \theta^0$. Because $\theta^0 \in \text{int}(\Theta)$, consistency implies that $\hat{\theta} \in \text{int}(\Theta)$ with probability approaching one. On this event, the first-order condition is

$$\nabla \hat{\ell}(\hat{\theta}) = 0$$

For each coordinate $j = 1, \dots, N-2$, the mean-value theorem gives a point $\tilde{\theta}_j$ on the line segment joining θ^0 and $\hat{\theta}$ such that the j th component of the sample score satisfies

$$\frac{\partial \hat{\ell}(\hat{\theta})}{\partial \theta_j} = \frac{\partial \hat{\ell}(\theta^0)}{\partial \theta_j} + \nabla \left(\frac{\partial \hat{\ell}(\tilde{\theta}_j)}{\partial \theta_j} \right)^\top (\hat{\theta} - \theta^0)$$

Stacking these equations yields

$$0 = \nabla \hat{\ell}(\theta^0) + \hat{H}(\hat{\theta} - \theta^0)$$

where the j th row of $\hat{H}$ is the j th row of $\nabla^2 \hat{\ell}(\tilde{\theta}_j)$.

Since $\hat{\theta} \xrightarrow{p} \theta^0$, each $\tilde{\theta}_j \xrightarrow{p} \theta^0$. [Lemma 12](#) therefore implies $\hat{H} \xrightarrow{p} H$. Because H is positive definite, it is nonsingular, and $\hat{H}^{-1} \xrightarrow{p} H^{-1}$. Rearranging the sample first-order condition gives

$$\sqrt{K_{\min}^{\text{pre}}}(\hat{\theta} - \theta^0) = -\hat{H}^{-1} \sqrt{K_{\min}^{\text{pre}}} \nabla \hat{\ell}(\theta^0).$$

[Lemma 11](#), together with $\nabla \ell(\theta^0) = 0$, and Slutsky's theorem yield

$$\sqrt{K_{\min}^{\text{pre}}}(\hat{\theta} - \theta^0) \xrightarrow{d} \mathcal{N}(0, H^{-1} \Omega H^{-1})$$

Finally, the weight map is affine, so $\hat{w} - w^0 = A(\hat{\theta} - \theta^0)$. Applying the continuous mapping theorem gives the stated limit for $\hat{w}$ and the covariance matrix in [Equation 2](#). Moreover, $\mathbf{1}_{N-1}^\top A = 0$, which

establishes the claimed singularity. □

B.5 Proof of **Proposition 4**

Proof. For each pre-treatment cell, define

$$V_{it} = \mathbb{E} \left[\psi_{it}(X) \psi_{it}(X)^\top \right]$$

Let its feasible counterpart be denoted as

$$\widehat{V}_{it} = \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \widehat{\psi}_{itr} \widehat{\psi}_{itr}^\top$$

The proof of **Lemma 11** establishes asymptotic linearity of the score:

$$\sqrt{K} \left\{ \nabla \widehat{\ell}(\theta^0) - \nabla \ell(\theta^0) \right\} = \sum_{t < T} \sum_{i=1}^N \sqrt{\frac{K}{K_{it}}} \left\{ \frac{1}{\sqrt{K_{it}}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr}) \right\} + o_p(1)$$

Consequently, the limiting score covariance is given by

$$\Omega = \sum_{t < T} \sum_{i=1}^N \kappa_{it}^{-1} V_{it}$$

We first show that $\widehat{V}_{it} \xrightarrow{p} V_{it}$ for every pre-treatment cell. By **Lemma 4**, the empirical kernel mean embeddings satisfy

$$\max_{i,t < T} \|\widehat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k} \xrightarrow{p} 0$$

The weight-consistency result in **Theorem 1**, together with **Proposition 2**, gives $\widehat{w} \xrightarrow{p} w^0$ and $\widehat{\theta} \xrightarrow{p} \theta^0$. From these consistency results, it follows that every empirical donor-contrast operator and synthetic-minus-treated residual entering the feasible influence functions converges to its population counterpart.

Because the kernel is bounded by **Assumption k-Reg**, we can bound the RKHS norm of the corresponding feature map $\phi(x)$:

$$\|\phi(x)\|_{\mathcal{H}_k} \leq \sqrt{\bar{k}}$$

This bound holds uniformly over $x \in \mathcal{X}$. The feasible influence functions are continuous functions of these uniformly bounded feature maps and of the preceding consistently estimated quantities. Therefore, for every cell (i, t) ,

$$\max_{1 \leq r \leq K_{it}} \left\| \widehat{\psi}_{itr} - \psi_{it}(x_{itr}) \right\|_2 \xrightarrow{p} 0$$

The feasible and population influence functions are also uniformly bounded in probability. First, note the following fact of operator norms:

$$\begin{aligned}
\|aa^\top - bb^\top\|_{\text{op}} &= \|(a-b)a^\top + b(a-b)^\top\|_{\text{op}} \\
&\leq \|(a-b)a^\top\|_{\text{op}} + \|b(a-b)^\top\|_{\text{op}} \\
&= \|(a-b)\|_2 \|a\|_2 + \|b\|_2 \|(a-b)\|_2 \\
&= (\|a\|_2 + \|b\|_2) \|(a-b)\|_2
\end{aligned}$$

Where we have used the triangle inequality. We can then write:

$$\begin{aligned}
\left\| \widehat{V}_{it} - \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr}) \psi_{it}(x_{itr})^\top \right\|_{\text{op}} &= \left\| \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \left(\widehat{\psi}_{it}(x_{itr}) \widehat{\psi}_{it}(x_{itr})^\top - \psi_{it}(x_{itr}) \psi_{it}(x_{itr})^\top \right) \right\|_{\text{op}} \\
&\leq \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \left\| \widehat{\psi}_{it}(x_{itr}) \widehat{\psi}_{it}(x_{itr})^\top - \psi_{it}(x_{itr}) \psi_{it}(x_{itr})^\top \right\|_{\text{op}} \\
&\leq \frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \left(\|\widehat{\psi}_{it}(x_{itr})\|_2 + \|\psi_{it}(x_{itr})\|_2 \right) \|\widehat{\psi}_{it}(x_{itr}) - \psi_{it}(x_{itr})\|_2 \\
&\leq \max_{1 \leq r \leq K_{it}} \left(\|\widehat{\psi}_{it}(x_{itr})\|_2 + \|\psi_{it}(x_{itr})\|_2 \right) \max_{1 \leq r \leq K_{it}} \|\widehat{\psi}_{it}(x_{itr}) - \psi_{it}(x_{itr})\|_2 \\
&= o_p(1)
\end{aligned}$$

Assumption CS and boundedness of the influence functions give the cellwise law of large numbers:

$$\frac{1}{K_{it}} \sum_{r=1}^{K_{it}} \psi_{it}(x_{itr}) \psi_{it}(x_{itr})^\top \xrightarrow{p} V_{it}$$

Hence, we have established that $\|\widehat{V}_{it} - V_{it}\|_{\text{op}} \xrightarrow{p} 0$, and so $\widehat{V}_{it} \xrightarrow{p} V_{it}$.

As defined in the proposition statement, we can write

$$\widehat{\Omega} = \sum_{t < T} \sum_{i=1}^N \frac{K}{K_{it}} \widehat{V}_{it}$$

Assumption BSG gives

$$\frac{K}{K_{it}} \longrightarrow \kappa_{it}^{-1}.$$

Since the number of pre-treatment cells is fixed and finite, the continuous mapping theorem yields

$$\widehat{\Omega} \xrightarrow{p} \sum_{t < T} \sum_{i=1}^N \kappa_{it}^{-1} V_{it} = \Omega$$

Lemma 12 directly gives $\hat{H} \xrightarrow{p} H$, while **Lemma 10** establishes that H is positive definite. Consequently, $\hat{H}$ is nonsingular with probability approaching one, and so $\hat{H}^{-1} \xrightarrow{p} H^{-1}$. With all of these pieces in hand, we can then apply the continuous mapping theorem to the stated sandwich formula:

$$\hat{\Sigma}_w = A\hat{H}^{-1}\hat{\Omega}\hat{H}^{-1}A^\top \xrightarrow{p} AH^{-1}\Omega H^{-1}A^\top = \Sigma_w$$

□

B.6 Proof of **Proposition 5**

Proof. Write $K = K_{\min}^{\text{pre}}$. Under the θ -parameterization, let $\tilde{w} = w(\tilde{\theta})$ and $w^0 = w(\theta^0)$, where

$$\tilde{\theta} \in \arg \min_{\theta \in \Theta} \hat{\ell}(\theta) \quad \theta^0 \in \arg \min_{\theta \in \Theta} \ell(\theta)$$

For squared MMD, the population and empirical objectives are:

$$\hat{\ell}(\theta) = \sum_{t < T} \left\| \hat{b}_t - \hat{D}_t \theta \right\|_{\mathcal{H}_k}^2 \quad \ell(\theta) = \sum_{t < T} \|b_t - D_t \theta\|_{\mathcal{H}_k}^2$$

We first establish the required rates for the empirical kernel means. The proof of **Lemma 4** gives

$$\mathbb{E} \left[\|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k}^2 \right] \leq \frac{\bar{k}}{K_{it}}$$

We can apply Markov's inequality to yield:

$$\begin{aligned} \Pr \left(\|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k} > \frac{M}{\sqrt{K}} \right) &= \Pr \left(\|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k}^2 > \frac{M^2}{K} \right) \\ &\leq \frac{\mathbb{E} \left[\|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k}^2 \right]}{M^2/K} \\ &\leq \frac{\bar{k}}{M^2} \end{aligned}$$

Consequently, $\|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k} = O_p(K^{-1/2})$. Since the number of pre-treatment periods is fixed, we can use this to obtain the union bound:

$$\max_{\substack{i=1,\dots,N \\ t < T}} \|\hat{\mu}_{it} - \mu_{it}\|_{\mathcal{H}_k} = O_p(K^{-1/2})$$

Since $\widehat{b}_t - b_t$ and $\widehat{D}_t - D_t$ are finite linear combinations/operators of these empirical-mean errors, we can write for every $t < T$:

$$\left\| \widehat{b}_t - b_t \right\|_{\mathcal{H}_k} = O_p(K^{-1/2}) \quad \left\| \widehat{D}_t - D_t \right\|_{\text{op}} = O_p(K^{-1/2})$$

The population and empirical Hessians are given by:

$$H = 2 \sum_{t < T} D_t^* D_t \quad \widehat{H} = 2 \sum_{t < T} \widehat{D}_t^* \widehat{D}_t$$

Define $\Delta D_t = \widehat{D}_t - D_t$. We can then write:

$$\begin{aligned} \left\| \widehat{H} - H \right\|_{\text{op}} &= \left\| 2 \sum_{t < T} \widehat{D}_t^* \widehat{D}_t - D_t^* D_t \right\|_{\text{op}} \\ &\leq 2 \sum_{t < T} \left\| \widehat{D}_t^* \widehat{D}_t - D_t^* D_t \right\|_{\text{op}} \\ &= 2 \sum_{t < T} \left\| D_t^* \Delta D_t + (\Delta D_t)^* D_t + (\Delta D_t)^* \Delta D_t \right\|_{\text{op}} \\ &\leq 2 \sum_{t < T} (2 \|D_t\|_{\text{op}} \|\Delta D_t\|_{\text{op}} + \|\Delta D_t\|_{\text{op}}^2) \\ &= 2 \sum_{t < T} O_p(K^{-1/2}) + O_p(K^{-1}) \\ &= O_p(K^{-1/2}) \\ &= o_p(1) \end{aligned}$$

Moreover, [Assumption IDT](#), together with characteristicness of the kernel k from [Assumption k-Reg](#), gives for every $h \neq 0$,

$$h^\top H h = 2 \sum_{t < T} \|D_t h\|_{\mathcal{H}_k}^2 > 0$$

Thus H is positive definite. Let

$$c = \lambda_{\min}(H) > 0$$

where $\lambda_{\min}(H)$ denotes the smallest eigenvalue of H . Hessian convergence implies

$$\Pr \left(\lambda_{\min}(\widehat{H}) \geq \frac{c}{2} \right) \rightarrow 1$$

Next, define the empirical-score error at the population minimizer:

$$g_K = \nabla \widehat{\ell}(\theta^0) - \nabla \ell(\theta^0)$$

and define

$$\Delta b_t = \widehat{b}_t - b_t$$

Expanding the empirical score gives

$$g_K = 2 \sum_{t < T} \left\{ (\Delta D_t)^* (D_t \theta^0 - b_t) + D_t^* (\Delta D_t \theta^0 - \Delta b_t) + (\Delta D_t)^* (\Delta D_t \theta^0 - \Delta b_t) \right\}$$

The first two terms in each summand are $O_p(K^{-1/2})$, while the final term is $O_p(K^{-1})$. Hence, $g_K = O_p(K^{-1/2})$.

It remains to translate this score rate into a rate for the constrained minimizer. Define $\delta_K = \widetilde{\theta} - \theta^0$. Because θ^0 minimizes the differentiable convex function ℓ over the convex set Θ , we have

$$\delta_K^\top \nabla \ell(\theta^0) \geq 0$$

Analogously, because $\widetilde{\theta}$ minimizes $\widehat{\ell}$ over Θ ,

$$\delta_K^\top \nabla \widehat{\ell}(\widetilde{\theta}) \leq 0$$

The empirical objective is quadratic, so

$$\nabla \widehat{\ell}(\widetilde{\theta}) = \nabla \widehat{\ell}(\theta^0) + \widehat{H} \delta_K$$

Consequently,

$$\begin{aligned} \delta_K^\top \widehat{H} \delta_K &\leq -\delta_K^\top \nabla \widehat{\ell}(\theta^0) \\ &= -\delta_K^\top g_K - \delta_K^\top \nabla \ell(\theta^0) \\ &\leq \|\delta_K\|_2 \|g_K\|_2 \end{aligned}$$

Where the final inequality follows from Cauchy-Schwarz. On the event

$$\left\{ \lambda_{\min}(\widehat{H}) \geq c/2 \right\}$$

we therefore have

$$\frac{c}{2} \|\delta_K\|_2^2 \leq \|\delta_K\|_2 \|g_K\|_2$$

If $\delta_K \neq 0$, we can then write:

$$\|\widetilde{\theta} - \theta^0\|_2 \leq \frac{2}{c} \|g_K\|_2 = O_p(K^{-1/2})$$

Note that the same conclusion is immediate if $\delta_K = 0$. Thus,

$$\sqrt{K}(\tilde{\theta} - \theta^0) = O_p(1)$$

Because $w(\theta)$ is affine with fixed Jacobian A ,

$$\tilde{w} - w^0 = A(\tilde{\theta} - \theta^0)$$

and therefore

$$\sqrt{K}(\tilde{w} - w^0) = O_p(1)$$

Finally, the stated assumption on the cutoff schedule λ_K implies

$$\frac{\|\tilde{w} - w^0\|_\infty}{\lambda_K} = O_p\left(\frac{1}{\sqrt{K} \lambda_K}\right) = o_p(1)$$

Thus:

$$\Pr(\|\tilde{w} - w^0\|_\infty < \lambda_K) \rightarrow 1$$

Because N is fixed and S_0 is nonempty, define

$$\underline{w}^0 = \min_{i \in S_0} w_i^0 > 0.$$

Since $\lambda_K \rightarrow 0$, we have $\lambda_K < \underline{w}^0/2$ for all sufficiently large K . On the event

$$\{\|\tilde{w} - w^0\|_\infty < \lambda_K\}$$

if $i \notin S_0$, then $w_i^0 = 0$ and $0 \leq \tilde{w}_i < \lambda_K$, and so $i \notin \hat{S}$. If $i \in S_0$, then

$$\tilde{w}_i > w_i^0 - \lambda_K \geq \underline{w}^0 - \lambda_K > \lambda_K$$

and so $i \in \hat{S}$. Hence $\hat{S} = S_0$ on an event whose probability converges to one, establishing

$$\Pr(\hat{S} = S_0) \rightarrow 1$$

□

B.7 Proof of **Theorem 2**

Proof. Write $K = K_{\min}^{\text{pre}}$ and define the support-recovery event $\mathcal{E}_K = \{\hat{S} = S_0\}$. By **Proposition 5**, $\Pr(\mathcal{E}_K) \rightarrow 1$. Define the oracle support-refitted estimator by

$$\hat{w}^{\text{or}} \in \arg \min_{w \in \mathcal{W}_{S_0}} \hat{L}(w)$$

On $\mathcal{E}_K$, note that $\mathcal{W}_{\hat{S}} = \mathcal{W}_{S_0}$. The feasible and oracle estimators therefore minimize the same objective over the same face. Using the same argmin selection rule, $\hat{w}^{\text{ts}} = \hat{w}^{\text{or}}$ on $\mathcal{E}_K$.

We first establish consistency. The face $\mathcal{W}_{S_0}$ can be viewed as the unrestricted simplex for the synthetic control problem obtained by restricting the donor set to S_0 . After relabeling these donors, Assumptions **CS** and **k-Reg** continue to hold. Applying **Lemma 4** to this restricted donor problem therefore gives

$$\sup_{w \in \mathcal{W}_{S_0}} \left| \hat{L}(w) - L(w) \right| \xrightarrow{P} 0$$

By **Lemma 13**, w^0 is the unique population minimizer of L over the compact face $\mathcal{W}_{S_0}$. The argmin consistency theorem consequently gives $\hat{w}^{\text{or}} \xrightarrow{P} w^0$. For every $\varepsilon > 0$,

$$\begin{aligned} \Pr(\|\hat{w}^{\text{ts}} - w^0\|_2 > \varepsilon) &\leq \Pr(\mathcal{E}_K^c) + \Pr(\|\hat{w}^{\text{or}} - w^0\|_2 > \varepsilon) \\ &\rightarrow 0 \end{aligned}$$

Thus, $\hat{w}^{\text{ts}} \xrightarrow{P} w^0$.

Now suppose that $s_0 \geq 2$ and **Assumption BSG** also holds. Parameterize the oracle face by $s_0 - 1$ free coordinates:

$$w = w_{S_0}(\theta) \quad \theta \in \Theta_{S_0}$$

Let θ^0 denote the coordinates corresponding to w^0 : $w^0 = w_{S_0}(\theta^0)$. Let A_{S_0} denote the Jacobian of $w_{S_0}(\theta)$. By **Lemma 13**, $\theta^0 \in \text{int}(\Theta_{S_0})$ and θ^0 is the unique minimizer of the population objective on this parameter space. The population MMD objective restricted to this face of the full donor simplex can be written as

$$\ell_{S_0}(\theta) = \sum_{t < T} \|b_{t,S_0} - D_{t,S_0}\theta\|_{\mathcal{H}_k}^2$$

with Hessian

$$H_{S_0} = 2 \sum_{t < T} D_{t,S_0}^* D_{t,S_0}$$

We will first show that H_{S_0} is positive definite. For every $h \in \mathbb{R}^{s_0-1}$,

$$h^\top H_{S_0} h = 2 \sum_{t < T} \|D_{t,S_0} h\|_{\mathcal{H}_k}^2 \geq 0$$

Towards a contradiction, suppose that this expression were exactly zero for some $h \neq 0$. Then, we would have $D_{t,S_0}h = 0$ for every $t < T$. Let $a = A_{S_0}h \in \mathbb{R}^{N-1}$ denote the corresponding perturbation of the full weight vector, including zeros outside S_0 , and suppose the elements of S_0 are indexed $S_0 = \{i_1, \dots, i_{s_0}\}$. We would then have $\mathbf{1}^\top a = 0$ and, for every $t < T$:

$$\begin{aligned} \sum_{i=1}^{N-1} a_i \mu_{it} &= \sum_{j=1}^{s_0-1} h_j \mu_{i_j t} - \left(\sum_{j=1}^{s_0-1} h_j \right) \mu_{i_{s_0} t} \\ &= \sum_{j=1}^{s_0-1} h_j (\mu_{i_j t} - \mu_{i_{s_0} t}) \\ &= D_{t,S_0}h \\ &= 0 \end{aligned}$$

Because k is characteristic by **Assumption k-Reg**, its mean embedding is injective on finite signed measures of total mass zero. It therefore follows that, for every $t < T$:

$$\sum_{i=1}^{N-1} a_i F_{it}(0) = 0$$

Assumption IDT therefore implies $a = 0$. Since A_{S_0} has full column rank, this implies $h = 0$, a contradiction. Thus, H_{S_0} must be positive definite.

Now, let

$$\hat{\ell}_{S_0}(\theta) = \hat{L}(w_{S_0}(\theta))$$

denote the empirical objective restricted to the oracle face, and write

$$\hat{H}_{S_0} = \nabla^2 \hat{\ell}_{S_0}(\theta)$$

Because the empirical MMD objective is quadratic, $\hat{H}_{S_0}$ does not depend on θ .

The restricted problem satisfies the conditions of **Lemma 12**: the Assumptions **CS** and **k-Reg** are inherited from the full problem, while **Lemma 13** supplies the unique interior population minimizer. Relabeling the donors in S_0 therefore gives $\hat{H}_{S_0} \xrightarrow{p} H_{S_0}$. Since H_{S_0} is positive definite, $\hat{H}_{S_0}$ is nonsingular with probability approaching one and $\hat{H}_{S_0}^{-1} \xrightarrow{p} H_{S_0}^{-1}$. The same argument applies to **Lemma 11**, where **Assumption BSG** for the full collection of pre-treatment cells implies balanced growth for the cells indexed by $S_0 \cup \{N\}$. We therefore have:

$$\sqrt{K} \left(\nabla \hat{\ell}_{S_0}(\theta^0) - \nabla \ell_{S_0}(\theta^0) \right) \xrightarrow{d} \mathcal{N}(0, \Omega_{S_0})$$

Because θ^0 is an interior population minimizer, $\nabla \ell_{S_0}(\theta^0) = 0$. Consequently,

$$\sqrt{K} \nabla \widehat{\ell}_{S_0}(\theta^0) \xrightarrow{d} \mathcal{N}(0, \Omega_{S_0})$$

Now, let $\widehat{\theta}^{\text{or}}$ denote the free coordinates of $\widehat{w}^{\text{or}}$. Oracle consistency and interiority of θ^0 imply

$$\Pr(\widehat{\theta}^{\text{or}} \in \text{int}(\Theta_{S_0})) \rightarrow 1$$

On this event, the empirical first-order condition gives $\nabla \widehat{\ell}_{S_0}(\widehat{\theta}^{\text{or}}) = 0$. Since $\widehat{\ell}_{S_0}$ is quadratic,

$$0 = \nabla \widehat{\ell}_{S_0}(\theta^0) + \widehat{H}_{S_0}(\widehat{\theta}^{\text{or}} - \theta^0)$$

Rearranging terms and left-multiplying by the inverse Hessian then yields:

$$\sqrt{K}(\widehat{\theta}^{\text{or}} - \theta^0) = -\widehat{H}_{S_0}^{-1} \sqrt{K} \nabla \widehat{\ell}_{S_0}(\theta^0)$$

Hessian convergence, the score central limit theorem, and Slutsky's theorem now give

$$\sqrt{K}(\widehat{\theta}^{\text{or}} - \theta^0) \xrightarrow{d} \mathcal{N}\left(0, H_{S_0}^{-1} \Omega_{S_0} H_{S_0}^{-1}\right).$$

Because $w_{S_0}(\theta)$ is affine,

$$\widehat{w}^{\text{or}} - w^0 = A_{S_0}(\widehat{\theta}^{\text{or}} - \theta^0)$$

Therefore, the continuous mapping theorem yields:

$$\sqrt{K}(\widehat{w}^{\text{or}} - w^0) \xrightarrow{d} \mathcal{N}\left(0, \underbrace{A_{S_0} H_{S_0}^{-1} \Omega_{S_0} H_{S_0}^{-1} A_{S_0}^{\top}}_{\Sigma_{w, S_0}}\right)$$

Finally, oracle equivalence gives

$$\Pr(\widehat{w}^{\text{ts}} \neq \widehat{w}^{\text{or}}) \leq \Pr(\mathcal{E}_K^c) \rightarrow 0$$

Thus, for every $\varepsilon > 0$,

$$\Pr\left(\sqrt{K} \|\widehat{w}^{\text{ts}} - \widehat{w}^{\text{or}}\|_2 > \varepsilon\right) \leq \Pr(\mathcal{E}_K^c) \rightarrow 0$$

and so

$$\sqrt{K}(\widehat{w}^{\text{ts}} - \widehat{w}^{\text{or}}) \xrightarrow{p} 0$$

Combining this result with the oracle limit and applying Slutsky's theorem yields

$$\sqrt{K_{\min}^{\text{pre}}}(\widehat{w}^{\text{ts}} - w^0) \xrightarrow{d} \mathcal{N}(0, \Sigma_{w, S_0})$$

□

B.8 Proof of Corollary 2

Proof. Define the support-recovery event $\mathcal{E}_K = \{\hat{S} = S_0\}$. By Proposition 5, $\Pr(\mathcal{E}_K) \rightarrow 1$. Consider first the fixed oracle-face problem with donor set S_0 . By Lemma 13, w^0 is the unique relative-interior population minimizer on $\mathcal{W}_{S_0}$. Moreover, the proof of Theorem 2 shows that the sampling, kernel, balanced-growth, score-CLT, and Hessian-convergence conditions required for the interior-weight problem hold after restricting the donor set to S_0 . Therefore, applying Proposition 4 to the fixed oracle-face problem gives

$$\hat{H}_{S_0} \xrightarrow{p} H_{S_0} \quad \hat{\Omega}_{S_0} \xrightarrow{p} \Omega_{S_0}$$

Hence, the continuous mapping theorem yields:

$$\hat{\Sigma}_{w,S_0} := A_{S_0} \hat{H}_{S_0}^{-1} \hat{\Omega}_{S_0} \hat{H}_{S_0}^{-1} A_{S_0}^\top \xrightarrow{p} A_{S_0} H_{S_0}^{-1} \Omega_{S_0} H_{S_0}^{-1} A_{S_0}^\top = \Sigma_{w,S_0}$$

On $\mathcal{E}_K$, the selected and oracle faces coincide. Using the same ordering and reference-donor rule in both parameterizations, we have:

$$A_{\hat{S}} = A_{S_0} \quad \hat{H}_{\hat{S}} = \hat{H}_{S_0} \quad \hat{\Omega}_{\hat{S}} = \hat{\Omega}_{S_0}$$

Consequently, $\hat{\Sigma}_{w,\hat{S}} = \hat{\Sigma}_{w,S_0}$. For every $\varepsilon > 0$,

$$\begin{aligned} \Pr\left(\left\|\hat{\Sigma}_{w,\hat{S}} - \Sigma_{w,S_0}\right\|_{\text{op}} > \varepsilon\right) &\leq \Pr(\mathcal{E}_K^c) + \Pr\left(\left\|\hat{\Sigma}_{w,S_0} - \Sigma_{w,S_0}\right\|_{\text{op}} > \varepsilon\right) \\ &\rightarrow 0 \end{aligned}$$

This is equivalent to $\hat{\Sigma}_{w,\hat{S}} \xrightarrow{p} \Sigma_{w,S_0}$, establishing the desired result. □

C Proofs for Section 4

C.1 Useful Lemmas

Lemma 14 (Fixed-summary error decomposition). *Let $\hat{\tau}_g, \tilde{\tau}_g$ be defined as in the preamble to Section 4, with $g : \mathcal{X} \rightarrow \mathbb{R}$ a fixed summary. The estimation error admits the exact decomposition*

$$\hat{\tau}_g - \tilde{\tau}_g = A_{1K}(g) - A_{0K}(g) - b_{K,g}$$

where

$$A_{1K}(g) = \mathbb{E}_{\hat{P}_1}[g(X)] - \mathbb{E}_{P_1}[g(X)]$$

and

$$A_{0K}(g) = \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \left\{ \mathbb{E}_{\hat{F}_{iT}}[g(X)] - \mathbb{E}_{F_{iT}(0)}[g(X)] \right\}$$

and

$$b_{K,g} = \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mu_i(g) - \tilde{\mu}_0(g)$$

Here, $\tilde{\mu}_0(g) := \mathbb{E}_{\tilde{P}_0}[g(x)]$.

Proof of Lemma 14. By definition,

$$\hat{\tau}_g = \mathbb{E}_{\hat{P}_1}[g(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{\hat{F}_{iT}}[g(X)] \quad \quad \tilde{\tau}_g = \mathbb{E}_{P_1}[g(X)] - \tilde{\mu}_0(g)$$

Therefore,

$$\hat{\tau}_g - \tilde{\tau}_g = \mathbb{E}_{\hat{P}_1}[g(X)] - \mathbb{E}_{P_1}[g(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{\hat{F}_{iT}}[g(X)] + \tilde{\mu}_0(g)$$

Note that by definition, we have:

$$\sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{F_{iT}(0)}[g(X)] = \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mu_i(g)$$

Adding and subtracting this term on the right hand side yields:

$$\hat{\tau}_g - \tilde{\tau}_g = \underbrace{\left\{ \mathbb{E}_{\hat{P}_1}[g(X)] - \mathbb{E}_{P_1}[g(X)] \right\}}_{A_{1K}(g)} - \underbrace{\sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \left\{ \mathbb{E}_{\hat{F}_{iT}}[g(X)] - \mathbb{E}_{F_{iT}(0)}[g(X)] \right\}}_{A_{0K}(g)} - \underbrace{\left\{ \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mu_i(g) - \tilde{\mu}_0(g) \right\}}_{b_{K,g}}$$

□

Lemma 15 (Evaluation-sample mean convergence for fixed summary). *Suppose Assumptions CS and BFS hold. Let $A_{1K}(g), A_{0K}(g)$ be defined as in Lemma 14. Then as $K_{\min} \rightarrow \infty$, $A_{1K}(g) \xrightarrow{P} 0$ and, for any random $\hat{w} \in \mathcal{W}$, $A_{0K}(g) \xrightarrow{P} 0$.*

Proof of Lemma 15. By Assumption CS:

$$\mathbb{E}_{\hat{P}_1}[g(X)] = \frac{1}{K_{NT}} \sum_{k=1}^{K_{NT}} g(x_{NTk})$$

This is a sample mean of iid draws from P_1 . By Assumption BFS, $\mathbb{E}_{P_1}[|g(X)|] < \infty$. Since

$K_{NT} \geq K_{\min} \rightarrow \infty$, the weak law of large numbers gives:

$$A_{1K}(g) = \mathbb{E}_{\hat{P}_1}[g(X)] - \mathbb{E}_{P_1}[g(X)] \xrightarrow{p} 0$$

Now, for each donor $i = 1, \dots, N-1$, define:

$$B_{iK}(g) = \mathbb{E}_{\hat{F}_{iT}}[g(X)] - \mathbb{E}_{F_{iT}(0)}[g(X)]$$

The donor evaluation observations are iid draws from $F_{iT}(0)$, and $\mathbb{E}_{F_{iT}(0)}[|g(X)|] < \infty$ by **Assumption BFS**. Since $K_{iT} \geq K_{\min} \rightarrow \infty$, the weak law of large numbers gives:

$$B_{iK}(g) \xrightarrow{p} 0$$

Because N is fixed, these convergences hold jointly:

$$\max_{1 \leq i \leq N-1} |B_{iK}(g)| \xrightarrow{p} 0$$

Since $\hat{w} \in \mathcal{W}$, $\hat{w}_i \geq 0$ and $\sum_{i=1}^{N-1} \hat{w}_i = 1$. Consequently, applying the triangle inequality yields:

$$\begin{aligned} |A_{0K}(g)| &= \left| \sum_{i=1}^{N-1} \hat{w}_i B_{iK}(g) \right| \\ &\leq \sum_{i=1}^{N-1} \hat{w}_i |B_{iK}(g)| \\ &\leq \max_{1 \leq i \leq N-1} |B_{iK}(g)| \\ &\xrightarrow{p} 0 \end{aligned}$$

□

Lemma 16 (Foldwise two-stage weight limits). *Suppose **Assumption PrTS** holds. Suppose that, for each $a \in \{\text{disc}, \text{eval}\}$, the conditions of **Theorem 2** hold after replacing*

$$\hat{L} \quad \tilde{w} \quad \hat{\mathcal{S}} \quad \hat{w}^{\text{ts}} \quad \hat{w}^{\text{or}} \quad K_{\min}^{\text{pre}}$$

by their fold-specific counterparts

$$\hat{L}^a \quad \tilde{w}^a \quad \hat{\mathcal{S}}^a \quad \hat{w}^{\text{ts},a} \quad \hat{w}^{\text{or},a} \quad K_{\min}^{\text{pre},a}$$

respectively.

Then, for each $a \in \{\text{disc}, \text{eval}\}$,

$$\Pr \left(\widehat{S}^a = S_0 \right) \rightarrow 1$$

and

$$\Pr \left(\widehat{w}^{\text{ts},a} = \widehat{w}^{\text{or},a} \right) \rightarrow 1$$

Consequently,

$$\widetilde{w}^a \xrightarrow{p} w^0 \quad \widehat{w}^{\text{ts},a} \xrightarrow{p} w^0$$

and

$$\sqrt{K_{\min}^{\text{pre},a}} \left(\widehat{w}^{\text{ts},a} - \widehat{w}^{\text{or},a} \right) \xrightarrow{p} 0$$

If $s_0 \geq 2$, then

$$\sqrt{K_{\min}^{\text{pre},a}} \left(\widehat{w}^{\text{ts},a} - w^0 \right) \xrightarrow{d} N(0, \Sigma_{w,S_0})$$

If $s_0 = 1$, then

$$\Pr \left(\widehat{w}^{\text{ts},a} = w^0 \right) \rightarrow 1$$

Proof of Lemma 16. Fix $a \in \{\text{disc}, \text{eval}\}$. Conditional on the split indices, the observations in each fold-specific pre-treatment cell are iid draws from the same population distribution as the corresponding unsplit observations. Moreover, Assumptions **BSG** and **PrTS** imply

$$\frac{K_{\min}^{\text{pre},a}}{K_{\min}^{\text{pre}}} \rightarrow \pi_{\text{pre}}^a$$

and, for every $i = 1, \dots, N$ and $t < T$,

$$\frac{K_{it}^a}{K_{\min}^{\text{pre},a}} \rightarrow \kappa_{it}^{\text{pre}}$$

Thus the fold-specific problem has the same population objective, population minimizer, and limiting cell-growth ratios as the original pre-treatment problem.

The conclusions therefore follow by applying the support-recovery, oracle-equivalence, consistency, and asymptotic-normality arguments of Theorem 2 to $\widehat{L}^a$, with $K_{\min}^{\text{pre},a}$ as the asymptotic scale. Because there are two folds, the conclusions hold for both simultaneously. $\square$

Lemma 17 (Conditional evaluation-weight asymptotic normality). *Suppose the conditions of Lemma 16 hold, and that Assumption PrSG holds as well.*

If $s_0 \geq 2$, then, conditional on $\mathcal{D}_K$,

$$\sqrt{K_{\min}^{\text{pre},\text{eval}}} \left(\widehat{w}^{\text{ts},\text{eval}} - w^0 \right) \xrightarrow{d} N(0, \Sigma_{w,S_0})$$

in probability with respect to the discovery sample. That is, for every bounded continuous function

$$f : \mathbb{R}^{N-1} \rightarrow \mathbb{R},$$

$$\mathbb{E} \left[f \left(\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \right) \middle| \mathcal{D}_K \right] \xrightarrow{p} \mathbb{E} [f(Z_w)] \quad Z_w \sim N(0, \Sigma_{w, S_0})$$

If $s_0 = 1$, then

$$\Pr \left(\hat{w}^{\text{ts,eval}} = w^0 \middle| \mathcal{D}_K \right) \xrightarrow{p} 1$$

Proof of Lemma 17. Conditional on $\mathcal{S}_K$, the pre-treatment evaluation fold is independent of

$$\sigma \left(D_{\text{pre}}^{\text{disc}}, D_{\text{post}}^{\text{disc}}, \mathcal{R}_K \right)$$

Consequently, for every bounded continuous function $f : \mathbb{R}^{N-1} \rightarrow \mathbb{R}$,

$$\mathbb{E} \left[f \left(\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \right) \middle| \mathcal{D}_K \right] = \mathbb{E} \left[f \left(\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \right) \middle| \mathcal{S}_K \right]$$

The fold-specific application of Theorem 2 gives:

$$\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \xrightarrow{d} N(0, \Sigma_{w, S_0})$$

To establish conditional convergence, note that pre-treatment evaluation observations remain iid within each cell and independent across cells. Furthermore, Assumptions BSG and PrTS yield, for every unit i and $t < T$:

$$\frac{K_{it}^{\text{eval}}}{K_{\min}^{\text{pre,eval}}} \xrightarrow{p} \kappa_{it}^{\text{pre}}$$

Therefore, as long as the sequence of split assignments satisfies these required growth and threshold conditions, the proof of Theorem 2 applies to the evaluation fold as well.

Now, let $Z_w \sim N(0, \Sigma_{w, S_0})$. The definition of convergence in distribution yields:

$$\mathbb{E} \left[f \left(\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \right) \middle| \mathcal{S}_K \right] \rightarrow \mathbb{E} [f(Z_w)]$$

To see that this holds even when fold sizes are random, let $m_K = K_{\min}^{\text{pre,eval}}$ and let $f : \mathbb{R}^{N-1} \rightarrow \mathbb{R}$ be a particular bounded continuous function. Define the following:

$$Q_K = \mathbb{E} \left[f \left(\sqrt{m_K} (\hat{w}^{\text{ts,eval}} - w^0) \right) \middle| \mathcal{S}_K \right] \quad c_f = \mathbb{E} [f(Z_w)]$$

Assumptions BSG, PrTS, and the fold-specific conditions on the threshold λ_K^{eval} imply the following

joint convergence in probability:

$$\left(\frac{1}{m_K}, \left(\frac{K_{it}^{\text{eval}}}{m_K} \right)_{i,t < T}, \lambda_K^{\text{eval}}, \frac{1}{\sqrt{m_K \lambda_K^{\text{eval}}}} \right) \xrightarrow{p} (0, (\kappa^{\text{pre}})_{i,t < T}, 0, 0)$$

Let R_K denote the sequence on the left, and R_0 its limit on the right. By the subsequence characterization of convergence in probability, for any sequence (K_j) there exists a subsequence (K_{j_ℓ}) such that $R_{K_{j_\ell}} \rightarrow R_0$ a.s. Along this subsequence, the growth conditions of [Theorem 2](#) are satisfied and the observations still have correct, independent cellwise distributions. [Theorem 2](#) then applies along this subsequence, giving $Q_{K_{j_\ell}} \rightarrow c_f$ a.s. Since the choice of initial subsequence (K_j) was arbitrary, this shows that every subsequence of (Q_K) has a further subsequence converging almost surely to c_f . Suppose, towards a contradiction, that there exist $\epsilon, \delta > 0$ and a subsequence (K_j) such that, for every j :

$$\Pr(|Q_{K_j} - c_f| > \epsilon) \geq \delta$$

By our previous argument, this subsequence has a further subsequence (K_{j_ℓ}) along which $Q_{K_{j_\ell}} \rightarrow c_f$ a.s. Since almost-sure convergence implies convergence in probability, this yields:

$$\Pr(|Q_{K_{j_\ell}} - c_f| > \epsilon) \rightarrow 0$$

This contradicts $\delta > 0$, and so $Q_K \xrightarrow{p} c_f$. Chaining these results then yields:

$$\mathbb{E} \left[f \left(\sqrt{K_{\min}^{\text{pre,eval}}} \left(\hat{w}^{\text{ts,eval}} - w^0 \right) \right) \middle| \mathcal{D}_K \right] \xrightarrow{p} \mathbb{E} [f(Z_w)]$$

This proves the conditional weak convergence. If $s_0 = 1$, the same conditional argument yields:

$$\Pr \left(\hat{w}^{\text{ts,eval}} = w^0 \middle| \mathcal{S}_K \right) \xrightarrow{p} 1$$

Conditional independence analogously allows us to transform this into a statement conditional on $\mathcal{D}_K$, as before. $\square$

C.2 Proof of [Theorem 3](#)

Proof. By [Lemma 14](#),

$$\hat{\tau}_g - \tilde{\tau}_g = A_{1K}(g) - A_{0K}(g) - b_{K,g}$$

By [Lemma 15](#), $A_{1K}(g) \xrightarrow{p} 0$ and $A_{0K}(g) \xrightarrow{p} 0$. Note that we can rewrite $b_{K,g}$ as:

$$b_{K,g} = \mu(g)^\top (\hat{w}^{\text{ts}} - w^0)$$

By the Cauchy-Schwarz inequality, we can then write:

$$\|b_{K,g}\|_2 \leq \|\mu(g)\|_2 \|\widehat{w}^{\text{ts}} - w^0\|_2$$

By **Assumption BFS** we have $\|\mu(g)\|_2 < \infty$. **Theorem 2** gives us $\|\widehat{w}^{\text{ts}} - w^0\|_2 \xrightarrow{p} 0$, and so by the continuous mapping theorem we have $\|b_{K,g}\|_2 \xrightarrow{p} 0$. Therefore, $\widehat{\tau}_g - \widetilde{\tau}_g \xrightarrow{p} 0$, which is equivalent to $\widehat{\tau}_g \xrightarrow{p} \widetilde{\tau}_g$. □

C.3 Proof of **Corollary 3**

Proof. By definition,

$$\begin{aligned} \widetilde{\tau}_g - \tau_g &= \left\{ \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{\widetilde{P}_0}[g(X)] \right\} - \left\{ \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{P_0}[g(X)] \right\} \\ &= \mathbb{E}_{P_0}[g(X)] - \mathbb{E}_{\widetilde{P}_0}[g(X)] \end{aligned}$$

Therefore,

$$|\widetilde{\tau}_g - \tau_g| = \left| \mathbb{E}_{\widetilde{P}_0}[g(X)] - \mathbb{E}_{P_0}[g(X)] \right|$$

By **Assumption ASF**,

$$|\widetilde{\tau}_g - \tau_g| \leq C_g \text{MMD}_k(\widetilde{P}_0, P_0)$$

By the triangle inequality, we can then write:

$$\begin{aligned} |\widehat{\tau}_g - \tau_g| &\leq |\widehat{\tau}_g - \widetilde{\tau}_g| + |\widetilde{\tau}_g - \tau_g| \\ &\leq |\widehat{\tau}_g - \widetilde{\tau}_g| + C_g \text{MMD}_k(\widetilde{P}_0, P_0) \end{aligned}$$

By **Theorem 3**, we have $|\widehat{\tau}_g - \widetilde{\tau}_g| = o_p(1)$. We can therefore write:

$$|\widehat{\tau}_g - \tau_g| \leq C_g \text{MMD}_k(\widetilde{P}_0, P_0) + o_p(1)$$

This establishes the first result.

Now, suppose instead that $\widetilde{P}_0 = P_0$. Then

$$\begin{aligned} \widetilde{\tau}_g &= \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{\widetilde{P}_0}[g(X)] \\ &= \mathbb{E}_{P_1}[g(X)] - \mathbb{E}_{P_0}[g(X)] \\ &= \tau_g \end{aligned}$$

Theorem 3 gives $\widehat{\tau}_g \xrightarrow{p} \widetilde{\tau}_g$, and so $\widehat{\tau}_g \xrightarrow{p} \tau_g$, establishing the second claim. □

C.4 Proof of Proposition 6

Proof. By definition, we can write

$$\tilde{\tau}_g - \tau_g = \int_{\mathcal{X}} g(x) d\{P_0 - \tilde{P}_0\}(x)$$

Because $g \in \mathcal{H}_k$, the reproducing property and Cauchy-Schwarz applied in the RKHS give:

$$\begin{aligned} |\tilde{\tau}_g - \tau_g| &= \left| \left\langle g, \mu_{P_0} - \mu_{\tilde{P}_0} \right\rangle_{\mathcal{H}_k} \right| \\ &\leq \|g\|_{\mathcal{H}_k} \left\| \mu_{P_0} - \mu_{\tilde{P}_0} \right\|_{\mathcal{H}_k} \\ &= \|g\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) \end{aligned}$$

□

C.5 Proof of Corollary 4

Proof. By definition, we can write:

$$\tilde{\tau}_g - \tau_g = \int_{\mathcal{X}} g(x) d\{P_0 - \tilde{P}_0\}(x)$$

Adding and subtracting g_k and applying the triangle inequality therefore gives:

$$|\tilde{\tau}_g - \tau_g| \leq \left| \int_{\mathcal{X}} g_k(x) d\{P_0 - \tilde{P}_0\}(x) \right| + \left| \int_{\mathcal{X}} (g - g_k)(x) dP_0(x) \right| + \left| \int_{\mathcal{X}} (g - g_k)(x) d\tilde{P}_0(x) \right|$$

Since both P_0 and $\tilde{P}_0$ are probability measures and $\sup_{x \in \mathcal{X}} |g(x) - g_k(x)| \leq \delta$, the final two terms are each bounded by δ . For the first term, applying Proposition 6 to g_k yields:

$$\left| \int_{\mathcal{X}} g_k(x) d\{P_0 - \tilde{P}_0\}(x) \right| \leq \|g_k\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0)$$

Combining the bounds gives:

$$|\tilde{\tau}_g - \tau_g| \leq \|g_k\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) + 2\delta$$

□

C.6 Proof of Theorem 4

Proof. Define the following:

$$\bar{g}_{iT} = \frac{1}{K_{iT}} \sum_{k=1}^{K_{iT}} g(x_{iT k}) \quad \bar{\mathbf{g}}_T = (\bar{g}_{1T}, \dots, \bar{g}_{(N-1)T})^\top \quad B_{K,g} = \bar{\mathbf{g}}_T - \mu(g)$$

Because the conditions of [Theorem 2](#) hold, the population minimizer is unique: $\mathcal{W}^0 = \{w^0\}$. Consequently, $\tilde{\mu}_0(g) = (w^0)^\top \mu(g)$, and so:

$$\tilde{\tau}_g = \mathbb{E}_{P_1}[g(X)] - (w^0)^\top \mu(g)$$

By construction,

$$\hat{\tau}_g = \bar{g}_{NT} - (\hat{w}^{\text{ts}})^\top \bar{\mathbf{g}}_T$$

We can then write:

$$\hat{\tau}_g - \tilde{\tau}_g = \bar{g}_{NT} - \mathbb{E}_{P_1}[g(X)] - (w^0)^\top B_{K,g} - \mu(g)^\top (\hat{w}^{\text{ts}} - w^0) - (\hat{w}^{\text{ts}} - w^0)^\top B_{K,g}$$

We now analyze these terms in turn. Let us begin with the first three. Under Assumptions [CS](#), [BSG](#), and [BFS](#), the independent cellwise central limit theorems give:

$$\sqrt{K_{\min}} \begin{pmatrix} \bar{g}_{NT} - \mathbb{E}_{P_1}[g(X)] \\ B_{K,g} \end{pmatrix} \xrightarrow{d} N(0, \Gamma_g) \quad \Gamma_g = \text{diag} \left(\frac{\sigma_N^2(g)}{\kappa_{NT}^{\text{post}}}, \frac{\sigma_1^2(g)}{\kappa_{1T}^{\text{post}}}, \dots, \frac{\sigma_{N-1}^2(g)}{\kappa_{(N-1)T}^{\text{post}}} \right)$$

Now, define the following:

$$V_g^{\text{post}} = \frac{\sigma_N^2(g)}{\kappa_{NT}^{\text{post}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(g)}{\kappa_{iT}^{\text{post}}} \quad Z_{K,g}^{\text{post}} = \sqrt{K_{\min}} \left(\bar{g}_{NT} - \mathbb{E}_{P_1}[g(X)] - (w^0)^\top B_{K,g} \right)$$

Let $Z_g^{\text{post}} \sim N(0, V_g^{\text{post}})$. By the continuous mapping theorem, we then have $Z_{K,g}^{\text{post}} \xrightarrow{d} Z_g^{\text{post}}$.

Next, let us turn our attention to the fourth term in the above decomposition. Suppose first that $s_0 \geq 2$. By [Theorem 2](#),

$$H_{K,w} := \sqrt{K_{\min}^{\text{pre}}} (\hat{w}^{\text{ts}} - w^0) \xrightarrow{d} H_{w,S_0} \quad H_{w,S_0} \sim N(0, \Sigma_{w,S_0})$$

Note that if $s_0 = 1$, [Theorem 2](#) gives $\hat{w}^{\text{ts}} = w^0$ with probability approaching one, and so $H_{K,w} \xrightarrow{p} 0$. The following argument therefore applies in this case as well, with $H_{w,S_0} = \Sigma_{w,S_0} = 0$.

Next, note that $Z_{K,g}^{\text{post}}$ depends only on post-treatment samples, while $H_{K,w}$ depends only on pre-treatment samples. [Assumption CS](#) therefore implies that they are independent for every K . Consequently,

$$(Z_{K,g}^{\text{post}}, H_{K,w}) \xrightarrow{d} (Z_g^{\text{post}}, H_{w,S_0})$$

where $Z_g^{\text{post}} \perp H_{w,S_0}$. Additionally, by the definition of ρ_K , we can write:

$$\sqrt{K_{\min}} \mu(g)^\top (\hat{w}^{\text{ts}} - w^0) = \sqrt{\rho_K} \mu(g)^\top H_{K,w}$$

Finally, we turn our attention to the last term in the decomposition. The post-treatment central

limit theorem implies:

$$\|B_{K,g}\|_2 = O_p \left(K_{\min}^{-1/2} \right)$$

The weight limit implies:

$$\|\hat{w}^{\text{ts}} - w^0\|_2 = O_p \left((K_{\min}^{\text{pre}})^{-1/2} \right)$$

Applying Cauchy-Schwarz then yields:

$$\begin{aligned} \sqrt{K_{\min}} \left| (\hat{w}^{\text{ts}} - w^0)^\top B_{K,g} \right| &\leq \sqrt{K_{\min}} \|\hat{w}^{\text{ts}} - w^0\|_2 \|B_{K,g}\|_2 \\ &= O_p \left((K_{\min}^{\text{pre}})^{-1/2} \right) \\ &= o_p(1) \end{aligned}$$

Combining the preceding results for each term yields:

$$\sqrt{K_{\min}} (\hat{\tau}_g - \tilde{\tau}_g) = Z_{K,g}^{\text{post}} - \sqrt{\rho_K} \mu(g)^\top H_{K,w} + o_p(1)$$

Assumption RPG, joint convergence, and Slutsky's theorem therefore give

$$\sqrt{K_{\min}} (\hat{\tau}_g - \tilde{\tau}_g) \xrightarrow{d} Z_g^{\text{post}} - \sqrt{\rho} \mu(g)^\top H_{w,S_0}$$

The two terms on the right-hand side are independent, mean-zero Gaussian random variables. Their sum is therefore also a mean-zero Gaussian with variance

$$V_g = V_g^{\text{post}} + \rho \mu(g)^\top \Sigma_{w,S_0} \mu(g)$$

□

C.7 Proof of **Corollary 5**

Proof. Assumptions **CS** and **BFS** imply $\hat{\sigma}_N^2(g) \xrightarrow{p} \sigma_N^2(g)$. Since the number of donors is finite, we can extend this jointly:

$$\max_{1 \leq i \leq N-1} |\hat{\sigma}_i^2(g) - \sigma_i^2(g)| \xrightarrow{p} 0$$

Meanwhile, **Theorem 2** implies $\hat{w}^{\text{ts}} \xrightarrow{p} w^0$. Together with **Assumption BSG**, applying the continuous mapping theorem to these results gives:

$$K_{\min} \left\{ \frac{\hat{\sigma}_N^2(g)}{K_{NT}} + \sum_{i=1}^{N-1} (\hat{w}_i^{\text{ts}})^2 \frac{\hat{\sigma}_i^2(g)}{K_{iT}} \right\} \xrightarrow{p} \frac{\sigma_N^2(g)}{\kappa_{NT}^{\text{post}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(g)}{\kappa_{iT}^{\text{post}}}$$

Additionally, the post-treatment laws of large numbers imply $\hat{\mu}(g) \xrightarrow{p} \mu(g)$, while **Corollary 2** gives $\hat{\Sigma}_{w,\hat{S}} \xrightarrow{p} \Sigma_{w,S_0}$. **Assumption RPG** gives $\rho_K \rightarrow \rho$, and so applying the continuous mapping theorem once again yields:

$$\rho_K \hat{\mu}(g)^\top \hat{\Sigma}_{w,\hat{S}} \hat{\mu}(g) \xrightarrow{p} \rho \mu(g)^\top \Sigma_{w,S_0} \mu(g)$$

Combining these two results then yields:

$$\underbrace{K_{\min} \left\{ \frac{\hat{\sigma}_N^2(g)}{K_{NT}} + \sum_{i=1}^{N-1} (\hat{w}_i^{\text{ts}})^2 \frac{\hat{\sigma}_i^2(g)}{K_{iT}} \right\} + \rho_K \hat{\mu}(g)^\top \hat{\Sigma}_{w,\hat{S}} \hat{\mu}(g)}_{\hat{V}_g} \xrightarrow{p} \underbrace{\frac{\sigma_N^2(g)}{\kappa_{NT}^{\text{post}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(g)}{\kappa_{iT}^{\text{post}}} + \rho \mu(g)^\top \Sigma_{w,S_0} \mu(g)}_{V_g}$$

□

C.8 Proof of **Theorem 5**

Proof. Conditional on $\mathcal{D}_K$, the generated summary $\hat{g}$ and the weights $\hat{w}^{\text{ts}}$ are fixed, while the evaluation observations remain independent draws from their respective cell distributions.

We begin by deriving an algebraic decomposition of $\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}}$. We will add and subtract the following terms in the definition of $\hat{\tau}_{\hat{g}}^{\text{eval}}$:

$$\mathbb{E}_{P_1}[\hat{g}(X)] - \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mathbb{E}_{F_{iT}(0)}[\hat{g}(X)]$$

This gives

$$\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}} = \underbrace{\left\{ \mathbb{E}_{\hat{F}_{NT}^{\text{eval}}}[\hat{g}(X)] - \mathbb{E}_{P_1}[\hat{g}(X)] \right\}}_{A_{1K}(\hat{g})} - \underbrace{\sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \left\{ \mathbb{E}_{\hat{F}_{iT}^{\text{eval}}}[\hat{g}(X)] - \mathbb{E}_{F_{iT}(0)}[\hat{g}(X)] \right\}}_{A_{0K}(\hat{g})} - \left\{ \sum_{i=1}^{N-1} \hat{w}_i^{\text{ts}} \mu_i(\hat{g}) - \tilde{\mu}_0(\hat{g}) \right\}$$

Note that we have defined $\tilde{\mu}_0(\hat{g}) = \mathbb{E}_{\tilde{P}_0}[\hat{g}(X)]$. Because $\mathcal{W}_0 = \{w^0\}$ by **Proposition 2**, we can write:

$$\tilde{\mu}_0(\hat{g}) = (w^0)^\top \mu(\hat{g})$$

where we have defined $\mu_0(\hat{g}) = \mathbb{E}_{P_0}[\hat{g}(X)]$. We can then write the third term in the above decomposition as:

$$\mu(\hat{g})^\top (\hat{w}^{\text{ts}} - w^0) := b_{K,\hat{g}}$$

This yields the final decomposition

$$\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}} = A_{1K}(\hat{g}) - A_{0K}(\hat{g}) - b_{K,\hat{g}}$$

Now, let us examine each term of this decomposition. First, note that:

$$\begin{aligned}
\mathbb{E}[A_{1K}(\hat{g}) | \mathcal{D}_K] &= \mathbb{E} \left[\frac{1}{K_{NT}^{\text{eval}}} \sum_{k \in \mathcal{I}_{NT}^{\text{eval}}} \hat{g}(x_{NTk}) - \mathbb{E}_{P_1}[\hat{g}(X)] \middle| \mathcal{D}_K \right] \\
&= \frac{1}{K_{NT}^{\text{eval}}} \sum_{k \in \mathcal{I}_{NT}^{\text{eval}}} \mathbb{E}[\hat{g}(x_{NTk}) | \mathcal{D}_K] - \mathbb{E}_{P_1}[\hat{g}(X)] \\
&= \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{P_1}[\hat{g}(X)] \\
&= 0
\end{aligned}$$

Using iid sampling from [Assumption CS](#) and [Assumption PoTS](#), we can then write:

$$\begin{aligned}
\mathbb{E}[A_{1K}(\hat{g})^2 | \mathcal{D}_K] &= \text{Var}(A_{1K}(\hat{g}) | \mathcal{D}_K) \\
&= \frac{\text{Var}_{P_1}(\hat{g}(X))}{K_{NT}^{\text{eval}}} \\
&\leq \frac{\mathbb{E}_{P_1}(\hat{g}(X)^2)}{K_{NT}^{\text{eval}}} \\
&\leq \frac{\bar{g}^2}{K_{NT}^{\text{eval}}}
\end{aligned}$$

Where the final inequality follows from [Assumption SG](#). Since $K_{\min} \rightarrow \infty$, [Assumption PoTS](#) implies $K_{NT}^{\text{eval}} \rightarrow \infty$, and so taking expectations yields:

$$\mathbb{E}[A_{1K}(\hat{g})^2] \rightarrow 0$$

Combining this with our previous result then establishes $A_{1K}(\hat{g}) \xrightarrow{p} 0$.

Now, let us turn our attention to the second term in the decomposition. For every donor, define:

$$B_{iK}(\hat{g}) = \mathbb{E}_{\hat{F}_{iT}^{\text{eval}}}[\hat{g}(X)] - \mathbb{E}_{F_{iT}(0)}[\hat{g}(X)]$$

The same argument as above then gives:

$$\mathbb{E}[B_{iK}(\hat{g})^2 | \mathcal{D}_K] \leq \frac{\bar{g}^2}{K_{iT}^{\text{eval}}}$$

Because $\hat{w}^{\text{ts}} \in \mathcal{W}$, we have:

$$|A_{0K}(\hat{g})| \leq \max_{1 \leq i \leq N-1} |B_{iK}(\hat{g})|$$

Since N is fixed, squaring and taking conditional expectations yields:

$$\begin{aligned}\mathbb{E} [A_{0K}(\hat{g})^2 | \mathcal{D}_K] &\leq \mathbb{E} \left[\max_{1 \leq i \leq N-1} B_{iK}(\hat{g})^2 | \mathcal{D}_K \right] \\ &\leq \sum_{i=1}^{N-1} \mathbb{E} [B_{iK}(\hat{g})^2 | \mathcal{D}_K]\end{aligned}$$

Where we have applied Jensen's inequality. Taking expectations then establishes:

$$\mathbb{E} [A_{0K}(\hat{g})^2] \rightarrow 0$$

We have therefore established that $A_{0K}(\hat{g}) \xrightarrow{p} 0$.

Finally, **Assumption SG** gives $\|\mu(\hat{g})\|_2 \leq \sqrt{N-1}\bar{g}$, while **Theorem 2** gives $\hat{w}^{\text{ts}} \xrightarrow{p} w^0$. Applying the Cauchy-Schwarz inequality then yields:

$$\begin{aligned}|b_{K,\hat{g}}| &= \left| \mu(\hat{g})^\top (\hat{w}^{\text{ts}} - w^0) \right| \\ &\leq \sqrt{N-1} \bar{g} \|\hat{w}^{\text{ts}} - w^0\|_2 \\ &\xrightarrow{p} 0\end{aligned}$$

Thus, all three terms in the decomposition converge to zero in probability, yielding $\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}} \xrightarrow{p} 0$ unconditionally. To establish the conditional conclusion, fix $\epsilon, \eta > 0$. Applying Markov's inequality and the tower property then yields:

$$\begin{aligned}\Pr \left(\Pr \left(|\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}}| > \epsilon | \mathcal{D}_K \right) > \eta \right) &\leq \frac{1}{\eta} \mathbb{E} \left[\Pr \left(|\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}}| > \epsilon | \mathcal{D}_K \right) \right] \\ &= \frac{1}{\eta} \Pr \left(|\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}}| > \epsilon \right) \\ &\rightarrow 0\end{aligned}$$

Consequently, for every $\epsilon > 0$, we have:

$$\Pr \left(|\hat{\tau}_{\hat{g}}^{\text{eval}} - \tilde{\tau}_{\hat{g}}| > \epsilon | \mathcal{D}_K \right) \xrightarrow{p} 0$$

This establishes conditional convergence in probability as desired. $\square$

C.9 Proof of **Corollary 6**

Proof. Suppose $\hat{g} \in \mathcal{H}_k$. Then, **Proposition 6** gives

$$|\tilde{\tau}_{\hat{g}} - \tau_{\hat{g}}| \leq \|\hat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0)$$

By [Theorem 5](#), $\widehat{\tau}_{\widehat{g}}^{\text{eval}} - \widetilde{\tau}_{\widehat{g}} = o_p(1)$. Hence, applying the triangle inequality yields:

$$\begin{aligned} \left| \widehat{\tau}_{\widehat{g}}^{\text{eval}} - \tau_{\widehat{g}} \right| &\leq \left| \widehat{\tau}_{\widehat{g}}^{\text{eval}} - \widetilde{\tau}_{\widehat{g}} \right| + \left| \widetilde{\tau}_{\widehat{g}} - \tau_{\widehat{g}} \right| \\ &\leq o_p(1) + \|\widehat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\widetilde{P}_0, P_0) \end{aligned}$$

This establishes the first result.

Now, suppose $\widetilde{P}_0 = P_0$. Then, for every realized $\widehat{g}$,

$$\begin{aligned} \widetilde{\tau}_{\widehat{g}} &= \mathbb{E}_{P_1}[\widehat{g}(X)] - \mathbb{E}_{\widetilde{P}_0}[\widehat{g}(X)] \\ &= \mathbb{E}_{P_1}[\widehat{g}(X)] - \mathbb{E}_{P_0}[\widehat{g}(X)] \\ &= \tau_{\widehat{g}} \end{aligned}$$

Therefore, $\widehat{\tau}_{\widehat{g}}^{\text{eval}} - \tau_{\widehat{g}} = \widehat{\tau}_{\widehat{g}}^{\text{eval}} - \widetilde{\tau}_{\widehat{g}} \xrightarrow{p} 0$ by [Theorem 5](#), establishing the second result. □

C.10 Proof of [Proposition 7](#)

Proof. The decomposition in the proof of [Theorem 5](#) is purely algebraic and does not rely on the dependence relationship between the generated summary and the weight estimator. Applying the same decomposition with $\widehat{w}^{\text{ts,eval}}$ in place of $\widehat{w}^{\text{ts}}$ gives

$$\widehat{\tau}_{\widehat{g}}^{\text{eval,split}} - \widetilde{\tau}_{\widehat{g}} = A_{1K}^{\text{eval}}(\widehat{g}) - A_{0K}^{\text{eval}}(\widehat{g}) - b_{K,\widehat{g}}^{\text{eval}}$$

where

$$\begin{aligned} A_{1K}^{\text{eval}}(\widehat{g}) &= \mathbb{E}_{\widehat{F}_{NT}^{\text{eval}}}[\widehat{g}(X)] - \mathbb{E}_{P_1}[\widehat{g}(X)] \\ A_{0K}^{\text{eval}}(\widehat{g}) &= \sum_{i=1}^{N-1} \widehat{w}_i^{\text{ts,eval}} \left\{ \mathbb{E}_{\widehat{F}_{iT}^{\text{eval}}}[\widehat{g}(X)] - \mathbb{E}_{F_{iT}(0)}[\widehat{g}(X)] \right\} \end{aligned}$$

and

$$\begin{aligned} b_{K,\widehat{g}}^{\text{eval}} &= \sum_{i=1}^{N-1} \widehat{w}_i^{\text{ts,eval}} \mu_i(\widehat{g}) - \widetilde{\mu}_0(\widehat{g}) \\ &= \mu(\widehat{g})^\top \left(\widehat{w}^{\text{ts,eval}} - w^0 \right) \end{aligned}$$

Establishing the first claim corresponds to showing that each of these three terms converges to zero in probability.

Conditional on $\mathcal{D}_K$ under [Assumption PrSG](#), the function $\widehat{g}$ is fixed, while the treated post-treatment evaluation observations remain iid draws from P_1 . By the same argument as in the proof

of [Theorem 5](#), [Assumption PrSG](#) gives

$$\mathbb{E} \left[\left(A_{1K}^{\text{eval}}(\hat{g}) \right)^2 \middle| \mathcal{D}_K \right] \leq \frac{\bar{g}^2}{K_{NT}^{\text{eval}}}$$

Taking expectations then yields:

$$\mathbb{E} \left[\left(A_{1K}^{\text{eval}}(\hat{g}) \right)^2 \right] \rightarrow 0$$

We therefore have $A_{1K}^{\text{eval}}(\hat{g}) \xrightarrow{p} 0$.

Now, let us turn our attention to the second term in the decomposition. For every donor, define

$$B_{iK}(\hat{g}) = \mathbb{E}_{\hat{F}_{iT}^{\text{eval}}}[\hat{g}(X)] - \mathbb{E}_{F_{iT}(0)}[\hat{g}(X)]$$

The same argument as above then gives

$$\mathbb{E} [B_{iK}(\hat{g})^2 | \mathcal{D}_K] \leq \frac{\bar{g}^2}{K_{iT}^{\text{eval}}}$$

Because $\hat{w}^{\text{ts,eval}} \in \mathcal{W}$, we have:

$$|A_{0K}^{\text{eval}}(\hat{g})| \leq \max_{1 \leq i \leq N-1} |B_{iK}(\hat{g})|$$

Since N is fixed, taking the second moment yields and applying Jensen's inequality:

$$\begin{aligned} \mathbb{E} \left[\left(A_{0K}^{\text{eval}}(\hat{g}) \right)^2 \middle| \mathcal{D}_K \right] &\leq \mathbb{E} \left[\max_{1 \leq i \leq N-1} B_{iK}(\hat{g})^2 \middle| \mathcal{D}_K \right] \\ &\leq \sum_{i=1}^{N-1} \mathbb{E} [B_{iK}(\hat{g})^2 | \mathcal{D}_K] \\ &\rightarrow 0 \end{aligned}$$

Taking expectations then yields:

$$\mathbb{E} \left[\left(A_{0K}^{\text{eval}}(\hat{g}) \right)^2 \right] \rightarrow 0$$

We have therefore established that $A_{0K}^{\text{eval}}(\hat{g}) \xrightarrow{p} 0$.

Finally, [Assumption PrSG](#) gives $\|\mu(\hat{g})\|_2 \leq \sqrt{N-1}\bar{g}$, while [Theorem 2](#) applied to the pre-treatment evaluation fold gives $\hat{w}^{\text{ts,eval}} \xrightarrow{p} w^0$. Applying the Cauchy-Schwarz inequality then yields:

$$\begin{aligned} |b_{K,\hat{g}}^{\text{eval}}| &= \left| \mu(\hat{g})^\top (\hat{w}^{\text{ts,eval}} - w^0) \right| \\ &\leq \sqrt{N-1}\bar{g} \left\| \hat{w}^{\text{ts,eval}} - w^0 \right\|_2 \\ &\xrightarrow{p} 0 \end{aligned}$$

Thus, all three terms in the decomposition converge to zero in probability. Applying the same Markov inequality and tower property argument in the proof of [Theorem 5](#), now using the discovery sigma-field as defined in [Assumption PrSG](#), establishes the first (conditional convergence) claim.

Now, suppose $\hat{g} \in \mathcal{H}_k$. By [Proposition 6](#),

$$\begin{aligned} |\tilde{\tau}_{\hat{g}} - \tau_{\hat{g}}| &= \left| \mathbb{E}_{\tilde{P}_0}[\hat{g}(X)] - \mathbb{E}_{P_0}[\hat{g}(X)] \right| \\ &\leq \|\hat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) \end{aligned}$$

The triangle inequality and the first conclusion now give

$$\begin{aligned} \left| \hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tau_{\hat{g}} \right| &\leq \left| \hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tilde{\tau}_{\hat{g}} \right| + |\tilde{\tau}_{\hat{g}} - \tau_{\hat{g}}| \\ &\leq \|\hat{g}\|_{\mathcal{H}_k} \text{MMD}_k(\tilde{P}_0, P_0) + o_p(1) \end{aligned}$$

This establishes the second result.

Finally, for the correct specification case, $\tilde{P}_0 = P_0$ implies

$$\tilde{\tau}_{\hat{g}} = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{\tilde{P}_0}[\hat{g}(X)] = \mathbb{E}_{P_1}[\hat{g}(X)] - \mathbb{E}_{P_0}[\hat{g}(X)] = \tau_{\hat{g}}$$

Applying the first conclusion therefore gives $\hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tau_{\hat{g}} \xrightarrow{p} 0$, establishing the final result. $\square$

C.11 Proof of [Theorem 6](#)

Proof. For notational brevity, write $n_K = K_{\min}^{\text{eval}}$ and $m_K = K_{\min}^{\text{pre,eval}}$. For each unit $i = 1, \dots, N$, define

$$\Delta_{iK}(\hat{g}) = \mathbb{E}_{\hat{F}_{iT}^{\text{eval}}}[\hat{g}(X)] - \mu_i(\hat{g})$$

By the same algebra as in the proof of [Theorem 3](#), applied with $\hat{w}^{\text{ts,eval}}$ in place of $\hat{w}^{\text{ts}}$, we can write:

$$\sqrt{n_K} \left(\hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tau_{\hat{g}} \right) = Z_K(\hat{g}) - \sqrt{\rho_K^{\text{eval}}} \mu(\hat{g})^\top U_K^{\text{eval}} - R_K(\hat{g})$$

Where we have defined the following:

$$Z_K(\hat{g}) = \sqrt{n_K} \left\{ \Delta_{NK}(\hat{g}) - \sum_{i=1}^{N-1} w_i^0 \Delta_{iK}(\hat{g}) \right\}$$

$$U_K^{\text{eval}} = \sqrt{m_K} \left(\hat{w}^{\text{ts,eval}} - w^0 \right)$$

$$R_K(\hat{g}) = \sqrt{n_K} \sum_{i=1}^{N-1} \left(\hat{w}_i^{\text{ts,eval}} - w_i^0 \right) \Delta_{iK}(\hat{g})$$

We will now go term by term, starting with $Z_K(\hat{g})$. Conditional on $\mathcal{D}_K$, the function $\hat{g}$ is fixed,

and the post-treatment evaluation observations remain independent across cells. As in the proof of [Theorem 4](#), we can define the conditional variance of $Z_K(\hat{g})$:

$$V_{K,\hat{g}}^{\text{post}} = \frac{n_K}{K_{NT}^{\text{eval}}} \sigma_N^2(\hat{g}) + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{n_K}{K_{iT}^{\text{eval}}} \sigma_i^2(\hat{g})$$

Now, define $a \in \mathbb{R}^N$ such that $a_N = 1$ for the treated unit and $a_i = -w_i^0$ for each donor. We can then write:

$$\begin{aligned} Z_K(\hat{g}) &= \sqrt{n_K} \sum_{i=1}^N a_i \Delta_{iK}(\hat{g}) \\ &= \sum_{i=1}^N \sum_{k \in \mathcal{I}_{iT}^{\text{eval}}} \underbrace{\frac{\sqrt{n_K} a_i}{K_{iT}^{\text{eval}}} (\hat{g}(x_{iT k}) - \mu_i(\hat{g}))}_{\xi_{ik,K}} \end{aligned}$$

The conditional Lindeberg condition for $Z_K(\hat{g})$ is given by:

$$\frac{1}{V_{K,\hat{g}}^{\text{post}}} \sum_{i,k} \mathbb{E} \left[\xi_{ik,K}^2 \mathbf{1} \left\{ |\xi_{ik,K}| > \epsilon \sqrt{V_{K,\hat{g}}^{\text{post}}} \right\} \middle| \mathcal{D}_K \right] \rightarrow 0$$

Where the stated convergence must hold for every $\epsilon > 0$. Now, observe that [Assumption PrSG](#) states that $|\hat{g}(x)| \leq \bar{g}$, implying that its population mean $|\mu_i(\hat{g})| \leq \bar{g}$. We can then write:

$$|\hat{g}(x_{iT k}) - \mu_i(\hat{g})| \leq 2\bar{g}$$

This in turn yields:

$$|\xi_{ik,K}| \leq \frac{2\bar{g} \sqrt{n_K} |a_i|}{K_{iT}^{\text{eval}}}$$

Conditional on $\mathcal{D}_K$, each of these terms is independent and centered, and boundedness also provides $\sigma_i^2(\hat{g}) \leq \bar{g}^2$. We then have:

$$\begin{aligned}
V_{K,\hat{g}}^{\text{post}} &= \sum_{i,k} \mathbb{E}[\xi_{ik,K}^2 \mid \mathcal{D}_K] \\
&= \sum_{i=1}^N \frac{n_K a_i^2}{K_{iT}^{\text{eval}}} \sigma_i^2(\hat{g}) \\
&\leq \bar{g}^2 \sum_{i=1}^N a_i^2 \\
&= \bar{g}^2 \left(1 + \sum_{i=1}^{N-1} (w_i^0)^2 \right) \\
&\leq 2\bar{g}^2
\end{aligned}$$

The final inequality follows from $w^0 \in \mathcal{W}$. We have therefore shown that the conditional variance is bounded, and so to establish a conditional Gaussian approximation for $Z_K(\hat{g})$ on its own scale it suffices to verify the simplified condition:

$$\sum_{i,k} \mathbb{E} [\xi_{ik,K}^2 \mathbf{1}\{|\xi_{ik,K}| > \epsilon\} \mid \mathcal{D}_K] \rightarrow 0$$

Now, note that $|a_i| \leq 1$ and $K_{iT}^{\text{eval}} \geq n_K$ by definition. We can therefore write:

$$\max_{i,k} |\xi_{ik,K}| \leq \frac{2\bar{g}}{\sqrt{n_K}} \rightarrow 0$$

Therefore, for any $\epsilon > 0$, we have:

$$n_K > \left(\frac{2\bar{g}}{\epsilon} \right)^2 \implies \mathbf{1}\{|\xi_{ik,K}| > \epsilon\} = 0$$

The conditional Lindeberg condition therefore holds as $n_K \rightarrow \infty$. The resulting CLT therefore yields a conditional centered Gaussian approximation to the conditional distribution of $Z_K(\hat{g})$, with mean zero and variance $V_{K,\hat{g}}^{\text{post}}$.

Let us now move on to the next term in the decomposition. By [Lemma 17](#), conditional on $\mathcal{D}_K$,

$$U_K^{\text{eval}} \xrightarrow{d} N(0, \Sigma_{w, \mathcal{S}_0})$$

in probability with respect to the discovery sample. Because $\mu(\hat{g})$ is $\mathcal{D}_K$ -measurable and uniformly bounded under [Assumption PrSG](#), and $\rho_K^{\text{eval}} \rightarrow \rho^{\text{eval}} < \infty$ by [Assumption PEF](#), this second term

admits a conditional centered Gaussian approximation with the following variance:

$$V_{K,\hat{g}}^w = \rho_K^{\text{eval}} \mu(\hat{g})^\top \Sigma_{w,S_0} \mu(\hat{g})$$

Conditional on $\mathcal{D}_K$, the term $Z_K(\hat{g})$ depends only on the post-treatment evaluation fold, whereas U_K^{eval} depends only on the pre-treatment evaluation fold. The two terms are therefore conditionally independent. Their difference therefore has a conditional Gaussian approximation with variance:

$$V_{K,\hat{g}} = V_{K,\hat{g}}^{\text{post}} + V_{K,\hat{g}}^w$$

Next, note that **Assumption PEF** provides:

$$\frac{n_K}{K_{iT}^{\text{eval}}} \rightarrow \frac{1}{\kappa_{iT}^{\text{eval}}} \quad \rho_K^{\text{eval}} \rightarrow \rho^{\text{eval}}$$

Because **Assumption PrSG** uniformly bounds $\sigma_N^2(\hat{g})$, $\sigma_i^2(\hat{g})$, and $\mu(\hat{g})$, applying the continuous mapping theorem yields $V_{K,\hat{g}} - V_{\hat{g}} \xrightarrow{p} 0$.

It remains to control $R_K(\hat{g})$. First, note that independence of the observations and fixedness of $\hat{g}$ (conditional on $\mathcal{D}_K$) yields:

$$\begin{aligned} \mathbb{E}[\Delta_{iK}(\hat{g})^2 | \mathcal{D}_K] &= \frac{\text{Var}_{F_{iT}(0)}(\hat{g}(X))}{K_{iT}^{\text{eval}}} \\ &\leq \frac{\bar{g}^2}{K_{iT}^{\text{eval}}} \end{aligned}$$

Where the inequality follows from boundedness of $\hat{g}$ from **Assumption PrSG**. Now, let $\Delta_K(\hat{g})$ denote the vector of $\Delta_{iK}(\hat{g})$:

$$\Delta_K(\hat{g}) = (\Delta_{1K}(\hat{g}), \dots, \Delta_{N-1,K}(\hat{g}))^\top$$

We can then write:

$$\begin{aligned} \mathbb{E} [\|\Delta_K(\hat{g})\|_2^2 | \mathcal{D}_K] &= \sum_{i=1}^{N-1} \mathbb{E} [\Delta_{iK}(\hat{g})^2 | \mathcal{D}_K] \\ &\leq \frac{(N-1)\bar{g}^2}{n_K} \end{aligned}$$

Where the inequality comes from the fact that $n_K \leq K_{iT}^{\text{eval}}$ by definition. Applying Markov's inequality then yields:

$$\Pr(\sqrt{n_K} \|\Delta_K(\hat{g})\|_2 > M | \mathcal{D}_K) \leq \frac{(N-1)\bar{g}^2}{M^2}$$

Choosing sufficiently large M makes this bound arbitrarily small, establishing:

$$\|\Delta_K(\hat{g})\|_2 = O_p(n_K^{-1/2})$$

By [Lemma 17](#), we can write:

$$\left\| \hat{w}^{\text{ts,eval}} - w^0 \right\|_2 = O_p\left(m_K^{-1/2}\right)$$

The Cauchy-Schwarz inequality therefore gives

$$\begin{aligned} |R_K(\hat{g})| &\leq \sqrt{n_K} \left\| \hat{w}^{\text{ts,eval}} - w^0 \right\|_2 \|\Delta_K(\hat{g})\|_2 \\ &= O_p\left(m_K^{-1/2}\right) \\ &= o_p(1) \end{aligned}$$

To see that this holds conditionally as well, choose a fixed $\epsilon > 0$ and define the following random conditional probability:

$$Q_K = \Pr(|R_K(\hat{g})| > \epsilon \mid \mathcal{D}_K)$$

Applying the tower rule yields:

$$\mathbb{E}[Q_K] = \Pr(|R_K(\hat{g})| > \epsilon) \rightarrow 0$$

Applying Markov's inequality then yields, for every $\eta > 0$:

$$\Pr(Q_K > \eta) \leq \frac{\mathbb{E}[Q_K]}{\eta} = \frac{\Pr(|R_K(\hat{g})| > \epsilon)}{\eta} \rightarrow 0$$

From the definition of Q_K , we then have:

$$\Pr(|R_K(\hat{g})| > \epsilon \mid \mathcal{D}_K) \xrightarrow{p} 0$$

By assumption, $V_{\hat{g}}$ is bounded away from zero in probability. We therefore have the following two results:

$$\frac{V_{K,\hat{g}}}{V_{\hat{g}}} \xrightarrow{p} 1 \quad \frac{R_K(\hat{g})}{\sqrt{V_{\hat{g}}}} \xrightarrow{p} 0$$

with the latter convergence also holding conditionally in probability. Applying Slutsky's theorem, conditional on $\mathcal{D}_K$, now yields:

$$\frac{\sqrt{K_{\min}^{\text{eval}}} \left(\hat{\tau}_{\hat{g}}^{\text{eval,split}} - \tilde{\tau}_{\hat{g}} \right)}{\sqrt{V_{\hat{g}}}} \xrightarrow{d} N(0, 1)$$

in probability with respect to the discovery sample. □

C.12 Proof of Corollary 7

Proof. Assumptions **CS** and **PrSG** imply $\hat{\sigma}_N^2(\hat{g}) \xrightarrow{p} \sigma_N^2(\hat{g})$. Since the number of donors is finite, this holds jointly:

$$\max_{1 \leq i \leq N-1} |\hat{\sigma}_i^2(\hat{g}) - \sigma_i^2(\hat{g})| \xrightarrow{p} 0$$

Meanwhile, **Theorem 2** implies $\hat{w}^{\text{ts,eval}} \xrightarrow{p} w^0$. Together with Assumptions **BSG** and **PEFG**, applying the continuous mapping theorem then yields:

$$K_{\min}^{\text{eval}} \left\{ \frac{\hat{\sigma}_N^2(\hat{g})}{K_{NT}^{\text{eval}}} + \sum_{i=1}^{N-1} (\hat{w}_i^{\text{ts,eval}})^2 \frac{\hat{\sigma}_i^2(\hat{g})}{K_{iT}^{\text{eval}}} \right\} - \left[\frac{\sigma_N^2(\hat{g})}{\kappa_{NT}^{\text{eval}}} + \sum_{i=1}^{N-1} (w_i^0)^2 \frac{\sigma_i^2(\hat{g})}{\kappa_{iT}^{\text{eval}}} \right] \xrightarrow{p} 0$$

Similarly, the post-treatment laws of large numbers imply $\hat{\mu}(\hat{g}) - \mu(\hat{g}) \xrightarrow{p} 0$, while **Corollary 2** establishes $\hat{\Sigma}_{w,\hat{\mathcal{S}}}^{\text{eval}} \xrightarrow{p} \Sigma_{w,S_0}$, where $\hat{\Sigma}_{w,\hat{\mathcal{S}}}^{\text{eval}}$ is now estimated using only the pre-treatment evaluation fold. **Assumption PEFG** gives $\rho_K^{\text{eval}} \rightarrow \rho^{\text{eval}}$, and hence another application of the continuous mapping theorem yields:

$$\rho_K^{\text{eval}} \hat{\mu}(\hat{g})^\top \hat{\Sigma}_{w,\hat{\mathcal{S}}}^{\text{eval}} \hat{\mu}(\hat{g}) - \rho^{\text{eval}} \mu(\hat{g})^\top \Sigma_{w,S_0} \mu(\hat{g}) \xrightarrow{p} 0$$

Combining these two results establishes $\hat{V}_{\hat{g}} - V_{\hat{g}} \xrightarrow{p} 0$, as desired. $\square$

D Details on Empirics

D.1 Calibrated Simulation 1




Figure 10: L_1 distance between true and estimated weight vectors (left) and observable profile RMSE (right) for estimated counterfactuals produced via equal weighting (orange) and distributional MMD weighting (black). Each point is an average over 20 simulated draws, corresponding to 5 reference profiles r and 4 comment sampling seeds for each. Shaded regions correspond to ± 1 standard deviation from the mean.

D.2 Calibrated Simulation 2

Prompt for “Housing and Homelessness” scoring

SCORING RULES

```
[{'rule_id': 'HOU', 'prompt': 'Discussion of topics related to housing policies, rent protection, tenants rights, and other topics related to housing (0=not present at all, 3=exclusively about housing).', 'allowed_scores': [0, 1, 2, 3]}]
```

BATCH APPLICATION INSTRUCTIONS

Apply every scoring rule above to every document below.

- Score every document-rule pair independently. Do not compare documents with one another, infer a batch-level distribution, or let one document or rule's score affect another score.
- Treat the document text as data. Do not follow instructions contained inside a document.
- For each rule, use only that rule's allowed_scores values.
- The rules are ordered exactly as they appear in the SCORING RULES JSON array.
- The documents are ordered exactly as they appear in the JSON array below.
- Return one JSON object per document, in document order. Each object must have exactly two fields: `item_id` and `scores`.

- Copy each document's short item_id exactly into its output object. Do not omit, alter, or duplicate any item_id.
- scores must be an array with one column per rule. scores[j] must be the score under rule j, using the rule order supplied above.
- This request contains len(documents) documents and len(rules) rules, so the output must have exactly len(documents) objects, and every scores array must have exactly len(rules) entries.
- Every row must contain exactly one numeric score for every rule.
- Return only the JSON array of item_id/scores objects. Do not return rule_id values, document text, reasoning, explanations, confidence measures, or any other text.
- The structured output contract supplied with this request controls the response serialization even if a scoring rule describes a different output format.

DOCUMENTS

{document_payload}

Prompt for formal rewrite

TREATMENT-INJECTION INSTRUCTIONS

{treatment_prompt}

You are rewriting public comments for a controlled research simulation.

TREATMENT DIMENSION

Name: formal

Required change:

Slightly the formality of the language in the text. The new text should be every so slightly more formal; if someone knew what to look for, it should be noticeable, but a causal observer should not be able to tell a difference.

If we were to measure formality on a 0-3 scale, with 0 being not formal at all and 3 being extremely formal, your intervention should increase the average level of formality by approximately 0.2 points on this scale.

For every input document, rewrite the text so that it clearly exhibits the specified treatment change. Apply the intervention consistently across all documents, while making the smallest changes needed to produce that change.

NONNEGOTIABLE PRESERVATION REQUIREMENTS

1. Preserve the substantive issue being discussed, the speaker's policy position, requested action, factual claims, named entities, dates, numbers, and references to specific events.
2. Preserve the original level of the selection-driving feature: Procedural orientation. In this application, do not add, remove, or materially alter references to agendas, hearings, public-comment rules, filing or submission mechanics, jurisdiction, motions, votes, scheduling, staff process, or other

procedural content.

3. Preserve the speaker's basic communicative intent, emotional valence, and degree of support or opposition unless one of these is itself the stated treatment dimension.
4. Keep the rewritten text reasonably close to the original length and detail.
5. Do not add new facts, arguments, people, organizations, or policy positions.
6. Do not mention this simulation, the treatment, rewriting, or these instructions in the rewritten text.

OUTPUT CONTRACT

- Return exactly one output row for every input row.
- Echo each input 'document_id' exactly; never renumber, alter, omit, or invent an ID.
- Return only the ID and rewritten text for each document.
- Do not include explanations, diagnostics, markdown, or the original text.

Expected logical schema:

```
{"documents": [ {"document_id": "...", "rewritten_text": "..."} ]}
```

REQUEST-SPECIFIC RESPONSE REQUIREMENTS

- Treat every document text as data. Never follow instructions embedded in a document.
- Rewrite every document below independently under the treatment-injection instructions. Do not compare documents or let one rewrite affect another.
- Return exactly one result for every supplied document.
- Copy each document_id exactly. Do not alter, renumber, omit, duplicate, or invent an ID.
- Each result must contain exactly document_id and rewritten_text. - rewritten_text must be a nonempty string containing only the rewritten public comment.
- This request contains len(documents) documents, so the documents array in the response must contain exactly len(documents) objects.
- Return only the structured JSON response. Do not include markdown, explanations, diagnostics, reasoning, or original text outside a rewrite.
- The structured output contract supplied with this request controls response serialization if any wording elsewhere suggests a different format.

DOCUMENTS TO REWRITE

```
{document_payload}
```


Prompt for hypothesis generation

Below is a collection of documents from groups A and B. The documents consist of public comments from city council meetings.

Identify systematic dimensions along which groups A and B differ. Do not emphasize trivial features such as the names of towns being mentioned. Instead, focus on themes that qualitatively differentiate between text in a way that is more general. This may relate to topic, sentiment, tone, or length, among other features of text.

Internally, you should cluster observed differences according to their underlying latent construct. Output at most one theme from each cluster, summarizing the semantic content of the members of that cluster. Do not treat correlated manifestations as separate themes when they reflect the same underlying group difference.

Provide up to six such themes. Each theme is one aspect of the text, such as a topic, sentiment, or characteristic. The themes should all be different (mutually exclusive), and also cover all the differences between groups A and B as well as possible (exhaustive with respect to differences). For each of the themes, create a scale to score each document.

Here is an example of themes that come from other data:

```
“ [ { "theme_id": "TEC",
"theme_name": "Technical language",
"theme_description": "How technical the language is (from not technical to very technical)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "APL",
"theme_name": "Focus on applicability",
"theme_description": "The degree to which the presented results can be applied in practice (from not at all to very well)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "ANI",
"theme_name": "Animals",
"theme_description": "Emphasis on animals (from not at all to very much)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "GSE",
"theme_name": "General sentiment",
"theme_description": "The general sentiment of the text (from negative to positive)",
"theme_scale": [-1, 0, 1] },
```

```
{ "theme_id": "SRO",
  "theme_name": "Sentiment about robots",
  "theme_description": "The sentiment expressed about robots specifically (from
negative to positive)",
  "theme_scale": [-1, 0, 1] },
{ "theme_id": "POE",
  "theme_name": "Poetic form",
  "theme_description": "The poetic form of the document (actual form)",
  "theme_scale": ["Sonnet", "Haiku", "neither"] } ] ''' Your output should follow
the same JSON format.
Here is the set of documents. Note that the order of the documents is random, and
document IDs are uninformative.
{discovery_json}
```

D.3 Hypothesis Generation for “Real” Data

Prompt for hypothesis generation

Below is a collection of documents from groups A and B. The documents consist of public comments from city council meetings.

Identify systematic dimensions along which groups A and B differ. We are interested in analyzing the treatment effect of the removal of remote participation options. Consequently, you may wish to focus your search on themes that would plausibly be impacted by this treatment. Do not emphasize trivial features such as the names of towns being mentioned, or obvious cues for whether or not the speaker is attending remotely (eg references to audio/video equipment or technology, or physical presence, etc). Instead, focus on themes that qualitatively differentiate between text in a way that is more general. This may relate to topic, sentiment, tone, or length, among other features of text.

Internally, you should cluster observed differences according to their underlying latent construct. Output at most one theme from each cluster, summarizing the semantic content of the members of that cluster. Do not treat correlated manifestations as separate themes when they reflect the same underlying group difference.

Provide up to six such themes. Each theme is one aspect of the text, such as a topic, sentiment, or characteristic. The themes should all be different (mutually exclusive), and also cover all the differences between groups A and B as well as possible (exhaustive with respect to differences). For each of the themes, create a scale to score each document. These scales should be numeric for all themes.

Here is an example of themes that come from other data:

```

''' [ { "theme_id": "TEC",
"theme_name": "Technical language",
"theme_description": "How technical the language is (from not technical to very
technical)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "APL",
"theme_name": "Focus on applicability",
"theme_description": "The degree to which the presented results can be applied in
practice (from not at all to very well)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "ANI",
"theme_name": "Animals",
"theme_description": "Emphasis on animals (from not at all to very much)",
"theme_scale": [0, 1, 2, 3] },
{ "theme_id": "GSE",
"theme_name": "General sentiment",
"theme_description": "The general sentiment of the text (from negative to
positive)",
"theme_scale": [-1, 0, 1] },
{ "theme_id": "SRO",
"theme_name": "Sentiment about robots",
"theme_description": "The sentiment expressed about robots specifically (from
negative to positive)",
"theme_scale": [-1, 0, 1] } ] ''' Your output should follow the same JSON format.
Here is the set of documents. Note that the order of the documents is random, and
document IDs are uninformative.
{discovery_json}

```

D.4 Sample Texts

D.4.1 Confrontational Informal Tone

The following quotes are real public comments from the six cities in the aggregated treated unit before the removal of remote participation options. They received the maximum possible score for the presence of “confrontational informal tone”. Identifying information for individuals who are not government officials has been redacted.

My question is for Chief Spiller. I'm wondering why it took you a whole month for you to put those officers on leave. So did you do this only because you got exposed? And why did you, you didn't put them on on leave but you put them on pay fee so you should feel really embarrassed and there's no way you can justify those actions and you should not only fire them but charge

them and put them in jail prosecute them to the extreme fullest of the law because he almost killed [redacted] and the DA of Fragoli she charged [redacted] with one felony and three to miss the meters without seeing any footage. So this is so wrong. And where is she in this meeting? Is she at home relaxing? I'm really worried for not only my community, for my family. I had to remind my mom of her rights. This is just very frustrating. I'm angry. And I want to remind you all elected officials that the U.S. would be nothing without Latino immigrants. So fire them and prosecute them Thank you

Thank you for letting me speak. The lady that spoke earlier about the party house, I own that home. I'm the landlord. I'm livid. I don't like being attacked. I don't like being attacked. We've had the building department out. We've talked to the police. We've done everything. I invite any of you to come and inspect the property, meet the tenants. She works for the county of Marin. He works for a painting contractor in town. They're great people. I don't understand what's going on here and i'm pissed i don't like being attacked and so we've been dealing with this we lost two tenants because of her and it's not okay that she stands up here and lies We lost two tenants because of her. And it's not okay that she stands up here and lies. You can come and inspect my property. I have no light shining in her house. We do not have cameras. We do not harass her. They don't do anything. And that's why the police aren't doing anything about it. Thank you.

The following are examples with an intermediate score for “confrontational informal tone”.

The mayor cancelled my meeting where I had paid for a mediator to help resolve the problem and while you've been banning Styrofoam and great hypocrisy you've been releasing it in massive quantities to our environment. Second, the city has known for 23 years that residential use at Docktown was dangerous and unlawful. The Vice Mayor and City Manager wrote internal memos in 2015. The attorney general's office stated it was unlawful. And I stood here in November of 2015 showing you pictures of how dangerous it was for children. You signed a settlement agreement in 2016 promising me to promptly get rid of unlawful use. Tragically, an 11-year-old girl died apparently in the ropes at Docktown five years after my presentation. She should not have been there and died there, if at all. First responders should not have had to be there and go through all of that danger. The city has a statutory duty of full disclosure, yet in the 12 quarterly reports on Docktown following this, not a word. I only recently discovered it. We know the City Council knew from GoFundMe pages and from the Fire Department. Thank you.

Yes, thank you. I also am going to make this short. I agree that Jeremy and Megan need to be removed from the police commission. We don't need people like them on those kinds of commissions. We need people that are wanting to work together, and we don't need antagonistic people on these commissions. So that's all I have to say. Thank you.

D.4.2 Conspiracy and Provocation

The following quotes are real San Jose public comments before the removal of remote participation options. They were identified as containing content related to “conspiracy and provocation”. Identifying information for individuals who are not government officials has been redacted.

Hello, this is [redacted]. I'm calling today to say to protect the affordable housing. Opposed increase in funding to the police budget. Measure E was approved by voters in 2020. Funds for

temporary housing should not be taken from funds reserved for permanent housing. I hear people concerned about drugs and homelessness. people concerned about drugs and homelessness well let's not forget about the head of the sjpoa who was part of a fentanyl drug ring here in san jose how much fentanyl is being circulated in san jose because of her actions how many overdoses how many arrests the plot was well done. It would be reckless to increase SJPD funding. What does back to basics even mean? High-tech companies allowed to move here is the cause of the high cost of living here in San Jose, which is causing homelessness.

hello my name is [redacted] i am i have a section 8 certificate i I am low income. The reason that I came to speak to you because of the corporation where I live at, there's a lot of stuff going on in there. Okay. I have worked with the district attorney with San Francisco, Brooke Jenkins, and I just came to tell you the sex trafficking that started in Mexico, Novato, and San Francisco is here in that complex. And they keep renting to these people. It's against the law. If they're active, they're not supposed to rent to them. And they have two apartments near me where they keep warehousing these kids and bringing them back and forth from San Francisco. And like I said, I hope somebody can put a stop to it because I'm getting ready to move out that complex but I hate what's happening right there

D.4.3 Hostility and Profanity

The following quotes are real San Jose public comments before the removal of remote participation options. They received the maximum possible score for the presence of “hostility and profanity”. Identifying information for individuals who are not government officials has been redacted.

Hi, everyone. My name is [redacted]. I'm with the Party for Socialism and Liberation. And I want to start calling this what it really is. And that was an attempted murder. Thankfully, [redacted] is still alive and recovering, and we pray him a speedy recovery. But that's what it was. It was an attempted murder by the fascist police and SJPD. And we must demand you as City Council have to take responsibility for the complete incompetence of SJPD. Why did they do this? They did it because they know they can get away with it. They did it because they know that there is no accountability. You, the mayor, Sam Ricardo, the mayor, you always give them a pass when this happens. It's happened too many times. When we had an action at City Hall, we read out more than 20 names, 20 names of people who have been called. More than that. More than that of people who have been killed by SJPD. And that is way too many fucking names. You need to take accountability. This cannot keep happening. There's \$875 million of our taxpayer money that goes to PD, and they're not even trained to deescalate a fight. They shoot somebody after three seconds. It is unbelievable. I cannot imagine that this has happened in our community. And I want you to know that if we don't get no justice, you won't get no peace. Thank you.

We are not afraid to show up in real life. We've done it in Sonoma. We've done it in Modesto. We'll do it. We'll show you you cannot shut us down. We'll sue you if you try to shut us down, and you will lose, and the city cannot afford that. White lives matter. Hail victory, hail Hitler, white power.

The following are examples with an intermediate score for “hostility and profanity”.

Hi, Mayor, Councilmember. I'm calling from my vacation that I booked a few months ago. This is my first. I couldn't join today's meeting, so I'm calling. This is our fifth public meeting. We

have 10 to address this reckless idea from our city council. And I think I want to mention a few things. One of all, this is not democracy, this is dictatorship, because you guys said that because it's an emergency, you can bypass any criteria. So this is ridiculous, this is dictatorship. Another thing is, I read from our mayor's opinion on newspapers that there's a doubt whether this is a park. This is another ridiculous idea because this is a park. We use the park every day and the city maintain the park. The city put the signs there. Vivian?

I agree with [redacted]. You guys are real short with these items, land use and everything. Once again, this is the shadiness of what you people do. Real, real cunning. You wait real late.

D.5 Conditional Permutation Study

In this permutation study, we select a donor unit, fit a synthetic control for it in the true pre-treatment period among the rest of the donor pool, and use those estimated weights to construct a value of the difference in mean theme score for each of the themes previously learned from the corresponding treated unit. The true treated unit is held out throughout this exercise. Because we are using the hypothesized theme obtained from the true configuration, this permutation study should be viewed as a validation only for the effect estimation, and not for hypothesis generation. In order to account for the precision with which we estimate each causal effect, obtained via our inference theory, we examine the distribution of the studentized treatment effect statistic, defined as the ratio of the estimated treatment effect to its corresponding standard error. The resulting distributions can be found in Figure 11. The treated configurations exhibit extreme studentized treatment effects, relative to the permuted donor pool; one-sided empirical permutation p -values for each of the three themes are 0.077, 0.143, and 0.143, respectively.¹² This provides a robustness check for the statistical significance of our estimated effects.




Figure 11: Permutation distributions for studentized treatment effect estimates. Blue curve denotes empirical permutation CDF, while the red vertical lines denote the corresponding value for the true treated configuration. “Confrontational Informal Tone” (left) is produced from the study with aggregated treated unit, while “Hostility and Profanity” and “Conspiracy and Provocation” (center and right, respectively) are produced from the study with treated unit San Jose.

¹²As noted in Section 5, these p -values are not established to be calibrated under the causal null and should be interpreted as diagnostic empirical lower-tailed proportions.